

1 Introduction

Human knowledge, both in children and in adults, seems to be organized into domains of expertise. Some of us develop competence in cooking, calculus, tennis, drawing pictures, or writing software, and virtually every child masters natural language, motor control, and intuitive theories of kinship, taxonomy, and physics. Over short timescales, people can solve new problems within their domain of expertise, and over longer time scales, they can grow into domain experts for new classes of problems. A long-term goal for machine intelligence is to build an AI that captures these two human abilities: both solving new problems, and expanding the domains of problems that can be solved in the first place.

We take the stance that an artificial agent with these humanlike abilities needs at least two ingredients. First it must acquire the underlying system of domain-specific concepts. For drawing pictures these concepts include polygons, spirals, arcs, and symmetries; for physics these concepts include dot products, vector fields, and conservation laws. Second it must learn to quickly deploy these concepts efficiently on new problems: at a glance, human domain experts can intuit which compositions of concepts are likely to solve the task at hand, even before they begin searching for a solution.

We introduce a computational model that jointly acquires these two kinds of domain knowledge. Because solutions to many kinds of problems can often be described as some kind of program [22, 45, 37], our system solves a specific problem by searching for a program. It gradually grows its domain-specific knowledge for a class of problems by assembling a library of code containing concepts useful for the domain, and by training a neural network to quickly infer, for a specific problem, what code concepts from its library are likely to solve it. We first investigate our model within classic program synthesis domains for manipulating sequences of numbers and text, and then consider visual and creative programs for drawing pictures and building towers out of toy blocks, altogether considering both deterministic and probabilistic programs that act both generatively (e.g., producing an artifact like an image or plan) and conditionally (e.g., mapping inputs to outputs).

We call our system ‘DreamCoder’ because it works through a novel kind of wake/sleep or ‘dream’ learning [16], iterating through a wake cycle – where it solves problems by writing programs – and a pair of sleep cycles, both of which are loosely biologically inspired. The first sleep cycle, which we refer to as **consolidation**, grows its library of code by replaying experiences from waking and consolidating them into new code abstractions. This cycle is inspired by the formation of abstractions during sleep memory consolidation [10]. The second sleep cycle, which we refer to as **dreaming**, improves the agents knowledge of how to write code by training a neural network to help search for programs. The neural net is trained on replayed experiences as well as ‘fantasies’, or samples, from the learned library. These two kinds of dreams are inspired by the distinct episodic replay and hallucination components of dream sleep [14].

DREAMCODER builds on several generations of research in AI, program synthesis, and cognitive science, with program induction being one of the oldest theoretical frameworks for concept learning within artificial intelligence [40], and the conceptually allied ‘Language of Thought Hypothesis’ being almost as old [13]. Recent neural program synthesis systems pair a fixed programming language to a learned neural network that guides program search [?, 3, 8], while recent symbolic AI research has developed frameworks for inferring reusable pieces of code shared in a multitask setting [11, 7, 26, 27]. A main goal of this work is to combine and refine these techniques with the intention of building agents that, like humans, develop domain expertise for new classes of problems, motivated by models from cognitive science that explain the flexibility and generality of human thinking in terms of program-like mental representations, referred to as a ‘language of thought’ [32].

Most prior program synthesis work assumes a fixed library of code components [18] or a hand-engineered domain-specific language (a ‘DSL’ [34, 15]). In the vocabulary of program synthesis, our approach is to learn an increasingly finely-tuned DSL (during consolidation) while jointly learning to deploy it efficiently on new problems (during dreaming). Each wake/sleep cycle creates new DSL components that build on components learned in earlier sleep cycles, growing a DSL with nested hierarchies of code. We identify this cumulative nesting of abstractions as a variety of deep representation learning [25]. Figure 2 diagrams a subset of these learned networks (the DSL). For example, the system learns to sort sequences of numbers by invoking a DSL component 4 layers deep, or draws the leftmost pairs of images in Figure 2 using a depth-3 component. For this reason we refer to DREAMCODER as an instance of ‘deep program learning’.

2 Old introduction

An old dream within AI is a machine that learns and reasons by writing its own programs. This vision stretches back to the 1960's [40] and, if fully realized, could bring us much closer to machines that learn and think like humans. Computational models of cognition often explain the flexibility and richness of human thinking in terms of program learning: from everyday thinking and problem solving (motor program induction as an account of recognition and generation of handwriting and speech [22]; functional programs as a model of natural language semantics [?]) to learning problems that unfold over longer developmental time scales: the child's acquisition of intuitive theories (of kinship, taxonomy, etc.) [45] and natural language grammar [37], to name just a few. An outstanding challenge, however, is to engineer program-learners that display the same level of domain-generalty as the humans they are meant to model.

Recent program-learning systems developed within the AI and machine learning community are impressive along many dimensions, authoring programs for problem domains like drawing pictures [?, 12], transforming text [15] and numerical sequences [3], robot motion planning [8], and reasoning over common sense knowledge bases [31]. These systems work in different ways, but typically hinge upon a carefully hand-engineered Domain Specific Language (DSL). The DSL restricts the space of programs to contain the kinds of concepts needed for one specific domain. For example, a picture-drawing DSL could include concepts like circles and spirals, and a DSL for numerical sequences could include sorting and reversing lists of numbers. Modern systems also learn how to efficiently deploy the DSL on new problems [9, 3, 19], but – unlike human learners – do not discover the underlying system of concepts needed to navigate the domain.

We contribute a program-induction system that learns the domain-specific concepts (DSL) while jointly learning how to use those concepts. This joint learning problem models two complementary notions of domain expertise: domain experts have at their disposal a powerful, yet specialized repertoire of concepts and abstractions (analogous to the DSL) while also having accurate intuitions about when and how to use those concepts to solve new problems. Representative domains, along with DSLs we learn for them, are shown in Figure 2.

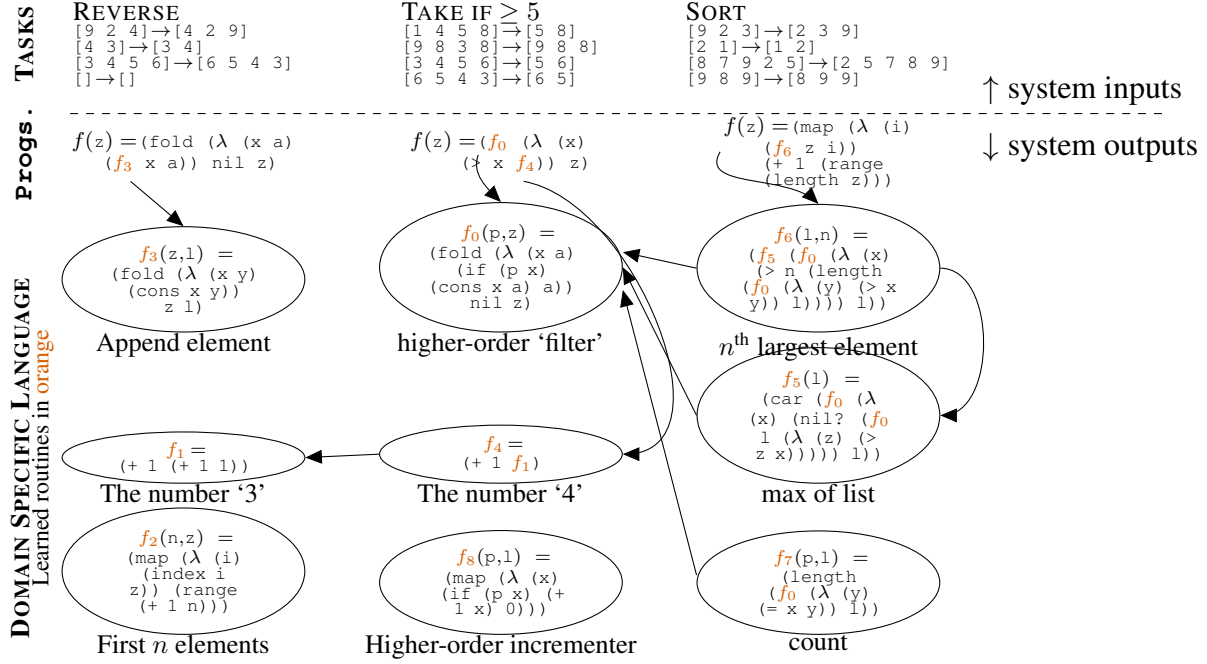
We call our system 'DreamCoder' because it acquires these two kinds of expertise through a novel kind of wake/sleep or 'dream' learning [16], iterating through a wake cycle – where it solves problems by writing programs – and a pair of sleep cycles, both of which are loosely biologically inspired by actual sleep. The first sleep cycle, which we refer to as **consolidation**, grows the DSL by replying experiences from waking and consolidating them into new code abstractions. This cycle is inspired by the formation of abstractions during sleep memory consolidation [10]. The second sleep cycle, which we refer to as **dreaming**, improves the agents knowledge of how to write code by training a neural network to help search for programs. The neural net is trained on replayed experiences as well as 'fantasies', or samples, from the DSL. These two kinds of dreams are inspired by the distinct episodic replay and hallucination components of dream sleep [14].

This dream-learning architecture brings together two lines of prior work, both of which have been separately influential within artificial intelligence. One line of work considers the problem of learning new concepts, abstractions, or 'options' from experience [7, 26, 41, ?, 42], while the other line of work considers the problem of learning how to deploy those concepts efficiently [9, 3, 19]. Our goal with DreamCoder is to show that the combination of these ideas is uniquely powerful, and pushes us toward program-writing systems that, like human learners, autonomously acquire the expertise needed to navigate a new domain of problems.

Each wake/sleep cycle creates new DSL components that build on components learned in earlier sleep cycles, growing a DSL with nested hierarchies of code. We identify this cumulative nesting of abstractions as a variety of deep representation learning [25]. Figure 2 diagrams a subset of these learned layers (the DSL). For example, the system learns to sort sequences of numbers by invoking a DSL component 4 layers deep, or draws the leftmost pairs of images in Figure 2 using a depth-3 component. For this reason we refer to DREAMCODER as an instance of 'deep program learning'.

Our goal with DREAMCODER is to engineer a system that develops domain expertise in humanlike ways. This involves learning both declarative and procedural knowledge, like a DSL and a recognition model, but includes other features of the development of expertise: humans can become experts in many fields, and so we evaluate our algorithm across six different problem domains; a human expert doesn't become an expert overnight, and needs more than a handful of example problems to learn from, but doesn't need millions of examples; similarly our algorithm's learning trajectory unfolds over a series of wake/sleep cycles requiring around a hundred problems per domain.

DOMAIN: LIST PROCESSING



DOMAIN: GRAPHICS PROGRAMMING

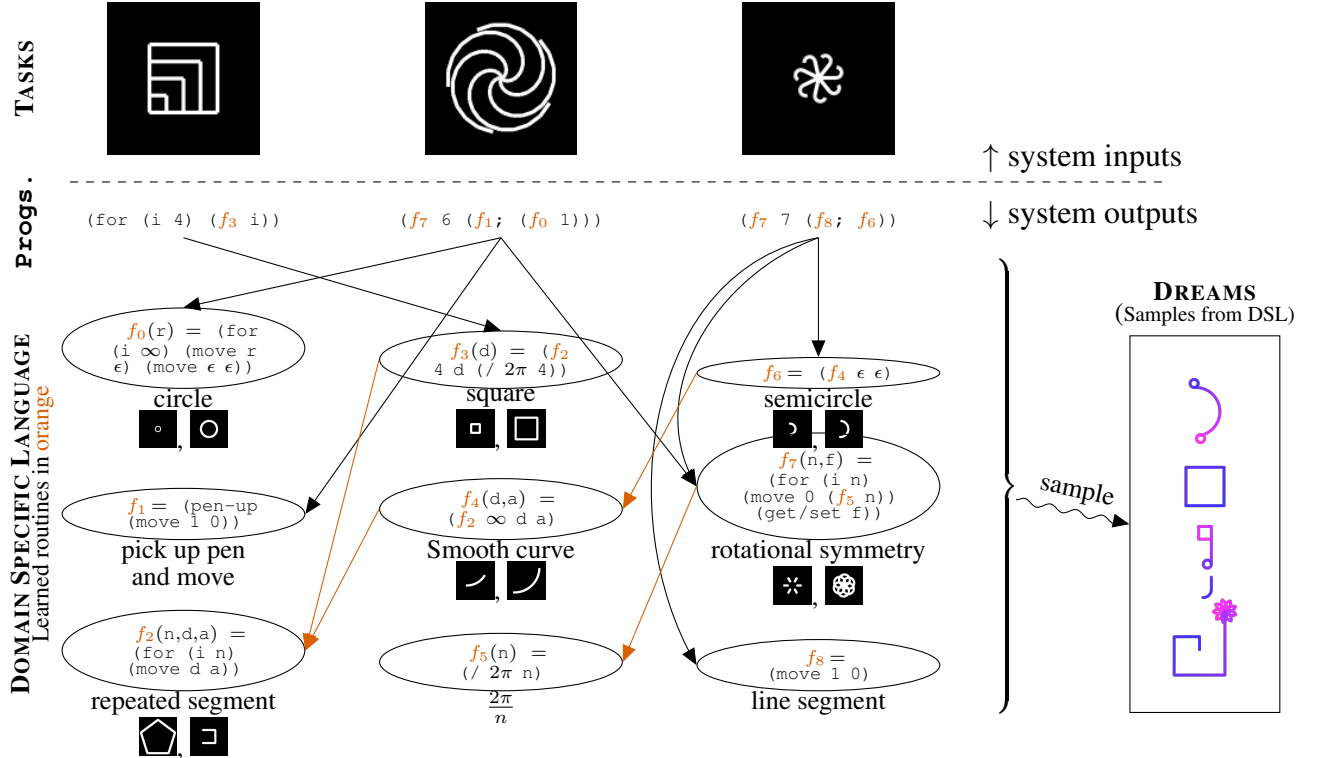


Figure 2: Two of the six domains we apply our system to. Agent observes tasks (top rows) which it solves by writing programs (middle rows) while jointly growing a library (DSL; bottom rows). Learned DSLs rediscover multiple higher-order functions (*filter* for list functions and rotational symmetry for generative graphics). Learned DSL components call each other (arrows).

3 Deep Wake/Sleep Program Induction

DREAMCODER takes as its goal to acquire domain-specific expertise, which it learns by solving a collection of programming **tasks**. It alternately finds programs that solve tasks (Wake – Figure 3 top); improves its DSL by analyzing programs found during waking (Consolidation – Figure 3 left); and trains a neural network that efficiently guides search for programs in the DSL (Dreaming – Figure 3 right). The learned DSL acts as a prior on programs likely to solve tasks in the domain, while the neural net looks at a specific task and produces a “posterior” for programs likely to solve that specific task (Figure 3 middle). The neural network thus functions as a **recognition model** supporting a form of approximate Bayesian program induction, jointly trained with a **generative model** for programs encoded in the DSL, in the spirit of the Helmholtz machine [16]. The recognition model ensures that searching for programs remains tractable even as the DSL (and hence the search space for programs) expands. The generative model, or DSL, distills out common abstractions across programs discovered during waking, growing a network of increasingly deep and specialized domain-specific concepts (Figure 2, bottom rows).

These wake sleep/cycles function as an approximate inference algorithm that observes a collection of tasks, written X , and infers both a program solving each task, as well as a distribution over programs encoded by a DSL, written \mathcal{D} . We equip \mathcal{D} with a learned weight vector θ , and together (\mathcal{D}, θ) define a generative model over programs (Appendix 2). Writing $Q(p|x)$ for the approximate posterior predicted by the recognition model, we iteratively (and approximately) solve for

$$p_x = \arg \max_p \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}^*, \theta^*] \quad \text{Wake}$$

$$\begin{aligned} \mathcal{D}^* &= \arg \max_{\mathcal{D}} \int \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_{p \text{ found during waking}} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] d\theta \\ \theta^* &= \arg \max_{\theta} \mathbb{P}[\mathcal{D}^*, \theta] \prod_{x \in X} \sum_{p \text{ found during waking}} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}^*, \theta] \end{aligned} \quad \text{Consolidation sleep}$$

$$Q(p|x) \approx \mathbb{P}[p|x, \mathcal{D}^*, \theta^*] \quad \text{Dream sleep}$$

which serves to maximize a lower bound on the posterior over (\mathcal{D}, θ) given X (Appendix A.1).

This 3-phase inference procedure works through two distinct kinds of bootstrapping. During each sleep cycle the next DSL bootstraps off the primitives learned during earlier cycles, growing an increasingly deep learned program representation. In tandem the DSL and recognition model bootstrap each other: a more finely tuned DSL yields richer dreams for the recognition model to learn from, while a more accurate recognition model solves more tasks during waking which then feed into the next DSL.

Waking consists of searching for task-specific programs with high posterior probability, or programs which are a priori likely and which solve a task. We find programs solving a task by enumerating programs from the DSL in decreasing order of their probability under the recognition model, and then check if a program p assigns positive probability to a task ($\mathbb{P}[x|p] > 0$). We represent programs as polymorphically typed λ -calculus expressions, an expressive formalism including conditionals, variables, higher-order recursive functions, and the ability to define new functions.

3.1 Wake: Solving tasks

During waking we enumerate programs from the DSL in decreasing order of their probability according to the recognition model, and then check if a program p assigns positive probability to a task ($\mathbb{P}[x|p] > 0$); if so, we incorporate p into the beam \mathcal{B}_x . We represent programs as polymorphically-typed λ -calculus expressions, a representation closely resembling Lisp and functional languages like Haskell and OCaml, including variables, conditionals, higher-order recursive functions, and the ability to create new functions. We ‘batch’ the tasks by randomly shuffling the training set and playing small minibatches of tasks to the agent during each wake cycle.

Why enumerate, when the program synthesis community has invented many sophisticated algorithms that search for programs? [?, ?, ?, ?, 34]. We have two reasons: (1) A key point of our work is that learning the DSL, along with a neural recognition model, can make program induction tractable, even if the search algorithm is very simple. (2) Enumeration is a general approach that can be applied to any program induction problem. Many of these more sophisticated approaches require special conditions on the space of programs.

3.2 Consolidation-Sleep: Growing a Domain Specific Language

The DSL offers a set of abstractions that allow an agent to concisely express solutions to the tasks at hand. We automatically discover these new abstractions by combining two ideas. First, we build on techniques from the

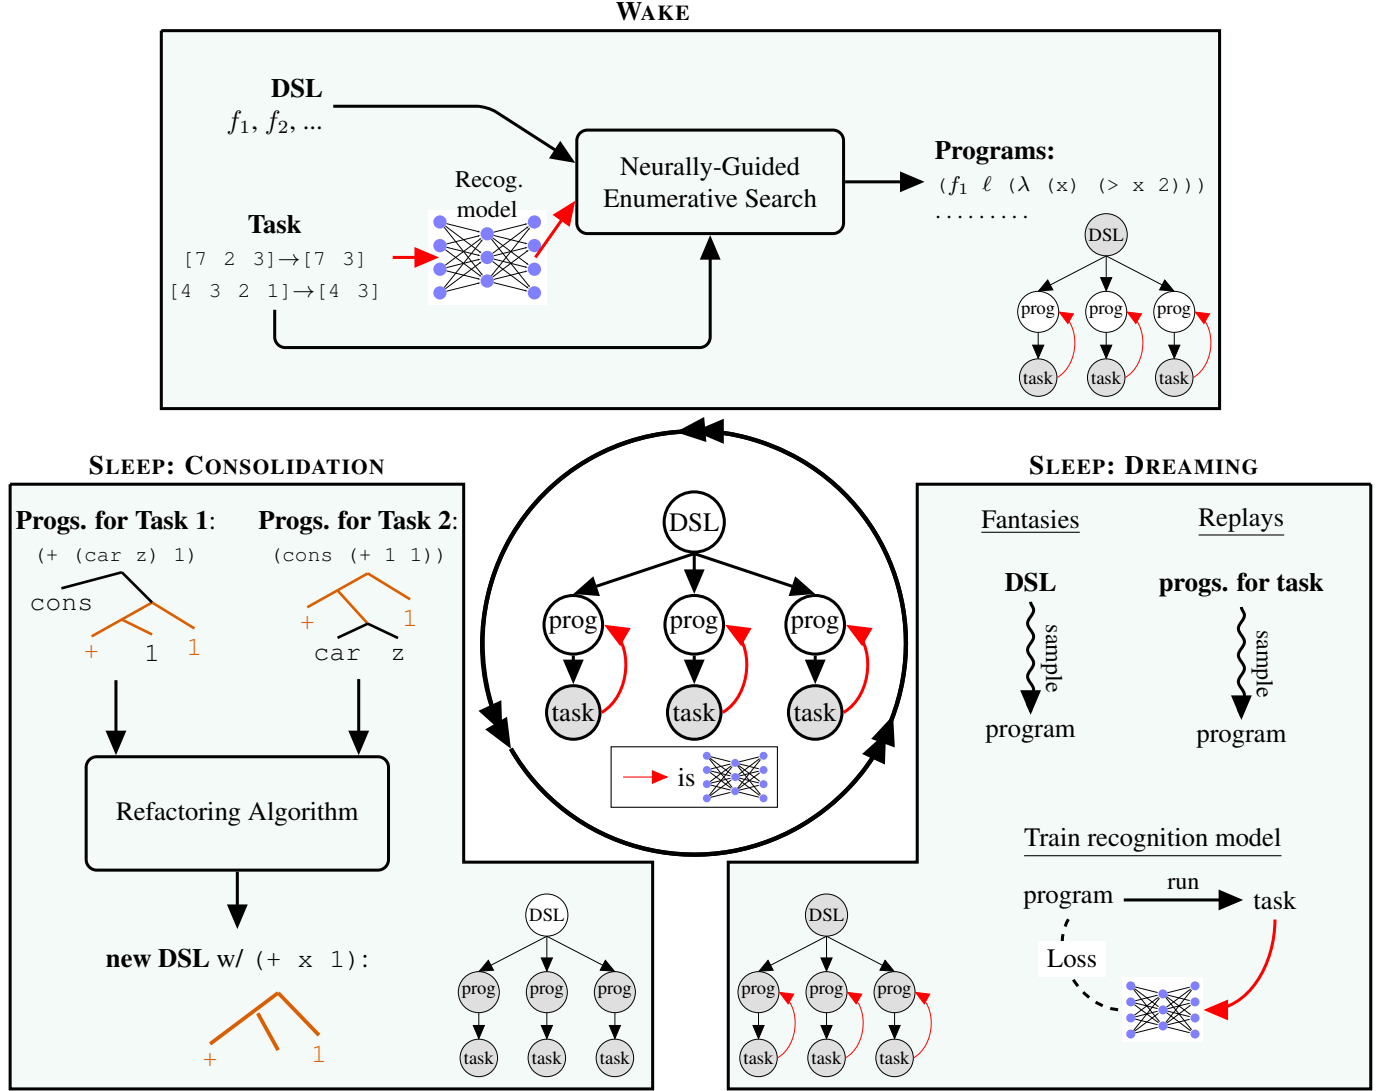


Figure 3: **Middle:** DREAMCODER as a graphical model. Agent observes programming tasks (e.g., input/outputs for list processing or images for graphics programs), which it explains with latent programs, while jointly inferring a latent Domain Specific Language (DSL) capturing cross-program regularities. A neural network, called the *recognition model* (red arrows) is trained to quickly infer programs with high posterior probability. **Top:** Wake phase infers programs while holding the DSL and recognition model fixed. **Left:** Sleep (Consolidation) phase infers DSL while holding the programs fixed by refactoring programs found during waking and extracting common components. **Right:** Sleep (Dreaming) phase trains recognition model to predict approximate posterior over programs conditioned on task. Trained on ‘Fantasies’ (programs sampled from DSL) & ‘Replays’ (programs found during waking).




Figure 4: Neurally-guided program inference pipeline. Recognition model outputs distribution over program $Q(p|x)$. Program output by enumerative search incorporated into beam if likelihood $\mathbb{P}[x|p] > 0$

programming languages community to develop a new algorithm for automatically refactoring programs, where this refactoring exposes common reused subexpressions across the programs found during waking. Second, we use this automatic refactoring process to search for DSLs that maximally compress these programs by incorporating reused subexpressions into the DSL.

Mathematically this compression takes the form of finding the DSL maximizing $\int \mathbb{P}[\mathcal{D}, \theta] \mathbb{P}[X|\mathcal{D}, \theta] d\theta$ (Sec. 3). We replace this marginal with an AIC approximation [1] and marginalize over refactorings of programs found during waking, minimizing the following expression, which can be interpreted as a kind of compression:

$$\underbrace{-\log \mathbb{P}[\mathcal{D}] + \min_{\theta} \left(-\log \mathbb{P}[\theta|\mathcal{D}] + \|\theta\|_0 + \sum_{x \in X} -\log \sum_{\substack{p \text{ a refactoring of } p' \\ p' \text{ found during waking}}} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] \right)}_{\text{Description length of } (\mathcal{D}, \theta)} \quad (1)$$

But a program has infinitely many possible refactorings, rendering Eq. 1 intractable. Rather than consider every refactoring we bound the number of λ -calculus evaluation steps separating a program from its refactoring. Now the number of refactorings is finite but astronomically large: Figure 5A diagrams a problem where the agent rediscovers the higher-order function `map` starting from the basics of Lisp and the Y-combinator, but where there are approximately 10^{14} possible refactorings – a quantity that grows exponentially both as a function of program size and a function of the bound on evaluation steps. How can we tame this combinatorial explosion?

To resolve this exponential growth we introduce a new data structure combining ideas from version space algebras [23, 30, 34] and equivalence graphs [43]. A version space is a tree-shaped data structure that compactly represents a large set of programs and supports efficient set operations like union, intersection, and membership checking, while equivalence graphs are data structures that track semantic equivalences between program subexpressions. In Appendix A.5.1, we give a dynamic program that takes as input a program and then outputs a version space containing its refactorings while tracking semantically equivalent subexpressions. Figure 5B diagrams a subtree of a version space containing refactorings of a small program. Our technique is substantially more efficient than explicitly representing the space of possible refactorings: for the example in Figure 5A, we represent the space of refactorings using a version space with 10^6 nodes, which encodes 10^{14} refactorings. Appendix A.5 specifies how we combine this probabilistic and symbolic machinery to update the DSL. At a high level, our approach is to search locally through the space of DSLs, proposing small changes until Eq. 1 fails to decrease.

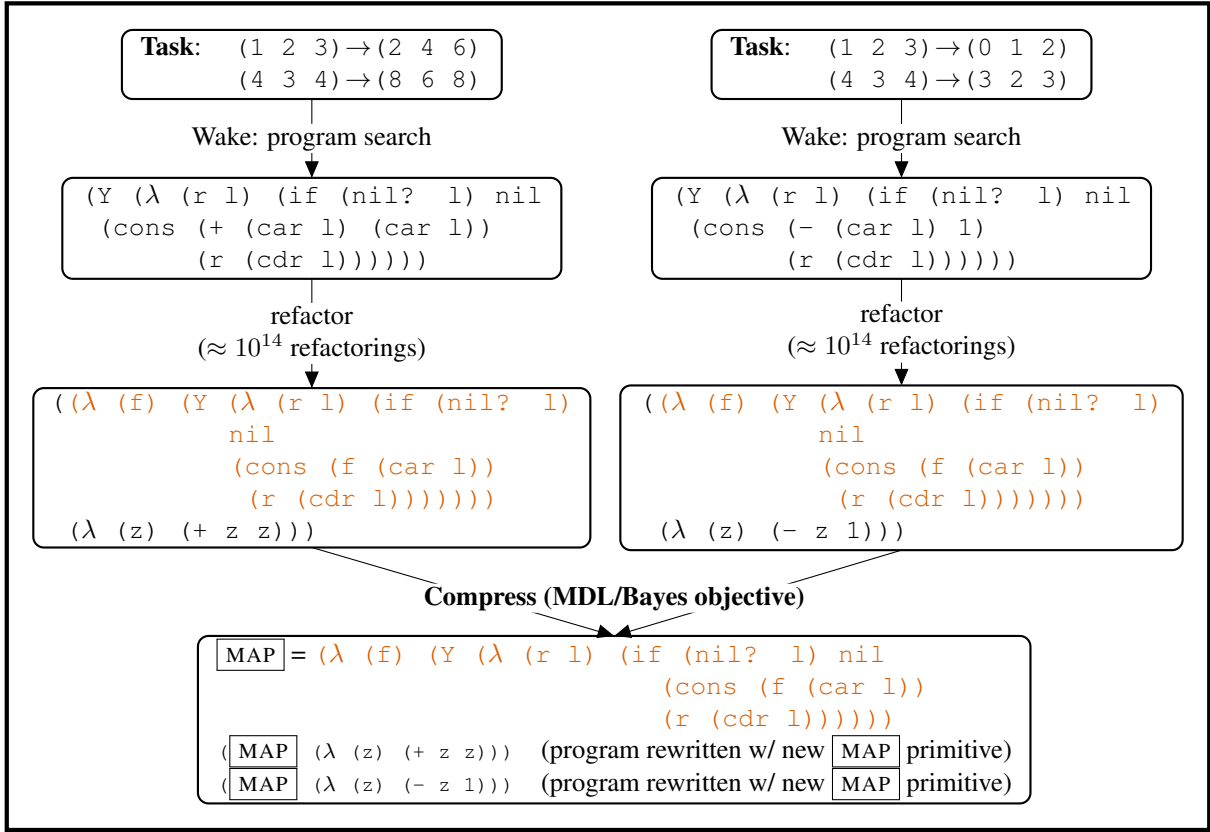
3.3 Dream Sleep: Training a Neural Recognition Model

During “dreaming” the system learns a recognition model that guides program search. It learns from (program, task) pairs drawn from two sources of self-supervised data: *replays* of programs discovered during waking, and *fantasies*, or programs drawn from the DSL. Replays ensure that the recognition model is trained on the actual tasks it needs to solve, and does not forget how to solve them. Fantasies ensure that the recognition model has a large and highly varied corpus of (program, task) pairs to learn from.

Formally, the recognition model $Q(p|x)$ should approximate the posterior $\mathbb{P}[p|\mathcal{D}, \theta, x]$. We can either train Q to perform full posterior inference by minimizing the expected KL-divergence, $\mathbb{E}[\text{KL}(\mathbb{P}[p|x, \mathcal{D}, \theta] || Q(p|x))]$, or we can train Q to perform MAP inference by maximizing $\mathbb{E}[\max_p \mathbb{P}[\cdot|x, \mathcal{D}, \theta] \log Q(p|x)]$, where in both cases the expectation is taken over tasks. Taking this expectation over the empirical distribution of tasks trains Q on replays; taking it over samples from the generative model trains Q on fantasies. We define a pair of alternative objectives for the recognition model, $\mathcal{L}^{\text{posterior}}$ and \mathcal{L}^{MAP} , which either train Q to perform full posterior inference or MAP inference, respectively. These objectives combine replays and fantasies:

$$\begin{aligned} \mathcal{L}^{\text{posterior}} &= \mathcal{L}_{\text{Replay}}^{\text{posterior}} + \mathcal{L}_{\text{Fantasy}}^{\text{posterior}} & \mathcal{L}^{\text{MAP}} &= \mathcal{L}_{\text{Replay}}^{\text{MAP}} + \mathcal{L}_{\text{Fantasy}}^{\text{MAP}} \\ \mathcal{L}_{\text{Replay}}^{\text{posterior}} &= \mathbb{E}_{x \sim X} \left[\sum_{p \in \mathcal{B}_x} \frac{\mathbb{P}[x, p|\mathcal{D}, \theta] \log Q(p|x)}{\sum_{p' \in \mathcal{B}_x} \mathbb{P}[x, p'|\mathcal{D}, \theta]} \right] & \mathcal{L}_{\text{Replay}}^{\text{MAP}} &= \mathbb{E}_{x \sim X} \left[\max_{\substack{p \in \mathcal{B}_x \\ p \text{ maxing } \mathbb{P}[\cdot|x, \mathcal{D}, \theta]}} \log Q(p|x) \right] \\ \mathcal{L}_{\text{Fantasy}}^{\text{posterior}} &= \mathbb{E}_{(p, x) \sim (\mathcal{D}, \theta)} [\log Q(p|x)] & \mathcal{L}_{\text{Fantasy}}^{\text{MAP}} &= \mathbb{E}_{x \sim (\mathcal{D}, \theta)} \left[\max_p \log Q(p|x) \right] \end{aligned}$$

A



B

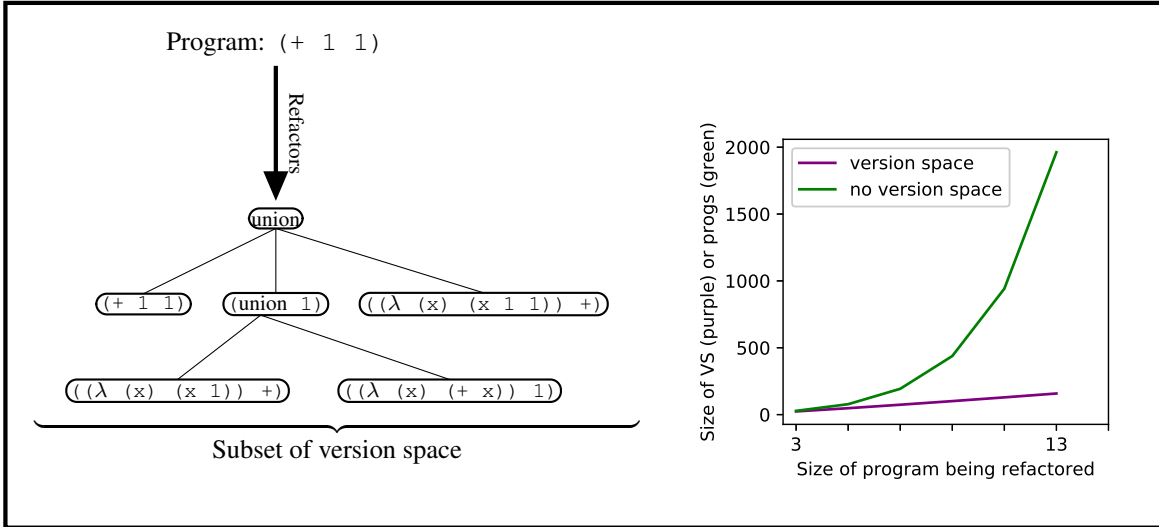


Figure 5: DSL learning as code refactoring. **Panel A:** During waking we discover programs for each task, then refactor the code from those programs to expose common subprograms (highlighted in orange). Common subprograms are incorporated into the DSL when they increase a Bayesian objective. Intuitively, these new DSL components best compress the programs found during waking. **Panel B:** # of possible refactorings grows exponentially with program size, so we represent refactorings using version spaces, which augment syntax trees with a *union* operator whose children are themselves version spaces. Right graph: version spaces are exponentially more efficient than explicitly constructing set of refactorings. In this graph, refactored programs are of the form $1 + 1 + \dots + 1$.

We maximize \mathcal{L}^{MAP} rather than $\mathcal{L}^{\text{posterior}}$ for two reasons: \mathcal{L}^{MAP} prioritizes the shortest program solving a task, thus more strongly accelerating enumerative search during waking; and, combined with our parameterization of Q , described next, we will show that \mathcal{L}^{MAP} forces the recognition model to break symmetries in the space of programs.

Parameterizing Q . The recognition model predicts a fixed-dimensional tensor encoding a distribution over subroutines in the DSL, conditioned on the local context in the syntax tree of the program. This local context consists of the parent node in the syntax tree, as well as which argument is being generated, functioning as a kind of ‘bigram’ model over trees. Figure 6 (left) diagrams this generative process. This parameterization confers three main advantages: (1) it supports fast enumeration and sampling of programs, because the recognition model only runs once per task, like in [3, 11, 29] – thus we can fall back on fast enumeration if the target program is unlike the training programs; (2) the recognition model provides fine-grained information about the structure of the target program, similar to [9, 46]; and (3) in conjunction with \mathcal{L}^{MAP} the recognition model learns to break symmetries in the space of programs.

Symmetry breaking. A good DSL not only exposes high-level building blocks, but also carefully restricts the ways in which those building blocks are allowed to compose. For example, a DSL for arithmetic should disallow adding zero, or force right-associative addition. A bigram parameterization of the recognition model, combined with the \mathcal{L}^{MAP} training objective, interact in a way that breaks symmetries like these, allowing the agent to more efficiently explore the space of programs. This interaction occurs because the bigram parameterization can disallow DSL primitives depending on their local syntactic context, while the \mathcal{L}^{MAP} objective forces all probability mass onto a single member of a set of syntactically distinct but semantically equivalent expressions (Appendix A.6). We experimentally confirm this symmetry-breaking behavior by training recognition models that minimize either $\mathcal{L}^{\text{MAP}}/\mathcal{L}^{\text{posterior}}$ and which use either a bigram parameterization/unigram¹ parameterization. Figure 6 (right) shows the result of training Q in these four regimes and then sampling programs. On this particular run, the combination of bigrams and \mathcal{L}^{MAP} learns to avoid adding zero and associate addition to the right — different random initializations lead to either right or left association.

3.3.1 Parameterizing Q

Broadly the literature contains two different approaches to parameterizing conditional distributions over programs. The first approach [9, 46] is to use a recurrent network to predict the entire program token-by-token, which has the advantage that, if the network is sufficiently powerful, it can completely solve the synthesis problem. The disadvantage is that these models can perform poorly at out-of-sample generalization [], which is critical for our setting, as the agent may need to solve new tasks that are qualitatively different from the tasks it has solved so far.


The second approach is to have Q predict a fixed-dimensional weight vector, which then biases a fast enumerator [3, 11] or sampler [29]. This approach can enjoy strong out-of-sample generalization, because it can fall back on enumeration or sampling when the target program is unlike the training programs. A main drawback is that the neural net is deliberately handicapped, and can only send so much information about the target program.

We adopt a middle ground between these two extremes. Our recognition model predicts a distribution over primitives in the DSL, conditioned on the local context in the syntax tree of the program. When predicting the next node to add to the syntax tree of a program, the recognition model conditions on the parent node, as well as which argument is being generated. This is a kind of ‘bigram’ model over trees, where the bigrams take the form of (parent, child, argument index). Figure 6 (left) diagrams this generative process and Algorithm 5 specifies a sampling procedure for $Q(\cdot|x)$. This parameterization confers three main advantages: (1) it supports fast enumeration and sampling of programs, because the recognition model only needs to run once for each task; (2) it allows the recognition model to provide fine-grained information about the structure of the target program; and (3) training this recognition model causes it to learn to break symmetries in the space of programs, described next.

3.3.2 Learning to break symmetries in program space

A good DSL not only exposes high-level building blocks, but also carefully restricts the ways in which those building blocks are allowed to compose. For example, a DSL for arithmetic should contain both addition and the number zero but disallow adding zero. These restrictions break symmetries in the space of programs. A bigram parameterization of the recognition model, combined with the \mathcal{L}^{MAP} training objective, interact in a way that breaks symmetries in the program space, allowing the agent to more efficiently explore the space of programs. This interaction occurs because the bigram parameterization can disallow DSL primitives depending on their local syntactic context, while the \mathcal{L}^{MAP}

¹In the unigram variant Q predicts a $|\mathcal{D}| + 1$ -dimensional vector: $Q(p|x) = \mathbb{P}[p|\mathcal{D}, \theta_i = Q_i(x)]$, and was used in our prior work [11]



	Unigram	Bigram
$\mathcal{L}^{\text{posterior}}$	<i>Three samples:</i>	<i>Three samples:</i>
	(+ 1 0)	0
	(+ (+ 0 0)	(+ (+ (+ 0 0)
	(+ 1 0))	(+ 0 1)) 1)
	(+ 1 1)	1
\mathcal{L}^{MAP}	63.0% right-associative	55.8% right-associative
	37.4% +0's	31.9% +0's
	<i>Three samples:</i>	<i>Three Samples:</i>
	1	(+ 1 (+ 1 (+ 1
	(+ 1 (+ 1 (+ (+ 1	(+ 1 (+ 1 1))))))
	(+ 1 1)) 1)))	0
	(+ (+ 1 1) 1)	(+ 1 (+ 1 (+ 1 1)))
	48.6% right-associative	97.9% right-associative
	0.5% +0's	2.5% +0's

Figure 6: **Left:** Bigram parameterization of distribution over programs predicted by recognition model. Here the program (syntax tree shown above) is $(\lambda (a) (+ 9 (* a a)))$. Each conditional distribution predicted by the recognition model is written $Q_{\text{parent,child,argument index}}(x)$, where x is a task. **Right:** Agent learns to break symmetries in program space only when using both bigram parameterization and \mathcal{L}^{MAP} objective, associating addition to the right and avoiding adding zero. % right-associative calculated by drawing 500 samples from Q . \mathcal{L}^{MAP} /Unigram agent incorrectly learns to never generate programs with 0's, while \mathcal{L}^{MAP} /Bigram agent correctly learns that 0 should only be disallowed as an argument of addition. Tasked with building programs from +, 1, and 0.

objective forces all probability mass onto a single member of a set of syntactically distinct but semantically equivalent expressions (Appendix A.6).

We experimentally confirm this symmetry-breaking behavior by training recognition models that minimize either $\mathcal{L}^{\text{MAP}}/\mathcal{L}^{\text{posterior}}$ and which use either a bigram parameterization/unigram² parameterization. Figure 6 shows the result of training Q in these four regimes for a DSL containing +, 0, and 1 and then sampling programs. On this particular run, the combination of bigrams and \mathcal{L}^{MAP} learns to avoid adding zero and associate addition to the right — different random initializations lead to either right or left association.

4 Experiments

4.1 Programs that manipulate sequences

We first apply DREAMCODER to two classic benchmark domains: list processing and text editing. In both cases we solve tasks specified by a input/output examples, starting with a generic functional programming basis: `foldr`, `unfold`, `if`, `map`, `length`, `index`, `=`, `+`, `-`, `0`, `1`, `cons`, `car`, `cdr`, `nil`, and `is-nil`.

4.1.1 List Processing

We took 218 list manipulation tasks from our previous work [11], each with 15 input/output examples. In solving these tasks, the system composed 16 new DSL components, and discovered multiple higher-order functions. Each round of memory consolidation built on components discovered in earlier sleep cycles — for example the agent first learns the higher-order function `filter`, uses `filter` to learn to take the maximum element of a list, then uses that routine to learn a new component for extracting the n^{th} largest element of a list, which it finally uses to solve a task involving sorting a list of numbers (Figure ??). This incremental, modular learning of deep hierarchies of DSL components occurs because of the alternation between code writing (during waking) and code refactoring (during the consolidation phase of sleep).

²In the unigram variant Q predicts a $|\mathcal{D}| + 1$ -dimensional vector: $Q(p|x) = \mathbb{P}[p|\mathcal{D}, \theta_i = Q_i(x)]$, and was used in our prior work [11]




Figure 7: Three domains where the agent infers a program from visual input. **Left:** 16 (out of 160) LOGO graphics tasks. Agent writes a program controlling a ‘pen’ that draws the target picture. **Middle:** 9 (out of 112) tower building tasks. Agent writes a program controlling a ‘hand’ that builds the target tower. **Right:** 16 (out of 200) symbolic regression tasks. Agent writes a program containing continuous real numbers that fits the points along the curve.

4.1.2 Text Editing

Synthesizing programs that edit text is a classic problem in the programming languages and AI literatures [23], and algorithms that synthesize text editing programs ship in Microsoft Excel [?]. This prior work uses hand-engineered DSLs and hand-engineered search strategies. Here, we will show that we can jointly learn both these ingredients and surpass the state-of-the-art domain-general program synthesizers on a standard text editing benchmark.

We trained our system on 128 automatically-generated text editing tasks, with 4 input/output examples each. We tested, but did not train, on the 108 text editing problems from the SyGuS [2] program synthesis competition. Before any learning, DREAMCODER solves 3.7% of the problems within 10 minutes with an average search time of 235 seconds. After learning, it solves 79.6%, and does so much faster, solving them in an average of 40 seconds. As of the 2017 SyGuS competition, the best-performing synthesizer (CVC4) solves 82.4% of the problems — but here, the competition conditions are 1 hour & 8 CPUs per problem, and with this more generous compute budget we surpass this previous result and solve 84.3% of the problems. SyGuS additionally comes with a different hand-engineered DSL *for each text editing problem*. Here we learned a single DSL that applied generically to all of the tasks, and perform comparably to the best prior work.

4.2 Programs from visual input

We consider three domains where the agent must infer a program from an image (Figure 7). First we consider programs that make plans and take actions: drawing pictures and building towers out of blocks (Sec. 4.2.1-4.2.2).

4.2.1 Programs that draw pictures

Procedural visual concepts are studied across AI and cognitive science — Bongard problems [5], Raven’s progressive matrices [35], and Lake et al.’s BPL model of omniglot [22] are prominent examples. Here we take inspiration from LOGO Turtle graphics [44], tasking our agent with drawing a modest corpus of images while equipping it with control over a ‘pen’, along with arithmetic operations on angles and distances.

Inside its learned DSL we find interpretable parametric drawing routines corresponding to the families of visual objects in its training data, like polygons, circles, and spirals (Figure 8, left) – without supervision the agent has learned the basic types of objects in its visual world. It additionally learns more abstract visual relationships, like rotational symmetry, which it models by incorporating a higher-order function into its DSL (Figure 8, right).

What does DREAMCODER dream of? Prior to learning samples from the DSL are simple and largely unstructured (Figure 9, left). After training the samples become richly structured (Figure 9, right), compositionally recombining latent building blocks and motifs acquired from the training data. This offers a visual window into how the generative model bootstraps recognition model training: as the DSL grows more finely tuned to the domain, the neural net gets richer and more highly varied training data.


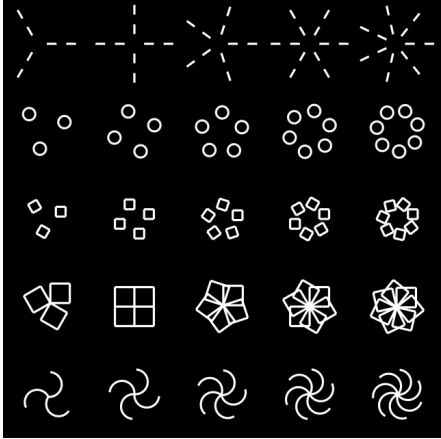
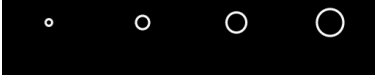




	Parametric drawing routines	Higher-order drawing routine
Semicircle:		<div>Rotational Symmetry:</div> 
Circles:		
Spiral:		
Greek Spiral:		
S-Curves:		
Polygons & Stars:		

Figure 8: Example primitives learned by DREAMCODER when trained on tasks in Figure ?? . Agent learns parametric routines for drawing families of curves (left) as well as subroutines that take entire programs as input (right). Each row of images on the left is the same code executed with different parameters. Each image on the right is the same code executed with different parameters and with a different subprogram provided as input.




Figure 9: Sixteen dreams, or samples, from the DSL before (left) and after (right) training on tasks in Figure ?? . Blue: where the agent started drawing. Pink: where the agent ended drawing.




Figure 10: **Left:** Three (out of 19) learned DSL components for building towers out of Lego-style blocks. These components act like parametric options [42], giving higher-level building blocks that the agent can use to plan. **Right:** 16 random samples, or ‘dreams’, from learned DSL.

4.2.2 Building towers out of ‘Lego’ blocks

Inspired by the classic AI ‘copy task’ — where an agent must look at an image of a tower made of toy blocks and re-create the tower [?] — we give DREAMCODER 112 tower ‘copy tasks’ (Figure ??). Here the agent observes both an image of a tower and the locations of each of its blocks, and must write a program that plans how a simulated hand would build the tower. These towers are built from Lego-style blocks that snap together on a discrete grid. The system starts with the same control flow primitives as with LOGO graphics, and learns parametric ‘options’ for building blocks towers (Figure 10), inferring concepts like arches, staircases, bridges and brick walls.

4.2.3 Symbolic Regression

Here, the agent observes points along the curve of a function, and must write a program that fits those points. We initially equip our learner with addition, multiplication, and division, and task it with solving 200 symbolic regression problems, each either a polynomial or rational function. The recognition model is a convnet that observes an image of the target function’s graph (Fig. 7, rightmost) — visually, different kinds of polynomials and rational functions produce different kinds of graphs, and so the convnet can look at a graph and predict what kind of function best explains it. A key difficulty, however, is that these problems are best solved with programs containing real numbers. Our solution to this difficulty is to enumerate programs with real-valued parameters, and then fit those parameters by automatically differentiating through the programs the system writes and use gradient descent to fit the parameters. We define the likelihood model, $\mathbb{P}[x|p]$, by assuming a Gaussian noise model for the input/output examples, and penalize the use of real-valued parameters using the BIC [4].

We learn a DSL containing 13 new functions, mainly templates for different pieces of polynomials or ratios of polynomials. The model also learns to find programs minimizing the number of continuous parameters — for example, learning to represent linear functions with `(\star real (+ \times real))`. This phenomenon arises from our Bayesian framing: both the generative model’s bias toward shorter programs, and the likelihood model’s BIC penalty.

4.2.4 Generative Modeling of Text

We investigate few-shot learning of generative models by tasking our agent with inferring a probabilistic regular expression from a small number (5) of strings - for example, observing the strings \$1.20 and \$9.42 and predicting the regex \$d

Observed text	Inferred regex	Held out testing text
---------------	----------------	-----------------------




Figure 11: Test set performance across wake/sleep iterations. Each curve is a run with a different random seed. Teal: Full model. Orange: Dreaming only (no DSL learning). Purple: Consolidation only (no recognition model). Search time plots show solid lines (time averaged over all tasks) and dotted lines (time averaged over solved tasks).

4.3 Quantitative Results on Held-Out Tasks

To evaluate the relative importance of DSL learning and recognition model training, we evaluate on held-out testing tasks for each of our domains, measuring both how many tasks are solved and how long it takes to solve them across successive wake/sleep iterations (Fig. 11). We always solve more held-out tasks – and generally solve them in less time – with both components combined. Why? One hypothesis is that some tasks are best solved by DSL learning and others by a neural network, and so including both of our sleep cycles takes the union of these sets of tasks. Another hypothesis is that the DSL and neural network interact synergistically, bootstrapping off each other to solve tasks for which neither alone suffice. We evaluate the relative weight of these interactions by computing the ratio of the tasks solved uniquely by the full model to

5 Discussion

5.1 Learning from Scratch

A long-standing dream within the program induction community is “learning from scratch”: starting with a *minimal* Turing-complete programming language, and then learning to solve a wide swath of induction problems [40, 38, 17, 41]. All existing systems, including ours, fall far short of this dream, and it is unclear (and we believe unlikely) that this dream could ever be fully realized. How far can we push in this direction? “Learning from scratch” is subjective, but a reasonable starting point is the set of primitives provided in 1959 Lisp [28]: these include conditionals, recursion, arithmetic, and the list operators `cons`, `car`, `cdr`, and `nil`. A basic first goal is to start with these primitives, and then recover a DSL that more closely resembles modern functional languages like Haskell and OCaml. Recall (Sec. 4.1) that we initially provided our system with functional programming routines like `map` and `fold`.

We ran the following experiment: DREAMCODER was given a subset of the 1959 Lisp primitives, and tasked with solving 18 programming exercises. A key difference between this setup and our previous experiments is that, for this experiment, the system is given primitive recursion, whereas previously we had sequestered recursion within higher-order functions like `map`, `fold`, and `unfold`.

After running for 93 hours on 48 CPUs, our algorithm solves these 18 exercises, along the way assembling a DSL with a modern repertoire of functional programming idioms and subroutines, including `map`, `fold`, `unfold`, `index`, `length`, and arithmetic operations like building lists of natural numbers between an interval (see Appendix A.9).

We believe that program learners should *not* start from scratch, but instead should start from a rich, domain-agnostic basis like those embodied in the standard libraries of modern languages. What this experiment shows is that DREAMCODER doesn't *need* to start from a rich basis, and can in principle recover many of the amenities of modern programming systems, provided it is given enough computational power and a suitable spectrum of tasks.

5.2 DREAMCODER and the Exploration-Compression family of algorithms

Our work sits within the Exploration-Compression (EC) family of algorithms. EC [7] is a program induction framework where an agent alternates between searching, or ‘exploring’, the space of programs, and then updating its search procedure by compressing programs found during exploration. DREAMCODER grew directly out of research on EC-style systems and both of our sleep phases can be interpreted as a kind of compression: consolidation aims to compactly refactor code, while the recognition model aims to encode a program in as few bits as possible, conditioned on a task. Our previous work, EC² [11], introduced neurally-guided search into the EC framework and served as the starting point for DREAMCODER. Our work here extends EC² by (1) introducing a new refactoring compression algorithm, (2) enriching the neural network with a new bigram-over-trees parameterization and loss function that together learn to break symmetries, accelerating program search during waking, and (3) demonstrating how this family of approaches can be applied to planning and generative modeling problems. Other offshoots of EC include the neurosymbolic framework in [24] and the hierarchical Bayesian program learner in [26]. Closely aligned ideas go even further back [41, 38].

5.3 Acquiring Domain Expertise

One interpretation of our system is as a model of the acquisition of domain expertise. Humans can acquire expertise across many domains – cooking, coding, music, architecture, painting, tennis, or calculus, to name a handful of examples, and every child develops expertise in natural language, intuitive physics, motor control, kinship relationships, and more. Domain experts learn domain-specific abstractions, similar to a DSL: for example, a expert chef knows what combinations of seasonings go together, or an expert mathematician knows a wide set of useful theorems and lemmas. Jointly, experts learn how to recognize when to use these abstractions to solve problems. Becoming an expert involves a learning trajectory that unfolds over relatively long time scales, but has modest data requirements relative to the dominant machine learning paradigms. For example, basic competency in cooking or coding might require on the order of learning a hundred recipes or solving a hundred programming exercises. It is this learning regime that we have targeted with DREAMCODER: Learning about a domain from at most several hundred tasks, but where the learning unfolds over many wake/sleep cycles.

5.4 Prospects for program Induction as part of the generic AI toolkit

Our aim with DREAMCODER is to chart a path by which program induction can become more broadly useful for AI. This means viewing the AI landscape through the lens of program learning, including the terrain considered here — simple kinds of generative modeling, inverse graphics, planning, and programming by example — but also many others like reinforcement learning, commonsense reasoning, natural language understanding, and causal inference. Can program induction rise to the challenge? We believe it can, provided we push jointly along many different axes of AI research; and provided we continue to integrate learning algorithms – both symbolic and neural, both top-down and bottom-up – into our artificial agents.

Our system, with its learned DSL and neural recognition model, is one embodiment of this hybrid symbolic/neural approach, and enjoys some success across small-scale problems in the different domains considered here. Scaling to larger problems, such as inferring 3-D object models (vs LOGO/Turtle), learning natural image grammars (vs),

References

- [1] Hirotugu Akaike. Information theory and an extension of the maximum likelihood principle. In *Selected papers of hirotugu akaike*, pages 199–213. Springer, 1998.
- [2] Rajeev Alur, Dana Fisman, Rishabh Singh, and Armando Solar-Lezama. Sygus-comp 2016: results and analysis. *arXiv preprint arXiv:1611.07627*, 2016.
- [3] Matej Balog, Alexander L Gaunt, Marc Brockschmidt, Sebastian Nowozin, and Daniel Tarlow. Deepcoder: Learning to write programs. *ICLR*, 2016.
- [4] Christopher M. Bishop. *Pattern Recognition and Machine Learning*. 2006.
- [5] M. M. Bongard. *Pattern Recognition*. Spartan Books, 1970.
- [6] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [7] Eyal Dechter, Jon Malmaud, Ryan P. Adams, and Joshua B. Tenenbaum. Bootstrap learning via modular concept discovery. In *IJCAI*, 2013.
- [8] Jacob Devlin, Rudy R Bunel, Rishabh Singh, Matthew Hausknecht, and Pushmeet Kohli. Neural program meta-induction. In *NIPS*, 2017.
- [9] Jacob Devlin, Jonathan Uesato, Surya Bhupatiraju, Rishabh Singh, Abdel-rahman Mohamed, and Pushmeet Kohli. Robustfill: Neural program learning under noisy i/o. *ICML*, 2017.
- [10] Yadin Dudai, Avi Karni, and Jan Born. The consolidation and transformation of memory. *Neuron*, 88(1):20 – 32, 2015.
- [11] Kevin Ellis, Lucas Morales, Mathias Sablé-Meyer, Armando Solar-Lezama, and Josh Tenenbaum. Library learning for neurally-guided bayesian program induction. In *NeurIPS*, 2018.
- [12] Kevin Ellis, Daniel Ritchie, Armando Solar-Lezama, and Joshua B Tenenbaum. Learning to infer graphics programs from hand-drawn images. *NIPS*, 2018.
- [13] Jerry A Fodor. *The language of thought*, volume 5. Harvard University Press, 1975.
- [14] Magdalena J Fosse, Roar Fosse, J Allan Hobson, and Robert J Stickgold. Dreaming and episodic memory: a functional dissociation? *Journal of cognitive neuroscience*, 15(1):1–9, 2003.
- [15] Sumit Gulwani. Automating string processing in spreadsheets using input-output examples. In *ACM SIGPLAN Notices*, volume 46, pages 317–330. ACM, 2011.
- [16] Geoffrey E Hinton, Peter Dayan, Brendan J Frey, and Radford M Neal. The ”wake-sleep” algorithm for unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.
- [17] Marcus Hutter. *Universal artificial intelligence: Sequential decisions based on algorithmic probability*. Springer Science & Business Media, 2004.
- [18] Susmit Jha, Sumit Gulwani, Sanjit A Seshia, and Ashish Tiwari. Oracle-guided component-based program synthesis. In *ICSE*, volume 1, pages 215–224. IEEE, 2010.
- [19] Ashwin Kalyan, Abhishek Mohta, Oleksandr Polozov, Dhruv Batra, Prateek Jain, and Sumit Gulwani. Neural-guided deductive search for real-time program synthesis from examples. *ICLR*, 2018.
- [20] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [21] J.D. Lafferty. *A Derivation of the Inside-outside Algorithm from the EM Algorithm*. Research report.

- [22] Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. Human-level concept learning through probabilistic program induction. *Science*, 350(6266):1332–1338, 2015.
- [23] Tessa Lau. *Programming by demonstration: a machine learning approach*. PhD thesis, 2001.
- [24] Miguel Lázaro-Gredilla, Dianhuan Lin, J Swaroop Guntupalli, and Dileep George. Beyond imitation: Zero-shot task transfer on robots by learning concepts as cognitive programs. *Science Robotics*, 4(26):eaav3150, 2019.
- [25] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015.
- [26] Percy Liang, Michael I. Jordan, and Dan Klein. Learning programs: A hierarchical bayesian approach. In *ICML*, 2010.
- [27] Dianhuan Lin, Eyal Dechter, Kevin Ellis, Joshua B. Tenenbaum, and Stephen Muggleton. Bias reformulation for one-shot function induction. In *ECAI 2014*, 2014.
- [28] John McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, 1960.
- [29] Aditya Menon, Omer Tamuz, Sumit Gulwani, Butler Lampson, and Adam Kalai. A machine learning framework for programming by example. In *ICML*, pages 187–195, 2013.
- [30] Tom M Mitchell. Version spaces: A candidate elimination approach to rule learning. In *Proceedings of the 5th international joint conference on Artificial intelligence-Volume 1*, pages 305–310. Morgan Kaufmann Publishers Inc., 1977.
- [31] Stephen H Muggleton, Dianhuan Lin, and Alireza Tamaddoni-Nezhad. Meta-interpretive learning of higher-order dyadic datalog: Predicate invention revisited. *Machine Learning*, 100(1):49–73, 2015.
- [32] Steven Thomas Piantadosi. *Learning and the language of thought*. PhD thesis, Massachusetts Institute of Technology, 2011.
- [33] Benjamin C. Pierce. *Types and programming languages*. MIT Press, 2002.
- [34] Oleksandr Polozov and Sumit Gulwani. Flashmeta: A framework for inductive program synthesis. *ACM SIGPLAN Notices*, 50(10):107–126, 2015.
- [35] Jean Raven et al. Raven progressive matrices. In *Handbook of nonverbal assessment*, pages 223–237. Springer, 2003.
- [36] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Pearson Education, 2 edition, 2003.
- [37] Ute Schmid and Emanuel Kitzelmann. Inductive rule learning on the knowledge level. *Cognitive Systems Research*, 12(3-4):237–248, 2011.
- [38] Jürgen Schmidhuber. Optimal ordered problem solver. *Machine Learning*, 54(3):211–254, 2004.
- [39] Jake Snell, Kevin Swersky, and Richard Zemel. Prototypical networks for few-shot learning. In *Advances in Neural Information Processing Systems*, 2017.
- [40] Ray J Solomonoff. A formal theory of inductive inference. *Information and control*, 7(1):1–22, 1964.
- [41] Ray J Solomonoff. A system for incremental learning based on algorithmic probability. Sixth Israeli Conference on Artificial Intelligence, Computer Vision and Pattern Recognition, 1989.
- [42] Martin Stolle and Doina Precup. Learning options in reinforcement learning. In *International Symposium on abstraction, reformulation, and approximation*, pages 212–223. Springer, 2002.
- [43] Ross Tate, Michael Stepp, Zachary Tatlock, and Sorin Lerner. Equality saturation: a new approach to optimization. In *ACM SIGPLAN Notices*, volume 44, pages 264–276. ACM, 2009.

- [44] David D. Thornburg. Friends of the turtle. *Compute!*, March 1983.
- [45] T. Ullman, N. D. Goodman, and J. B. Tenenbaum. Theory learning as stochastic search in the language of thought. *Cognitive Development*, 27(4):455–480, 2012.
- [46] Maksym Zavershynskiy, Alex Skidanov, and Illia Polosukhin. Naps: Natural program synthesis dataset. In *ICML*, 2018.

A Appendix

A.1 Probabilistic Formulation of DREAMCODER

Our objective is to infer the maximum a posteriori DSL \mathcal{D} and parameters θ . Writing J for the joint probability of (\mathcal{D}, θ) , this corresponds to solving

$$J(\mathcal{D}, \theta) \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_p \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \int J(\mathcal{D}, \theta) d\theta \quad \theta^* = \arg \max_{\theta} J(\mathcal{D}^*, \theta) \quad (2)$$

where $\mathbb{P}[x|p]$ scores the likelihood of a task $x \in X$ given a program p .³

Evaluating Eq. 2 entails marginalizing over the infinite set of all programs – which is impossible. We make a particle-based approximation to Eq. 2 and instead marginalize over a finite **beam** of programs, with one beam per task, collectively written $\{\mathcal{B}_x\}_{x \in X}$. This particle-based approximation is written $\mathcal{L}(\mathcal{D}, \theta, \{\mathcal{B}_x\})$ and acts as a lower bound on the joint density:

$$J(\mathcal{D}, \theta) \geq \mathcal{L}(\mathcal{D}, \theta, \{\mathcal{B}_x\}) \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_{p \in \mathcal{B}_x} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta], \text{ where } |\mathcal{B}_x| \text{ is small} \quad (3)$$

In all of our experiments we set the maximum beam size $|\mathcal{B}_x|$ to 5.

Wake and sleep cycles correspond to alternate maximization of \mathcal{L} w.r.t. $\{\mathcal{B}_x\}_{x \in X}$ (**Wake**) and (\mathcal{D}, θ) (**Consolidation**):

Wake: Maxing \mathcal{L} w.r.t. the beams. Here (\mathcal{D}, θ) is fixed and we want to find new programs to add to the beams so that \mathcal{L} increases the most. \mathcal{L} most increases by finding programs where $\mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] \propto \mathbb{P}[p|x, \mathcal{D}, \theta]$ is large, i.e., programs with high posterior probability, which is the search objective during waking.

Sleep (Consolidation): Maxing \mathcal{L} w.r.t. the DSL. Here $\{\mathcal{B}_x\}_{x \in X}$ is held fixed and the problem is to search the discrete space of DSLs and find one maximizing $\int \mathcal{L} d\theta$, and then update θ to $\arg \max_{\theta} \mathcal{L}(\mathcal{D}, \theta, \{\mathcal{B}_x\})$.

Finding programs solving tasks is difficult because of the infinitely large, combinatorial search landscape. We ease this difficulty by training a neural recognition model, $Q(p|x)$, during the **Dreaming** phase: Q is trained to approximate the posterior over programs, $Q(p|x) \approx \mathbb{P}[p|x, \mathcal{D}] \propto \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}]$. Thus training the neural network amortizes the cost of finding programs with high posterior probability.

Sleep (Dreaming): tractably maxing \mathcal{L} w.r.t. the beams. Here we train $Q(p|x)$ to assign high probability to programs p where $\mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$ is large, because incorporating those programs into the beams will most increase \mathcal{L} .

A.2 DREAMCODER pseudocode

Algorithm 1 specifies how we integrate wake and sleep cycles.

A.3 Generative model over programs

Algorithm 2 gives a stochastic procedure for drawing samples from $\mathbb{P}[\cdot|\mathcal{D}, \theta]$. It takes as input the desired type of the unknown program, and performs type inference during sampling to ensure that the program has the desired type. It also maintains a *environment* mapping variables to types, which ensures that lexical scoping rules are obeyed.

³For example, for list processing, the likelihood is 1 if the program predicts the observed outputs on the observed inputs, and 0 otherwise; when learning a generative model or probabilistic program, the likelihood is the probability of the program sampling the observation.

Algorithm 1 Full DREAMCODER algorithm

```
1: function DREAMCODER( $\mathcal{D}, X$ ):
2: Input: Initial DSL  $\mathcal{D}$ , tasks  $X$ 
3: Output: Infinite stream of DSLs, recognition models, and beams
4: Hyperparameters: Batch size  $B$ , enumeration timeout  $T$ , maximum beam size  $F$ 
5:  $\theta \leftarrow$  uniform distribution
6:  $\mathcal{F}_x \leftarrow \emptyset, \forall x \in X$  ▷ Initialize beams to be empty
7: while true do ▷ Loop over epochs
8:   shuffle  $\leftarrow$  random permutation of  $X$  ▷ Randomize minibatches
9:   while shuffle is not empty do ▷ Loop over minibatches
10:    batch  $\leftarrow$  first  $B$  elements of shuffle ▷ Next minibatch of tasks
11:    shuffle  $\leftarrow$  shuffle with first  $B$  elements removed
12:     $\forall x \in$  batch:  $\mathcal{F}_x \leftarrow \mathcal{F}_x \cup \{p \mid p \in \text{enumerate}(\mathbb{P}[\cdot|\mathcal{D}, \theta], T) \text{ if } \mathbb{P}[x|p] > 0\}$  ▷ Wake
13:    Train  $Q(\cdot|\cdot)$  to minimize  $\mathcal{L}^{\text{MAP}}$  across all  $\{\mathcal{F}_x\}_{x \in X}$  ▷ Dream Sleep
14:     $\forall x \in$  batch:  $\mathcal{F}_x \leftarrow \mathcal{F}_x \cup \{p \mid p \in \text{enumerate}(Q(\cdot|x), T) \text{ if } \mathbb{P}[x|p] > 0\}$  ▷ Wake
15:     $\forall x \in$  batch:  $\mathcal{F}_x \leftarrow$  top  $F$  elements of  $\mathcal{F}_x$  as measured by  $\mathbb{P}[\cdot|x, \mathcal{D}, \theta]$  ▷ Keep top  $F$  programs
16:     $\mathcal{D}, \theta, \{\mathcal{F}_x\}_{x \in X} \leftarrow \text{CONSOLIDATE}(\mathcal{D}, \theta, \{\mathcal{F}_x\}_{x \in X})$  ▷ Consolidation Sleep
17:    yield  $(\mathcal{D}, \theta), Q, \{\mathcal{F}_x\}_{x \in X}$  ▷ Yield the updated DSL, recognition model, and solutions found to tasks
18:  end while
19: end while
```

A.4 Enumerative program search

Our current implementation of DREAMCODER takes the simple and generic strategy of enumerating programs in descending order of their probability under either (\mathcal{D}, θ) or $Q(p|x)$. Algorithm 2 and 5 specify procedures for sampling from these distributions, but not for enumerating from them. We combine two different enumeration strategies, which allowed us to build a massively parallel program enumerator:

- **Best-first search:** Best-first search maintains a heap of partial programs ordered by their probability — here a partial program means a program whose syntax tree may contain unspecified ‘holes’. Best-first search is guaranteed to enumerate programs in decreasing order of their probability, and has memory requirements that in general grow exponentially as a function of the description length of programs in the heap (thus linearly as a function of run time).
- **Depth-first search:** Depth first search recursively explores the space of execution traces through Algorithm 2 and 5, equivalently maintaining a stack of partial programs. In general it does not enumerate programs in decreasing order of probability, but has memory requirements that grow linearly as a function of the description length of the programs in the stack (thus logarithmically as a function of run time).

Our parallel enumeration algorithm (Algorithm 3) first performs a best-first search until the best-first heap is much larger than the number of CPUs. At this point, it switches to performing many depth-first searches in parallel, initializing a depth first search with one of the entries in the best-first heap. Because depth-first search does not produce programs in decreasing order of their probability, we wrap this entire procedure up into an outer loop that first enumerates programs whose description length is between 0 to Δ , then programs with description length between Δ and 2Δ , then 2Δ to 3Δ , etc., until a timeout is reached. This is similar in spirit to iterative deepening depth first search [36].

A.5 Details of DSL learning

Algorithm 4 specifies our DSL learning procedure. This integrates two toolkits: the machinery of version spaces and equivalence graphs (Appendix A.5.1) along with a probabilistic objective favoring compressive DSLs. The functions $I\beta(\cdot)$ and REFACTOR construct a version space from a program and extract the shortest program from a version space, respectively (Algorithm 4, lines 5-6, 14; Appendix A.5.1). To define the prior distribution over (\mathcal{D}, θ) (Algorithm 4, lines 7-8), we penalize the syntactic complexity of the λ -calculus expressions in the DSL, defining $\mathbb{P}[\mathcal{D}] \propto \exp(-\lambda \sum_{p \in \mathcal{D}} \text{size}(p))$ where $\text{size}(p)$ measures the size of the syntax tree of program p , and λ controls how strongly we regularize the size of the DSL. We place a symmetric Dirichlet prior over the weight vector θ .

Algorithm 2 Generative model over programs

```

1: function sample( $\mathcal{D}, \theta, \tau$ ):
2: Input: DSL ( $\mathcal{D}, \theta$ ), type  $\tau$ 
3: Output: a program whose type unifies with  $\tau$ 
4: return sample'( $\mathcal{D}, \theta, \emptyset, \tau$ )

5: function sample'( $\mathcal{D}, \theta, \mathcal{E}, \tau$ ):
6: Input: DSL ( $\mathcal{D}, \theta$ ), environment  $\mathcal{E}$ , type  $\tau$  ▷ Environment  $\mathcal{E}$  starts out as  $\emptyset$ 
7: Output: a program whose type unifies with  $\tau$ 
8: if  $\tau = \alpha \rightarrow \beta$  then ▷ Function type — start with a lambda
9:   var  $\leftarrow$  an unused variable name
10:  body  $\sim$  sample'( $\mathcal{D}, \theta, \{\text{var} : \alpha\} \cup \mathcal{E}, \beta$ ) ▷ Recursively sample function body
11:  return (lambda (var) body)
12: else ▷ Build an application to give something w/ type  $\tau$ 
13:  primitives  $\leftarrow \{p \mid p : \tau' \in \mathcal{D} \cup \mathcal{E} \text{ if } \tau \text{ can unify with } \text{yield}(\tau')\}$  ▷ Everything in scope w/ type  $\tau$ 
14:  variables  $\leftarrow \{p \mid p \in \text{primitives and } p \text{ a variable}\}$ 
15:  Draw  $e \sim \text{primitives}$ , w.p.  $\propto \begin{cases} \theta_e & \text{if } e \in \mathcal{D} \\ \theta_{\text{var}}/|\text{variables}| & \text{if } e \in \mathcal{E} \end{cases}$ 
16:  Unify  $\tau$  with  $\text{yield}(\tau')$ . ▷ Ensure well-typed program
17:   $\{\alpha_k\}_{k=1}^K \leftarrow \text{args}(\tau')$ 
18:  for  $k = 1$  to  $K$  do ▷ Recursively sample arguments
19:     $a_k \sim \text{sample}'(\mathcal{D}, \theta, \mathcal{E}, \alpha_k)$ 
20:  end for
21:  return ( $e \ a_1 \ a_2 \ \dots \ a_K$ )
22: end if

where:
23:  $\text{yield}(\tau) = \begin{cases} \text{yield}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ \tau & \text{otherwise.} \end{cases}$  ▷ Final return type of  $\tau$ 
24:  $\text{args}(\tau) = \begin{cases} [\alpha] + \text{args}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ [] & \text{otherwise.} \end{cases}$  ▷ Types of arguments needed to get something w/ type  $\tau$ 

```

Algorithm 3 Parallel enumerative program search algorithm

```
1: function enumerate( $\mu, T, \text{CPUs}$ ):
2:   Input: Distribution over programs  $\mu$ , timeout  $T$ , CPU count
3:   Output: stream of programs in approximately descending order of probability under  $\mu$ 
4:   Hyperparameter: nat increase rate  $\Delta$  ▷ We set  $\Delta = 1.5$ 
5:   lowerBound  $\leftarrow$  0
6:   while total elapsed time  $< T$  do
7:     heap  $\leftarrow$  newMaxHeap() ▷ Heap for best-first search
8:     heap.insert(priority = 0, value = empty syntax tree) ▷ Initialize heap with start state of search space
9:     while  $0 < |\text{heap}| \leq 10 \times \text{CPUs}$  do ▷ Each CPU will get approximately 10 jobs (a partial program)
10:      priority, partialProgram  $\leftarrow$  heap.popMaximum()
11:      if partialProgram is finished then ▷ Nothing more to fill in in the syntax tree
12:        if lowerBound  $\leq -\text{priority} < \text{lowerBound} + \Delta$  then
13:          yield partialProgram
14:        end if
15:      else
16:        for child  $\in$  children(partialProgram) do ▷ children( $\cdot$ ) fills in next random choice in syntax tree.
17:          if  $-\log \mu(\text{child}) < \text{lowerBound} + \Delta$  then ▷ Child's description length small enough
18:            heap.insert(priority =  $\log \mu(\text{child})$ , value = child)
19:          end if
20:        end for
21:      end if
22:    end while
23:    yield from ParallelMapCPUs(depthFirst( $\mu, T - \text{elapsed time}, \text{lowerBound}, \cdot$ ), heap.values())
24:    lowerBound  $\leftarrow$  lowerBound +  $\Delta$  ▷ Push up lower bound on MDL by  $\Delta$ 
25:  end while

26: function depthFirst( $\mu, T, \text{lowerBound}, \text{partialProgram}$ ): ▷ Each worker does a depth first search. Enumerates
    completions of partialProgram whose MDL is between lowerBound and lowerBound +  $\Delta$ 
27:  stack  $\leftarrow$  [partialProgram]
28:  while total elapsed time  $< T$  and stack is not empty do
29:    partialProgram  $\leftarrow$  stack.pop()
30:    if partialProgram is finished then
31:      if lowerBound  $\leq -\log \mu(\text{partialProgram}) < \text{lowerBound} + \Delta$  then
32:        yield partialProgram
33:      end if
34:    else
35:      for child  $\in$  children(partialProgram) do
36:        if  $-\log \mu(\text{child}) < \text{lowerBound} + \Delta$  then ▷ Child's description length small enough
37:          stack.push(child)
38:        end if
39:      end for
40:    end if
41:  end while
```

To appropriately score each proposed \mathcal{D} we must reestimate the weight vector θ (Algorithm 4, line 7). Although this may seem very similar to estimating the parameters of a probabilistic context free grammar, for which we have effective approaches like the Inside/Outside algorithm [21], our DSLs are context-sensitive due to the presence of variables in the programs and also due to the polymorphic typing system. Appendix A.5.4 derives a tractable MAP estimator for θ .

Algorithm 4 DSL Induction Algorithm

```

1: Input: Set of beams  $\{\mathcal{B}_x\}$ 
2: Output: DSL  $\mathcal{D}$ , weight vector  $\theta$ 
3:  $\mathcal{D} \leftarrow$  every primitive in  $\{\mathcal{B}_x\}$ 
4: while true do
5:    $\forall p \in \bigcup_x \mathcal{B}_x : v_p \leftarrow I\beta(p)$  ▷ Construct a version space for each program
6:   Define  $L(\mathcal{D}', \theta) = \prod_x \sum_{p \in \mathcal{B}_x} \mathbb{P}[x|p] \mathbb{P}[\text{REFACTOR}(v_p|\mathcal{D}')|\mathcal{D}', \theta]$  ▷ Likelihood if  $(\mathcal{D}', \theta)$  were the DSL
7:   Define  $\theta^*(\mathcal{D}') = \arg \max_{\theta} \mathbb{P}[\theta|\mathcal{D}'] L(\mathcal{D}', \theta)$  ▷ MAP estimate of  $\theta$ 
8:   Define  $\text{score}(\mathcal{D}') = \log \mathbb{P}[\mathcal{D}'] + L(\mathcal{D}', \theta^*(\mathcal{D}')) - \|\theta^*(\mathcal{D}')\|_0$  ▷ objective function
9:    $\text{components} \leftarrow \{\text{REFACTOR}(v|\mathcal{D}) : \forall x, \forall p \in \mathcal{B}_x, \forall v \in \text{children}(v_p)\}$  ▷ Propose many new DSL components
10:   $\text{proposals} \leftarrow \{\mathcal{D} \cup \{c\} : \forall c \in \text{components}\}$  ▷ Propose many new DSLs
11:   $\mathcal{D}' \leftarrow \arg \max_{\mathcal{D}' \in \text{proposals}} \text{score}(\mathcal{D}')$  ▷ Get highest scoring new DSL
12:  if  $\text{score}(\mathcal{D}') < \text{score}(\mathcal{D})$  return  $\mathcal{D}, \theta^*(\mathcal{D})$  ▷ No changes to DSL led to a better score
13:   $\mathcal{D} \leftarrow \mathcal{D}'$  ▷ Found better DSL. Update DSL.
14:   $\forall x : \mathcal{B}_x \leftarrow \{\text{REFACTOR}(v_p|\mathcal{D}) : p \in \mathcal{B}_x\}$  ▷ Refactor beams in terms of new DSL
15: end while

```

A.5.1 Refactoring code with version spaces

Formally, a version space is either:

- A deBuijn⁴ index: written $\$i$, where i is a natural number
- An abstraction: written λv , where v is a version space
- An application: written $(f x)$, where both f and x are version spaces
- A union: $\uplus V$, where V is a set of version spaces
- The empty set, \emptyset
- The set of all λ -calculus expressions, Λ

The purpose of a version space is to compactly represent a set of programs. We refer to this set as the **extension** of the version space:

Definition 1. The **extension** of a version space v is written $\llbracket v \rrbracket$ and is defined recursively as:

$$\begin{aligned}
\llbracket \$i \rrbracket &= \{\$i\} & \llbracket \lambda v \rrbracket &= \{\lambda e : e \in \llbracket v \rrbracket\} & \llbracket (v_1 v_2) \rrbracket &= \{(e_1 e_2) : e_1 \in \llbracket v_1 \rrbracket, e_2 \in \llbracket v_2 \rrbracket\} \\
\llbracket \uplus V \rrbracket &= \{e : v \in V, e \in \llbracket v \rrbracket\} & \llbracket \emptyset \rrbracket &= \emptyset & \llbracket \Lambda \rrbracket &= \Lambda
\end{aligned}$$

Version spaces also support efficient membership checking, which we write as $e \in \llbracket v \rrbracket$. Important for our purposes, it is also efficient to refactor the members of a version space's extension in terms of a new DSL. We define $\text{REFACTOR}(v|\mathcal{D})$ inductively as:

$$\text{REFACTOR}(v|\mathcal{D}) = \begin{cases} e, & \text{if } e \in \mathcal{D} \text{ and } e \in \llbracket v \rrbracket. \text{ Exploits the fact that } \llbracket e \rrbracket \in v \text{ can be efficiently computed.} \\ \text{REFACTOR}'(v|\mathcal{D}), & \text{otherwise.} \end{cases}$$

⁴deBuijn indices are an alternative way of naming variables in λ -calculus. When using deBuijn indices, λ -abstractions are written *without* a variable name, and variables are written as the count of the number of λ -abstractions up in the syntax tree the variable is bound to. For example, $\lambda x. \lambda y. (x y)$ is written $\lambda \lambda (\$1 \$0)$ using deBuijn indices. See [33] for more details.

$$\begin{aligned}
\text{REFACTOR}'(e|\mathcal{D}) &= e, \text{ if } e \text{ is a leaf} & \text{REFACTOR}'(\lambda b|\mathcal{D}) &= \lambda \text{REFACTOR}(b|\mathcal{D}) \\
\text{REFACTOR}'(f \ x|\mathcal{D}) &= \text{REFACTOR}(f|\mathcal{D}) \text{ REFACTOR}(x|\mathcal{D}) & \text{REFACTOR}'(\oplus V|\mathcal{D}) &= \arg \min_{e \in \{\text{REFACTOR}(v|\mathcal{D}) : v \in V\}} \text{size}(e|\mathcal{D})
\end{aligned}$$

where $\text{size}(e|\mathcal{D})$ for program e and DSL \mathcal{D} is the size of the syntax tree of e , when members of \mathcal{D} are counted as having size 1. Concretely, $\text{REFACTOR}(v|\mathcal{D})$ calculates $\arg \min_{p \in \llbracket v \rrbracket} \text{size}(p|\mathcal{D})$.

Recall that our goal is to define an operator over version spaces, $I\beta_n$, which calculates the set of n -step refactorings of a program p , e.g., the set of all programs p' where $p' \xrightarrow{\leq n \text{ times}} q \xrightarrow{\leq n \text{ times}} \dots \xrightarrow{\leq n \text{ times}} p$, where $a \rightarrow b$ is the standard notation for a rewriting to b according to the standard rewrite rules of λ -calculus [33].

We define this operator in terms of another operator, $I\beta'$, which performs a single step of refactoring:

$$I\beta_n(v) = \oplus \left\{ \underbrace{I\beta'(I\beta'(I\beta'(\dots v)))}_{i \text{ times}} : 0 \leq i \leq n \right\}$$

where

$$I\beta'(u) = \oplus \{ (\lambda b)v : v \mapsto b \in S_0(u) \} \cup \begin{cases} \text{if } u \text{ is a primitive or index or } \emptyset: & \emptyset \\ \text{if } u \text{ is } \Lambda: & \{ \Lambda \} \\ \text{if } u = \lambda b: & \{ \lambda I\beta'(b) \} \\ \text{if } u = (f \ x): & \{ (I\beta'(f) \ x), (f \ I\beta'(x)) \} \\ \text{if } u = \oplus V: & \{ I\beta'(u') \mid u' \in V \} \end{cases}$$

where we have defined $I\beta'$ in terms of another operator, $S_k : \text{VS} \rightarrow 2^{\text{VS} \times \text{VS}}$, whose purpose is to construct the set of substitutions that are refactorings of a program in a version space. We define S as:

$$S_k(v) = \{ \downarrow_0^k v \mapsto \$k \} \cup \begin{cases} \text{if } v \text{ is primitive:} & \{ \Lambda \mapsto v \} \\ \text{if } v = \$i \text{ and } i < k: & \{ \Lambda \mapsto \$i \} \\ \text{if } v = \$i \text{ and } i \geq k: & \{ \Lambda \mapsto \$(i+1) \} \\ \text{if } v = \lambda b: & \{ v' \mapsto \lambda b' : v' \mapsto b' \in S_{k+1}(b) \} \\ \text{if } v = (f \ x): & \{ v_1 \cap v_2 \mapsto (f' \ x') : v_1 \mapsto f' \in S_k(f), v_2 \mapsto x' \in S_k(x) \} \\ \text{if } v = \oplus V: & \bigcup_{v' \in V} S_n(v') \\ \text{if } v \text{ is } \emptyset: & \emptyset \\ \text{if } v \text{ is } \Lambda: & \{ \Lambda \mapsto \Lambda \} \end{cases}$$

$$\downarrow_c^k \$i = \$i, \text{ when } i < c$$

$$\downarrow_c^k \$i = \$(i - k), \text{ when } i \geq c + k$$

$$\downarrow_c^k \$i = \emptyset, \text{ when } c \leq i < c + k$$

$$\downarrow_c^k \lambda b = \lambda \downarrow_{c+1}^k b$$

$$\downarrow_c^k (f \ x) = (\downarrow_c^k f \ \downarrow_c^k x)$$

$$\downarrow_c^k \oplus V = \oplus \{ \downarrow_c^k v \mid v \in V \}$$

$$\downarrow_c^k v = v, \text{ when } v \text{ is a primitive or } \emptyset \text{ or } \Lambda$$

where \uparrow^k is the shifting operator [33], which adds k to all of the free variables in a λ -expression or version space, and we have defined a new operator, \downarrow , whose purpose is to undo the action of \uparrow . We have written definitions recursively, but implement them using a dynamic program: we hash cons each version space, and only calculate the operators $I\beta_n$, $I\beta'$, and S_k once per each version space.

We now formally prove that $I\beta$ exhaustively enumerates the space of possible refactorings. Our approach is to first prove that S_k exhaustively enumerates the space of possible substitutions that could give rise to a program. The following pair of technical lemmas are useful; both are easily proven by structural induction.

Lemma 1. *Let e be a program or version space and n, c be natural numbers. Then $\uparrow_{n+c}^{-1} \uparrow_c^{n+1} e = \uparrow_c^n e$, and in particular $\uparrow_n^{-1} \uparrow^{n+1} e = \uparrow^n e$.*

Lemma 2. *Let e be a program or version space and n, c be natural numbers. Then $\downarrow_c^n \uparrow_c^n e = e$, and in particular $\downarrow^n \uparrow^n e = e$.*

Theorem 1. Consistency of S_n .

If $(v \mapsto b) \in S_n(u)$ then for every $v' \in v$ and $b' \in b$ we have $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' \in u$.

Proof. Suppose $b = \$n$ and therefore, by the definition of S_n , also $v = \downarrow_0^n u$. Invoking Lemmas 1 and 2 we know that $u = \uparrow_n^{-1} \uparrow^{n+1} v$ and so for every $v' \in v$ we have $\uparrow_n^{-1} \uparrow^{n+1} v' \in u$. Because $b = \$n = b'$ we can rewrite this to $\uparrow_n^{-1} [\$n \mapsto \uparrow^{n+1} v'] b' \in u$.

Otherwise assume $b \neq \$n$ and proceed by structural induction on u :

- If $u = \$i < n$ then we have to consider the case that $v = \Lambda$ and $b = u = \$i = b'$. Pick $v' \in \Lambda$ arbitrarily. Then $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' = \uparrow_n^{-1} \$i = \$i \in u$.
- If $u = \$i \geq n$ then we have to consider the case that $v = \Lambda$ and $b = \$(i+1) = b'$. Pick $v' \in \Lambda$ arbitrarily. Then $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' = \uparrow_n^{-1} \$(i+1) = \$i \in u$.
- If u is primitive then we have to consider the case that $v = \Lambda$ and $b = u = b'$. Pick $v' \in \Lambda$ arbitrarily. Then $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' = \uparrow_n^{-1} u = u \in u$.
- If u is of the form λa , then $S_n(u) \subset \{v \mapsto \lambda b \mid (v \mapsto b) \in S_{n+1}(a)\}$. Let $v \mapsto \lambda b \in S_n(u)$. By induction for every $v' \in v$ and $b' \in b$ we have $\uparrow_{n+1}^{-1} [\$n \mapsto \uparrow^{2+n} v'] b' \in a$, which we can rewrite to $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] \lambda b' \in \lambda a = u$.
- If u is of the form $(f \ x)$ then $S_n(u) \subset \{v_f \cap v_x \mapsto (b_f \ b_x) \mid (v_f \mapsto b_f) \in S_n(f), (v_x \mapsto b_x) \in S_n(x)\}$. Pick $v' \in v_f \cap v_x$ arbitrarily. By induction for every $v'_f \in v_f, v'_x \in v_x, b'_f \in b_f, b'_x \in b_x$ we have $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'_f] b'_f \in f$ and $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'_x] b'_x \in x$. Combining these facts gives $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] (b'_f \ b'_x) \in (f \ x) = u$.
- If u is of the form $\oplus U$ then pick $(v \mapsto b) \in S_n(u)$ arbitrarily. By the definition of S_n there is a z such that $(v \mapsto b) \in S_n(z)$, and the theorem holds immediately by induction.
- If u is \emptyset or Λ then the theorem holds vacuously.

□

Theorem 2. Completeness of S_n .

If there exists programs v' and b' , and a version space u , such that $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' \in u$, then there also exists $(v \mapsto b) \in S_n(u)$ such that $v' \in v$ and $b' \in b$.

Proof. As before we first consider the case that $b' = \$n$. If so then $\uparrow_n^{-1} \uparrow^{1+n} v' \in u$ or (invoking Lemma 1) that $\uparrow^n v' \in u$ and (invoking Lemma 2) that $v' \in \downarrow^n u$. From the definition of S_n we know that $(\downarrow^n u \mapsto \$n) \in S_n(u)$ which is what was to be shown.

Otherwise assume that $b' \neq \$n$. Proceeding by structural induction on u :

- If $u = \$i$ then, because b' is not $\$n$, we have $\uparrow_n^{-1} b' = \$i$. Let $b' = \$j$, and so

$$i = \begin{cases} j & \text{if } j < n \\ j - 1 & \text{if } j > n \end{cases}$$

where $j = n$ is impossible because by assumption $b' \neq \$n$.

If $j < n$ then $i = j$ and so $u = b'$. By the definition of S_n we have $(\Lambda \mapsto \$i) \in S_n(u)$, completing this inductive step because $v' \in \Lambda$ and $b' \in \$i$. Otherwise assume $j > n$ and so $\$i = \$(j-1) = u$. By the definition of S_n we have $(\Lambda \mapsto \$(i+1)) \in S_n(u)$, completing this inductive step because $v' \in \Lambda$ and $b' = \$j = \$(i+1)$.

- If u is a primitive then, because b' is not $\$n$, we have $\uparrow_n^{-1} b' = u$, and so $b' = u$. By the definition of S_n we have $(\Lambda \mapsto u) \in S_n(u)$ completing this inductive step because $v' \in \Lambda$ and $b' = u$.
- If u is of the form λa then, because of the assumption that $b' \neq \$n$, we know that b' is of the form $\lambda c'$ and that $\lambda \uparrow_{n+1}^{-1} [\$ (n+1) \mapsto \uparrow^{2+n} v'] c' \in \lambda a$. By induction this means that there is a $(v \mapsto c) \in S_{n+1}(a)$ satisfying $v' \in v$ and $c' \in c$. By the definition of S_n we also know that $(v \mapsto \lambda c) \in S_n(u)$, completing this inductive step because $b' = \lambda c' \in \lambda c$.
- If u is of the form $(f \ x)$ then, because of the assumption that $b' \neq \$n$, we know that b' is of the form $(b'_f \ b'_x)$ and that both $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b'_f \in f$ and $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b'_x \in x$. Invoking the inductive hypothesis twice gives a $(v_f \mapsto b_f) \in S_n(f)$ satisfying $v' \in v_f$, $b'_f \in b_f$ and a $(v_x \mapsto b_x) \in S_n(x)$ satisfying $v' \in v_x$, $b'_x \in b_x$. By the definition of S_n we know that $(v_f \cap v_x \mapsto b_f \ b_x) \in S_n(u)$ completing the inductive step because v' is guaranteed to be in both v_f and v_x and we know that $b' = (b'_f \ b'_x) \in (b_f \ b_x)$.
- If u is of the form $\uplus U$ then there must be a $z \in U$ such that $\uparrow_n^{-1} [\$n \mapsto \uparrow^{1+n} v'] b' \in z$. By induction there is a $(v \mapsto b) \in S_n(z)$ such that $v' \in v$ and $b' \in v$. By the definition of S_n we know that $(v \mapsto b)$ is also in $S_n(u)$ completing the inductive step.
- If u is \emptyset or Λ then the theorem holds vacuously.

□

From these results the consistency and completeness of $I\beta$ follows:

Theorem 3. Consistency of $I\beta'$.

If $p \in \llbracket I\beta'(u) \rrbracket$ then there exists $p' \in \llbracket u \rrbracket$ such that $p \longrightarrow p'$.

Proof. Proceed by structural induction on u . If $p \in \llbracket I\beta'(u) \rrbracket$ then, from the definition of $I\beta'$ and $\llbracket \cdot \rrbracket$, at least one of the following holds:

- Case $p = (\lambda b')v'$ where $v' \in v$, $b' \in b$, and $v \mapsto b \in S_0(u)$: From the definition of β -reduction we know that $p \longrightarrow \uparrow^{-1} [\$0 \mapsto \uparrow^1 v'] b'$. From the consistency of S_n we know that $\uparrow^{-1} [\$0 \mapsto \uparrow^1 v'] b' \in u$. Identify $p' = \uparrow^{-1} [\$0 \mapsto \uparrow^1 v'] b'$.
- Case $u = \lambda b$ and $p = \lambda b'$ where $b' \in \llbracket I\beta'(b) \rrbracket$: By induction there exists $b'' \in \llbracket b \rrbracket$ such that $b' \longrightarrow b''$. So $p \longrightarrow \lambda b''$. But $\lambda b'' \in \llbracket \lambda b \rrbracket = \llbracket u \rrbracket$, so identify $p' = \lambda b''$.
- Case $u = (f \ x)$ and $p = (f' \ x')$ where $f' \in \llbracket I\beta'(f) \rrbracket$ and $x' \in \llbracket x \rrbracket$: By induction there exists $f'' \in \llbracket f \rrbracket$ such that $f' \longrightarrow f''$. So $(f' \ x') \longrightarrow (f'' \ x')$. But $(f'' \ x') \in \llbracket (f \ x) \rrbracket = \llbracket u \rrbracket$, so identify $p' = (f'' \ x')$.
- Case $u = (f \ x)$ and $p = (f' \ x')$ where $x' \in \llbracket I\beta'(x) \rrbracket$ and $f' \in \llbracket f \rrbracket$: Symmetric to the previous case.
- Case $u = \uplus U$ and $p \in \llbracket I\beta'(u') \rrbracket$ where $u' \in U$: By induction there is a $p' \in \llbracket u' \rrbracket$ satisfying $p' \longrightarrow p$. But $\llbracket u' \rrbracket \subseteq \llbracket u \rrbracket$, so also $p' \in \llbracket u \rrbracket$.
- Case u is an index, primitive, \emptyset , or Λ : The theorem holds vacuously.

□

Theorem 4. Completeness of $I\beta'$.

Let $p \longrightarrow p'$ and $p' \in \llbracket u \rrbracket$. Then $p \in \llbracket I\beta'(u) \rrbracket$.

Proof. Structural induction on u . If $u = \uplus V$ then there is a $v \in V$ such that $p' \in \llbracket v \rrbracket$; by induction on v combined with the definition of $I\beta'$ we have $p \in \llbracket I\beta'(v) \rrbracket \subseteq \llbracket I\beta'(u) \rrbracket$, which is what we were to show. Otherwise assume that $u \neq \uplus V$.

From the definition of $p \longrightarrow p'$ at least one of these cases must hold:

- Case $p = (\lambda b')v'$ and $p' = \uparrow^{-1} [\$0 \mapsto \uparrow^1 v'] b'$: Using the fact that $\uparrow^{-1} [\$0 \mapsto \uparrow^1 v'] b' \in \llbracket u \rrbracket$, we can invoke the completeness of S_n to construct a $(v \mapsto b) \in S_0(u)$ such that $v' \in \llbracket v \rrbracket$ and $b' \in \llbracket b \rrbracket$. Combine these facts with the definition of $I\beta'$ to get $p = (\lambda b')v' \in \llbracket (\lambda b)v \rrbracket \subseteq \llbracket I\beta'(u) \rrbracket$.

- Case $p = \lambda b$ and $p' = \lambda b'$ where $b \rightarrow b'$: Because $p' = \lambda b' \in \llbracket u \rrbracket$ and by assumption $u \neq \uplus V$, we know that $u = \lambda v$ and $b' \in \llbracket v \rrbracket$. By induction $b \in \llbracket I\beta'(v) \rrbracket$. Combine with the definition of $I\beta'$ to get $p = \lambda b \in \llbracket \lambda I\beta'(v) \rrbracket \subseteq \llbracket I\beta'(u) \rrbracket$.
- Case $p = (f \ x)$ and $p' = (f' \ x)$ where $f \rightarrow f'$: Because $p' = (f' \ x) \in \llbracket u \rrbracket$ and by assumption $u \neq \uplus V$ we know that $u = (a \ b)$ where $f' \in \llbracket a \rrbracket$ and $x \in \llbracket b \rrbracket$. By induction on a we know $f \in \llbracket I\beta'(a) \rrbracket$. Therefore $p = (f \ x) \in \llbracket (I\beta'(a) \ b) \rrbracket \subseteq \llbracket I\beta'((a \ b)) \rrbracket \subseteq \llbracket I\beta'(u) \rrbracket$.
- Case $p = (f \ x)$ and $p' = (f \ x')$ where $x \rightarrow x'$: Symmetric to the previous case.

□

Finally we have our main result:

Theorem 5. Consistency and completeness of $I\beta_n$. Let p and p' be programs. Then $p \xrightarrow{\leq n \text{ times}} q \rightarrow \dots \rightarrow p'$ if and only if $p \in \llbracket I\beta_n(p') \rrbracket$.

Proof. Induction on n .

If $n = 0$ then $\llbracket I\beta_n(p') \rrbracket = \{p'\}$ and $p = p'$; the theorem holds immediately. Assume $n > 0$.

If $p \xrightarrow{\leq n \text{ times}} q \rightarrow \dots \rightarrow p'$ then $q \xrightarrow{\leq n-1 \text{ times}} p'$; induction on n gives $q \in \llbracket I\beta_{n-1}(p') \rrbracket$. Combined with $p \rightarrow q$ we can invoke the completeness of $I\beta'$ to get $p \in \llbracket I\beta'(I\beta_{n-1}(p')) \rrbracket \subseteq \llbracket I\beta_n(p') \rrbracket$.

If $p \in \llbracket I\beta_n(p') \rrbracket$ then there exists a $i \leq n$ such that $p \in \llbracket I\beta'(I\beta'(I\beta'(\dots p')) \rrbracket$ (i times). If $i = 0$ then $p = p'$ and p reduces to p' in $0 \leq n$ steps. Otherwise $i > 0$ and $p \in \llbracket I\beta'(I\beta'(I\beta'(\dots p')) \rrbracket$ ($i-1$ times). Invoking the consistency of $I\beta'$ we know that $p \rightarrow q$ for a program $q \in \llbracket I\beta'(I\beta'(\dots p')) \rrbracket$ ($i-1$ times) $\subseteq \llbracket I\beta_{i-1}(p') \rrbracket$. By induction $q \xrightarrow{\leq i-1 \text{ times}} p'$, which combined with $p \rightarrow q$ gives $p \xrightarrow{\leq i \leq n \text{ times}} q \rightarrow \dots \rightarrow p'$. □

A.5.2 Tracking equivalences

A.5.3 Computational complexity of DSL learning

How long does each update to the DSL in Algorithm 4 take? Constructing the version spaces takes time linear in the number of programs (written P) in the beams (Algorithm 4, line 5), and, in the worst case, exponential time as a function of the number of refactoring steps n — but we bound the number of steps to be a small number (typically $n = 3$). Writing V for the number of version spaces, this means that V is $O(P2^n)$. The number of proposals (line 10) is linear in the number of distinct version spaces, so is $O(V)$. For each proposal we have to refactor every program (line 6), so this means we spend $O(V^2) = O(P^22^n)$ per DSL update. In practice this quadratic dependence on P (the number of programs) is prohibitively slow. We now describe a linear time approximation to the refactor step in Algorithm 4 based on beam search.

For each version space v we calculate a *beam*, which is a function from a DSL \mathcal{D} to a shortest program in $\llbracket v \rrbracket$ using primitives in \mathcal{D} . Our strategy will be to only maintain the top B shortest programs in the beam; throughout all of the experiments in this paper, we set $B = 10^6$, and in the limit $B \rightarrow \infty$ we recover the exact behavior of REFACTOR. The following recursive equations define how we calculate these beams; the set ‘proposals’ is defined in line 10 of

Algorithm 4, and \mathcal{D} is the current DSL:

$$\begin{aligned} \text{beam}_v(\mathcal{D}') &= \begin{cases} \text{if } \mathcal{D}' \in \text{dom}(b_v): & b_v(\mathcal{D}') \\ \text{if } \mathcal{D}' \notin \text{dom}(b_v): & \text{REFACTOR}(v, \mathcal{D}) \end{cases} \\ b_v &= \text{the } B \text{ pairs } (\mathcal{D}' \mapsto p) \text{ in } b'_v \text{ where the syntax tree of } p \text{ is smallest} \\ b'_v(\mathcal{D}') &= \begin{cases} \text{if } \mathcal{D}' \in \text{proposals and } e \in \mathcal{D}' \text{ and } e \in v: & e \\ \text{otherwise if } v \text{ is a primitive or index:} & v \\ \text{otherwise if } v = \lambda b: & \lambda \text{beam}_b(\mathcal{D}') \\ \text{otherwise if } v = (f \ x): & (\text{beam}_f(\mathcal{D}') \ \text{beam}_x \mathcal{D}') \\ \text{otherwise if } v = \oplus V: & \arg \min_{e \in \{b'_{v'}(\mathcal{D}') : v' \in V\}} \text{size}(e|\mathcal{D}') \end{cases} \end{aligned}$$

We calculate $\text{beam}_v(\cdot)$ for each version space using dynamic programming. Using a minheap to represent $\text{beam}_v(\cdot)$, this takes time $O(VB \log B)$, replacing the quadratic dependence on V (and therefore the number of programs, P) with a $B \log B$ term, where the parameter B can be chosen freely, but at the cost of a less accurate beam search.

After performing this beam search, we take only the top I proposals as measured by $-\sum_x \min_{p \in \mathcal{F}_x} \text{beam}_{v_p}(\mathcal{D}')$. We set $I = 300$ in all of our experiments, so $I \ll B$. The reason why we only take the top I proposals (rather than take the top B) is because parameter estimation (estimating θ for each proposal) is much more expensive than performing the beam search — so we perform a very wide beam search and then at the very end trim the beam down to only $I = 300$ proposals. Next, we describe our MAP estimator for the continuous parameters (θ) of the DSL.

A.5.4 Estimating the continuous parameters θ of a DSL

We use an EM algorithm to estimate the continuous parameters of the DSL, i.e. θ . Suppressing dependencies on \mathcal{D} , the EM updates are

$$\theta = \arg \max_{\theta} \log P(\theta) + \sum_x \mathbb{E}_{q_x} [\log \mathbb{P}[p|\theta]] \quad (4)$$

$$q_x(p) \propto \mathbb{P}[x|p] \mathbb{P}[p|\theta] \mathbb{1}[p \in \mathcal{F}_x] \quad (5)$$

In the M step of EM we will update θ by instead maximizing a lower bound on $\log \mathbb{P}[p|\theta]$, making our approach an instance of Generalized EM.

We write $c(e, p)$ to mean the number of times that primitive e was used in program p ; $c(p) = \sum_{e \in \mathcal{D}} c(e, p)$ to mean the total number of primitives used in program p ; $c(\tau, p)$ to mean the number of times that type τ was the input to sample in Algorithm 2 while sampling program p . Jensen’s inequality gives a lower bound on the likelihood:

$$\begin{aligned} & \sum_x \mathbb{E}_{q_x} [\log \mathbb{P}[p|\theta]] = \\ & \sum_{e \in \mathcal{D}} \log \theta_e \sum_x \mathbb{E}_{q_x} [c(e, p_x)] - \sum_{\tau} \mathbb{E}_{q_x} \left[\sum_x c(\tau, p_x) \right] \log \sum_{\substack{e: \tau' \in \mathcal{D} \\ \text{unify}(\tau, \tau')}} \theta_e \\ & = \sum_e C(e) \log \theta_e - \beta \sum_{\tau} \frac{\mathbb{E}_{q_x} [\sum_x c(\tau, p_x)]}{\beta} \log \sum_{\substack{e: \tau' \in \mathcal{D} \\ \text{unify}(\tau, \tau')}} \theta_e \\ & \geq \sum_e C(e) \log \theta_e - \beta \log \sum_{\tau} \frac{\mathbb{E}_{q_x} [\sum_x c(\tau, p_x)]}{\beta} \sum_{\substack{e: \tau' \in \mathcal{D} \\ \text{unify}(\tau, \tau')}} \theta_e \\ & = \sum_e C(e) \log \theta_e - \beta \log \sum_{\tau} \frac{R(\tau)}{\beta} \sum_{\substack{e: \tau' \in \mathcal{D} \\ \text{unify}(\tau, \tau')}} \theta_e \end{aligned}$$

where we have defined

$$\begin{aligned} C(e) &\triangleq \sum_x \mathbb{E}_{q_x} [c(e, p_x)] \\ R(\tau) &\triangleq \mathbb{E}_{q_x} \left[\sum_x c(\tau, p_x) \right] \\ \beta &\triangleq \sum_{\tau} \mathbb{E}_{q_x} \left[\sum_x c(\tau, p_x) \right] \end{aligned}$$

Crucially it was defining β that let us use Jensen's inequality. Recalling from the main paper that $P(\theta) \triangleq \text{Dir}(\alpha)$, we have the following lower bound on M-step objective:

$$\sum_e (C(e) + \alpha) \log \theta_e - \beta \log \sum_{\tau} \frac{R(\tau)}{\beta} \sum_{\substack{e: \tau' \in \mathcal{D} \\ \text{unify}(\tau, \tau')}} \theta_e \quad (6)$$

Differentiate with respect to θ_e , where $e : \tau$, and set to zero to obtain:

$$\frac{C(e) + \alpha}{\theta_e} \propto \sum_{\tau'} \mathbb{1} [\text{unify}(\tau, \tau')] R(\tau') \quad (7)$$

$$\theta_e \propto \frac{C(e) + \alpha}{\sum_{\tau'} \mathbb{1} [\text{unify}(\tau, \tau')] R(\tau')} \quad (8)$$

The above is our estimator for θ_e . The above estimator has an intuitive interpretation. The quantity $C(e)$ is the expected number of times that we used e . The quantity $\sum_{\tau'} \mathbb{1} [\text{unify}(\tau, \tau')] R(\tau')$ is the expected number of times that we *could* have used e . The hyperparameter α acts as pseudocounts that are added to the number of times that we used each primitive, and are not added to the number of times that we could have used each primitive.

We are only maximizing a lower bound on the log posterior; when is this lower bound tight? This lower bound is tight whenever all of the types of the expressions in the DSL are not polymorphic, in which case our DSL is equivalent to a PCFG and this estimator is equivalent to the inside/outside algorithm. Polymorphism introduces context-sensitivity to the DSL, and exactly maximizing the likelihood with respect to θ becomes intractable, so for domains with polymorphic types we use this estimator.

A.6 Recognition model training

Recall that our goal is to maximize either $\mathcal{L}^{\text{posterior}}$ or \mathcal{L}^{MAP} , defined as:

$$\begin{aligned} \mathcal{L}^{\text{posterior}} &= \mathcal{L}_{\text{Replay}}^{\text{posterior}} + \mathcal{L}_{\text{Fantasy}}^{\text{posterior}} & \mathcal{L}^{\text{MAP}} &= \mathcal{L}_{\text{Replay}}^{\text{MAP}} + \mathcal{L}_{\text{Fantasy}}^{\text{MAP}} \\ \mathcal{L}_{\text{Replay}}^{\text{posterior}} &= \mathbb{E}_{x \sim X} \left[\sum_{p \in \mathcal{F}_x} \frac{\mathbb{P}[x, p | \mathcal{D}, \theta] \log Q(p|x)}{\sum_{p' \in \mathcal{F}_x} \mathbb{P}[x, p' | \mathcal{D}, \theta]} \right] & \mathcal{L}_{\text{Replay}}^{\text{MAP}} &= \mathbb{E}_{x \sim X} \left[\max_{\substack{p \in \mathcal{F}_x \\ p \text{ maxing } \mathbb{P}[\cdot | x, \mathcal{D}, \theta]}} \log Q(p|x) \right] \\ \mathcal{L}_{\text{Fantasy}}^{\text{posterior}} &= \mathbb{E}_{(p, x) \sim (\mathcal{D}, \theta)} [\log Q(p|x)] & \mathcal{L}_{\text{Fantasy}}^{\text{MAP}} &= \mathbb{E}_{x \sim (\mathcal{D}, \theta)} \left[\max_{\substack{p \\ p \text{ maxing } \mathbb{P}[\cdot | x, \mathcal{D}, \theta]}} \log Q(p) \right] \end{aligned}$$

The fantasy objectives are essential for data efficiency: all of our experiments train DREAMCODER on only a few hundred tasks, which is too little for a high-capacity neural network. Once we bootstrap a (\mathcal{D}, θ) , we can draw unlimited samples from (\mathcal{D}, θ) and train Q on those samples. But, evaluating $\mathcal{L}_{\text{Fantasy}}$ involves drawing programs from the current DSL, running them to get their outputs, and then training Q to regress from the input/outputs to the program. Since these programs map inputs to outputs, we need to sample the inputs as well. Our solution is to sample the inputs from the empirical observed distribution of inputs in X .

The $\mathcal{L}_{\text{Fantasy}}^{\text{MAP}}$ objective involves finding the MAP program solving a task drawn from the DSL. To make this tractable, rather than *sample* programs as training data for $\mathcal{L}_{\text{Fantasy}}^{\text{MAP}}$, we *enumerate* programs in decreasing order of their prior probability, tracking, for each dreamed task x , the set of enumerated programs maximizing $\mathbb{P}[x, p | \mathcal{D}, \theta]$.

We parameterize Q using a bigram model over syntax trees. Formally, Q predicts a $(|\mathcal{D}| + 2) \times (|\mathcal{D}| + 1) \times A$ -dimensional tensor, where A is the maximum arity⁵ of any primitive in the DSL. Slightly abusing notation, we write this tensor as $Q_{ijk}(x)$, where x is a task, $i \in \mathcal{D} \cup \{\text{start}, \text{var}\}$, $j \in \mathcal{D} \cup \{\text{var}\}$, and $k \in \{1, 2, \dots, A\}$. The output $Q_{ijk}(x)$ controls the probability of sampling primitive j given that i is the parent node in the syntax tree and we are sampling the k^{th} argument. Algorithm 5 specifies a procedure for drawing samples from $Q(\cdot|X)$.

Algorithm 5 Drawing from distribution over programs predicted by recognition model. Compare w/ Algorithm 2

```

1: function recognitionSample( $Q, x, \mathcal{D}, \tau$ ):
2:   Input: recognition model  $Q$ , task  $x$ , DSL  $\mathcal{D}$ , type  $\tau$ 
3:   Output: a program whose type unifies with  $\tau$ 
4:   return recognitionSample'( $Q, x, \text{start}, 1, \mathcal{D}, \emptyset, \tau$ )

5: function recognitionSample'( $Q, x, \text{parent}, \text{argumentIndex}, \mathcal{D}, \mathcal{E}, \tau$ ):
6:   Input: recognition model  $Q$ , task  $x$ , DSL  $\mathcal{D}$ ,  $\text{parent} \in \mathcal{D} \cup \{\text{start}, \text{var}\}$ ,  $\text{argumentIndex} \in \mathbb{N}$ , environment  $\mathcal{E}$ , type  $\tau$ 
7:   Output: a program whose type unifies with  $\tau$ 
8:   if  $\tau = \alpha \rightarrow \beta$  then                                     ▷ Function type — start with a lambda
9:      $\text{var} \leftarrow$  an unused variable name
10:     $\text{body} \sim$  recognitionSample'( $Q, x, \text{parent}, \text{argumentIndex}, \mathcal{D}, \{\text{var} : \alpha\} \cup \mathcal{E}, \beta$ )
11:    return (lambda ( $\text{var}$ )  $\text{body}$ )
12:   else                                                         ▷ Build an application to give something w/ type  $\tau$ 
13:      $\text{primitives} \leftarrow \{p | p : \tau' \in \mathcal{D} \cup \mathcal{E} \text{ if } \tau \text{ can unify with } \text{yield}(\tau')\}$       ▷ Everything in scope w/ type  $\tau$ 
14:      $\text{variables} \leftarrow \{p | p \in \text{primitives and } p \text{ a variable}\}$ 
15:     Draw  $e \sim \text{primitives}$ , w.p.  $\propto \begin{cases} Q_{\text{parent}, e, \text{argumentIndex}}(x) & \text{if } e \in \mathcal{D} \\ Q_{\text{parent}, \text{var}, \text{argumentIndex}}(x) / |\text{variables}| & \text{if } e \in \mathcal{E} \end{cases}$ 
16:     Unify  $\tau$  with  $\text{yield}(\tau')$ .                                     ▷ Ensure well-typed program
17:      $\text{newParent} \leftarrow \begin{cases} e & \text{if } e \in \mathcal{D} \\ \text{var} & \text{if } e \in \mathcal{E} \end{cases}$ 
18:      $\{\alpha_k\}_{k=1}^K \leftarrow \text{args}(\tau')$ 
19:     for  $k = 1$  to  $K$  do                                           ▷ Recursively sample arguments
20:        $a_k \sim$  recognitionSample'( $Q, x, \text{newParent}, k, \mathcal{D}, \mathcal{E}, \alpha_k$ )
21:     end for
22:     return ( $e \ a_1 \ a_2 \ \dots \ a_K$ )
23:   end if

```

Symmetry breaking. Why does the combination of \mathcal{L}^{MAP} and the bigram parameterization lead to symmetry breaking? The reason is twofold: (1) the objective \mathcal{L}^{MAP} prefers symmetry breaking recognition models; and (2) the bigram parameterization permits certain kinds of symmetry breaking. To sharpen these intuitions, we prove (Theorem 6) that any global optimizer of \mathcal{L}^{MAP} breaks symmetries, and then give a concrete worked out example contrasting the behavior of \mathcal{L}^{MAP} and $\mathcal{L}^{\text{posterior}}$.

⁵The arity of a function is the number of arguments that it takes as input.

Theorem 6. Let $\mu(\cdot)$ be a distribution over tasks and let $Q^*(\cdot|\cdot)$ be a task-conditional distribution over programs satisfying

$$Q^* = \arg \max_Q \mathbb{E}_\mu \left[\max_{p \text{ maximizing } \mathbb{P}[\cdot|x, \mathcal{D}, \theta]} \log Q(p|x) \right]$$

where (\mathcal{D}, θ) is a generative model over programs. Pick a task x where $\mu(x) > 0$. Partition Λ into expressions that are observationally equivalent under x :

$$\Lambda = \bigcup_i \mathcal{E}_i^x \text{ where for any } p_1 \in \mathcal{E}_i^x \text{ and } p_2 \in \mathcal{E}_j^x: \mathbb{P}[x|p_1] = \mathbb{P}[x|p_2] \iff i = j$$

Then there exists an equivalence class \mathcal{E}_i^x that gets all the probability mass of Q^* – e.g., $Q^*(p|x) = 0$ whenever $p \notin \mathcal{E}_i^x$ – and there exists a program in that equivalence class which gets all of the probability mass assigned by $Q^*(\cdot|x)$ – e.g., there is a $p \in \mathcal{E}_i^x$ such that $Q^*(p|x) = 1$ – and that program maximizes $\mathbb{P}[\cdot|x, \mathcal{D}, \theta]$.

Proof. We proceed by defining the set of “best programs” – programs maximizing the posterior $\mathbb{P}[\cdot|x, \mathcal{D}, \theta]$ – and then showing that a best program satisfies $Q^*(p|x) = 1$. Define the set of best programs \mathcal{B}_x for the task x by

$$\mathcal{B}_x = \left\{ p \mid \mathbb{P}[p|x, \mathcal{D}, \theta] = \max_{p' \in \Lambda} \mathbb{P}[p'|x, \mathcal{D}, \theta] \right\}$$

For convenience define

$$f(Q) = \mathbb{E}_\mu \left[\max_{p \in \mathcal{B}_x} \log Q(p|x) \right]$$

and observe that $Q^* = \arg \max_Q f(Q)$.

Suppose by way of contradiction that there is a $q \notin \mathcal{B}_x$ where $Q^*(q|x) = \epsilon > 0$. Let $p^* = \arg \max_{p \in \mathcal{B}_x} \log Q^*(p|x)$. Define

$$Q'(p|x) = \begin{cases} 0 & \text{if } p = q \\ Q^*(p|x) + \epsilon & \text{if } p = p^* \\ Q^*(p|x) & \text{otherwise.} \end{cases}$$

Then

$$f(Q') - f(Q^*) = \mu(x) \left(\max_{p \in \mathcal{B}_x} \log Q'(p|x) - \max_{p \in \mathcal{B}_x} \log Q^*(p|x) \right) = \mu(x) (\log(Q^*(p^*|x) + \epsilon) - \log Q^*(p^*|x)) > 0$$

which contradicts the assumption that Q^* maximizes $f(\cdot)$. Therefore for any $p \notin \mathcal{B}_x$ we have $Q^*(p|x) = 0$.

Suppose by way of contradiction that there are two distinct programs, q and r , both members of \mathcal{B}_x , where $Q^*(q|x) = \alpha > 0$ and $Q^*(r|x) = \beta > 0$. Let $p^* = \arg \max_{p \in \mathcal{B}_x} \log Q^*(p|x)$. If $p^* \notin \{q, r\}$ then define

$$Q'(p|x) = \begin{cases} 0 & \text{if } p \in \{q, r\} \\ Q^*(p|x) + \alpha + \beta & \text{if } p = p^* \\ Q^*(p|x) & \text{otherwise.} \end{cases}$$

Then

$$\begin{aligned} f(Q') - f(Q^*) &= \mu(x) \left(\max_{p \in \mathcal{B}_x} \log Q'(p|x) - \max_{p \in \mathcal{B}_x} \log Q^*(p|x) \right) \\ &= \mu(x) (\log(Q^*(p^*|x) + \alpha + \beta) - \log Q^*(p^*|x)) > 0 \end{aligned}$$

which contradicts the assumption that Q^* maximizes $f(\cdot)$. Otherwise assume $p^* \in \{q, r\}$. Without loss of generality let $p^* = q$. Define

$$Q'(p|x) = \begin{cases} 0 & \text{if } p = r \\ Q^*(p|x) + \beta & \text{if } p = p^* \\ Q^*(p|x) & \text{otherwise.} \end{cases}$$

Then

$$f(Q') - f(Q^*) = \mu(x) \left(\max_{p \in \mathcal{B}_x} \log Q'(p|x) - \max_{p \in \mathcal{B}_x} \log Q^*(p|x) \right) = \mu(x) (\log(Q^*(p^*|x) + \beta) - \log Q^*(p^*|x)) > 0$$

which contradicts the assumption that Q^* maximizes $f(\cdot)$. Therefore $Q^*(p|x) > 0$ for at most one $p \in \mathcal{B}_x$. But we already know that $Q^*(p|x) = 0$ for any $p \notin \mathcal{B}_x$, so it must be the case that $Q^*(\cdot|x)$ places all of its probability mass on exactly one $p \in \mathcal{B}_x$. Call that program p^* .

Because the equivalence classes $\{\mathcal{E}_i^x\}$ form a partition of Λ we know that p^* is a member of exactly one equivalence class; call it \mathcal{E}_i^x . Let $q \in \mathcal{E}_j^x \neq \mathcal{E}_i^x$. Then because the equivalence classes form a partition we know that $q \neq p^*$ and so $Q^*(q|x) = 0$, which was our first goal: *any* program not in \mathcal{E}_i^x gets no probability mass.

Our second goal — that there is a member of \mathcal{E}_i^x which gets all the probability mass assigned by $Q^*(\cdot|x)$ — is immediate from $Q^*(p^*|x) = 1$.

Our final goal — that p^* maximizes $\mathbb{P}[\cdot|x, \mathcal{D}, \theta]$ — follows from the fact that $p^* \in \mathcal{B}_x$. \square

Notice that Theorem 6 makes no guarantees as to the cross-task systematicity of the symmetry breaking; for example, an optimal recognition model could associate addition to the right for one task and associate addition to the left on another task. *Systematic* breaking of symmetries must arise only as a consequence as the network architecture (i.e., it is more parsimonious to break symmetries the same way for every task than it is to break them differently for each task).

As a concrete example of symmetry breaking, consider an agent tasked with writing programs built from addition and the constants zero and one. A bigram parameterization of Q allows it to represent the fact that it should never add zero ($Q_{+,0,0} = Q_{+,0,1} = 0$) or that addition should always associate to the right ($Q_{+,+,0} = 0$). The \mathcal{L}^{MAP} training objective encourages learning these canonical forms. Consider two recognition models, Q_1 and Q_2 , and two programs in frontier \mathcal{F}_x , $p_1 = (+ (+ 1 1) 1)$ and $p_2 = (+ 1 (+ 1 1))$, where

$$\begin{aligned} Q_1(p_1|x) &= \frac{\epsilon}{2} & Q_1(p_2|x) &= \frac{\epsilon}{2} \\ Q_2(p_1|x) &= 0 & Q_2(p_2|x) &= \epsilon \end{aligned}$$

i.e., Q_2 breaks a symmetry by forcing right associative addition, but Q_1 does not, instead splitting its probability mass equally between p_1 and p_2 . Now because $\mathbb{P}[p_1|\mathcal{D}, \theta] = \mathbb{P}[p_2|\mathcal{D}, \theta]$ (Algorithm 2), we have

$$\begin{aligned} \mathcal{L}_{\text{real}}^{\text{posterior}}(Q_1) &= \frac{\mathbb{P}[p_1|\mathcal{D}, \theta] \log \frac{\epsilon}{2} + \mathbb{P}[p_2|\mathcal{D}, \theta] \log \frac{\epsilon}{2}}{\mathbb{P}[p_1|\mathcal{D}, \theta] + \mathbb{P}[p_2|\mathcal{D}, \theta]} = \log \frac{\epsilon}{2} \\ \mathcal{L}_{\text{real}}^{\text{posterior}}(Q_2) &= \frac{\mathbb{P}[p_1|\mathcal{D}, \theta] \log 0 + \mathbb{P}[p_2|\mathcal{D}, \theta] \log \epsilon}{\mathbb{P}[p_1|\mathcal{D}, \theta] + \mathbb{P}[p_2|\mathcal{D}, \theta]} = +\infty \\ \mathcal{L}_{\text{real}}^{\text{MAP}}(Q_1) &= \log Q_1(p_1) = \log Q_1(p_2) = \log \frac{\epsilon}{2} \\ \mathcal{L}_{\text{real}}^{\text{MAP}}(Q_2) &= \log Q_2(p_2) = \log \epsilon \end{aligned}$$

So \mathcal{L}^{MAP} prefers Q_2 (the symmetry breaking recognition model), while $\mathcal{L}^{\text{posterior}}$ reverses this preference.

How would this example work out if we did not have a bigram parameterization of Q ? With a unigram parameterization, Q_2 would be impossible to express, because it depends on local context within the syntax tree of a program. So even though the objective function would prefer symmetry breaking, a simple unigram model lacks the expressive power to encode it.

To be clear, our recognition model does not learn to break *every* possible symmetry in every possible DSL. But in practice we found that a bigrams combined with \mathcal{L}^{MAP} works well, and we use with this combination throughout the paper.




Figure 12: Full set of LOGO graphics tasks that we apply our system to




Figure 13: Full set of tower building tasks that we apply our system to

A.7 Full set of LOGO tasks

A.8 Full set of tower tasks

A.9 Learning from Scratch: Tasks and DSL

We gave our system the following primitives: `if`, `=`, `>`, `+`, `-`, `0`, `1`, `cons`, `car`, `cdr`, `nil`, and `is-nil`, all of which are present in some form in McCarthy’s 1959 Lisp [28].⁶ We furthermore allowed functions to call themselves, which we modeled using the Y combinator. We did not use the recognition model for this experiment: a bottom-up pattern recognizer is of little use for acquiring this abstract knowledge from less than two dozen problems.

Figure 14 shows the full set of tasks and the learned DSL.

A.10 Learned DSLs

Here we present representative DSLs learned by our model. DSL primitives discovered by the algorithm are prefixed with `#`. Variables are prefixed with `$`, and we adopt De Bruijn indices to model bound variables [33].

A.10.1 List processing

```
#(λ (λ (map (λ (index $0 $2)) (range ($0 (+ 1))))))
#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0))))))
#(+ 1 (+ 1 1))
#(λ (λ (fold $1 (cons $0 empty) (λ (λ (cons $1 $0))))))
#(+ 1 #(+ 1 (+ 1 1)))
#(λ (map (λ (if ($1 $0) (+ $0 1) 0))))
#(λ (cdr (cdr $0)))
#(λ (λ (fold (#(λ (cdr (cdr $0))) (#(λ (cdr (cdr $0))) $1)) $0 (λ (λ (cdr (#(λ (λ (fold $1 (cons
→ $0 empty) (λ (λ (cons $1 $0)))))) $0 (car $0))))))
#(λ (car (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $0 (λ (empty? (#(λ (λ
→ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $1 (λ (gt? $0 $1))))))
#(λ (λ (length (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $1 (λ (eq? $0
→ $1))))))
#(λ (λ (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $1 (λ (is-prime (+ $1 (mod
→ $0 #(+ 1 #(+ 1 (+ 1 1)))))))))
#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0))))))
#(λ (map (λ (mod $0 $1))))
#(λ (map (λ (eq? (mod $0 $1) 0))))
#(λ (gt? (mod $0 #(+ 1 (+ 1 1))) 0))
#(λ (λ (#(λ (car (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $0 (λ (empty?
→ (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $1 (λ (gt? $0 $1))))))
→ (#(λ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $1 (λ (gt? $1 (length (#(λ
→ (λ (fold $1 empty (λ (λ (if ($2 $1) (cons $1 $0) $0)))))) $2 (λ (gt? $1 $0)))))))))
```

A.10.2 Text editing

```
#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0))))))
#(λ (λ (cons '.' (cons $0 $1))))
#(λ (λ (map (λ (index $0 $2)) (range $0))))
#(+ 1)
#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ (#(+ 1) (if
→ (char-eq? $2 $1) 0 $0))))))
#(λ (λ (cons (car $0) (#(λ (λ (cons '.' (cons $0 $1))) (cons '.' empty) (car $1))))))
#(λ (λ (map (λ (if (char-eq? $2 $0) $1 $0))))
#(λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) (cons LPAREN $0) (cons RPAREN empty)))
#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $0 (length (cdr (cdr $1))))))
#(λ (λ (cdr (fold (#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ
→ (#(+ 1) (if (char-eq? $2 $1) 0 $0)))))) $1 $0) $1 (λ (λ (cdr $0))))))
#(λ (#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $0 (length (cdr (cdr $1)))))) (cdr (cdr
→ $0))))
#(#(+ 1) 1)
#(λ (λ (cons (car $1) (cons $0 empty))))
```

⁶McCarthy’s first version of Lisp used `cond` instead of `if`. Because we are using a typed language, we instead used `if`, because Lisp-style `cond` is unwieldy to express as a function in typed languages.

Programs & Tasks	<pre> [1 9]→2 [5 3 8]→3 f(ℓ) = (f₅ ℓ) [0 1 1 0 0]→[1 1] [9 0 8]→[9 8] f(ℓ) = (f₃ (eq? 0) ℓ) [2 1 4]→[2 1 4 0] [9 8]→[9 8 0] f(ℓ) = (f₀ cons ℓ (cons 0 nil)) [2 5 6 0 6]→19 [9 2 7 6 3]→27 f(ℓ) = (f₀ + ℓ 0) [4 2 6 4]→[-4 -2 -6 -4] [2 3 0 7]→[-2 -3 -0 -7] f(ℓ) = (f₂ (- 0) ℓ) [1 5 2 9]→[1 2] [3 8 1 3 1 2]→[3 1 1] f(ℓ) = (f₆ (λ (l) (cdr (cdr l))) empty? ℓ) 3→[0 1 2 3] 2→[0 1 2] f(n) = (f₉ (+ 1 n)) 0, [9 2 3]→9 3, [0 2 8 4 5 6]→4 f(n,l) = (f₁₀ l n) 3→[-3 -2 -1] 4→[-4 -3 -2 -1] f(n) = (f₈ 0 n) </pre>	<pre> [[2 1] []]→[1 0] [[[] [] [9 8 9 9]]→[0 0 4] f(ℓ) = (f₂ f₄ ℓ) [1 -1 0 2]→[1 2] [9 -5 5 0 8]→[9 5 8] f(ℓ) = (f₃ (gt? 1) ℓ) [2 1 4]→[2 1] [9 8]→[9] f(ℓ) = (f₆ cdr (λ (z) (empty? (cdr z))) ℓ) [4 2 6 4]→[8 4 12 8] [2 3 0 7]→[4 6 0 14] f(ℓ) = (f₂ (λ (x) (+ x x)) ℓ) [4 2 6 4]→[5 3 7 5] [2 3 0 7]→[3 4 1 8] f(ℓ) = (f₂ (+ 1) ℓ) 3→[0 1 2] 2→[0 1] f(n) = (f₉ n) [9 2]→[9 9 2 2] [1 2 3 4]→[1 1 2 2 3 3 4 4] f(l) = (f₀ (λ (a x) (cons x (cons x a))) l nil) 1, [9 2 3]→9 4, [0 2 8 1 5 6]→1 f(n,l) = (f₁₀ l (+ 1 n)) 2→[2 1] 4→[4 3 2 1] f(n) = (f₈ 0 (- 0 n)) </pre>
	<pre> f₀(f,l,x) = (if (empty? l) x (f (car l) (f₀ (cdr l)))) (f₀: fold) f₂(f,l) = (f₀ nil l (λ (x a) (cons (f x) a))) (f₂: map) f₄(f,p,n) = (f₁ p f (+ 1) n) (f₄: generalization of range) f₆(n,p,l) = (f₁ p n car l) (f₆: specialization of unfold) f₈(n,m) = (f₄ (λ (x) (- n x)) (eq? 0) m) (f₇: count downwards) f₁₀(l,n) = (car (f₀ (λ (a x) (cdr a)) (f₉ n) l)) (f₁₀: index) </pre>	<pre> f₁(p,f,n,x) = (if (p x) nil (cons (f x) (f₁ (n x)))) (f₁: unfold) f₃(f,l) = (f₀ nil l (λ (x a) (if (f x) a (cons x a)))) (f₃: filter) f₅(l) = (f₀ (λ (a x) (+ 1 a)) l 0) (f₅: length) f₇(p) = (f₄ (λ (x) x) p 0) (f₇: another generalization of range) f₉(n) = (f₇ (eq? n)) (f₉: range) </pre>

Figure 14: Bootstrapping a standard library of functional programming routines, starting from recursion along with primitive operations found in 1959 Lisp. Complete set of tasks and learned DSL shown. Learned DSL components are numbered in the order that they are learned, *i.e.*, the agent first learns fold, then unfold, then uses fold to define map, etc.

```

#(λ (λ (#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ (#(+ 1)
  ↳ (if (char-eq? $2 $1) 0 $0)))))) (#(λ (λ (cdr (fold (#(λ (λ (#(λ (λ (map (λ (index $0
  ↳ $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ (#(+ 1) (if (char-eq? $2 $1) 0 $0)))))) $1
  ↳ $0) $1 (λ (λ (cdr $0)))))) (#(λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) (cons LPAREN
  ↳ $0) (cons RPAREN empty))) $1) $0) $0)))
#(λ (λ (λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) $1 (cons $0 $2))))
#(λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) $0 STRING))
#(#+ 1) #(#+ 1) 1))
#(λ (#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ (#(+ 1) (if
  ↳ (char-eq? $2 $1) 0 $0)))))) (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) STRING $0)
  ↳ SPACE)))
#(λ (λ (#(λ (λ (λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) $1 (cons $0 $2)))) (cons $0 $1)))
#(λ (λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) STRING)
#(λ (λ (#(λ (λ (cons (car $0) (#(λ (λ (cons '.' (cons $0 $1)))) (cons '.' empty) (car $1))))
  ↳ (#(λ (λ (cdr (fold (#(λ (λ (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1)
  ↳ 0 (λ (λ (#(+ 1) (if (char-eq? $2 $1) 0 $0)))))) $1 $0) $1 (λ (λ (cdr $0)))))) $1 $0)
  ↳ $1)))
#(#+ 1) #(#+ 1) #(#+ 1) 1))
#(λ (λ (λ (#(λ (λ (λ (fold $1 $0 (λ (λ (cons $1 $0)))))) $1 (#(λ (λ (cdr (fold (#(λ (λ (#(λ (λ (map
  ↳ (λ (index $0 $2)) (range $0)))) $1 (fold (cdr $1) 0 (λ (λ (#(+ 1) (if (char-eq? $2 $1) 0
  ↳ $0)))))) $1 $0) $1 (λ (λ (cdr $0)))))) (#(λ (#(λ (λ (fold $1 $0 (λ (λ (cons $1
  ↳ $0)))))) $0 STRING)) (#(λ (λ (map (λ (index $0 $2)) (range $0)))) $2 #(+ 1) #(+ 1)
  ↳ #(+ 1) 1)))) $0))))

```

A.10.3 Graphics

```

#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0))))))
#(logo_DIVA logo-UA 4)
#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) logo_IFTY)
#(logo_PT (λ (logo_FWRT logo_UL logo_ZA $0)))
#(λ (λ (λ (logo_forLoop $1 (λ (λ (logo_FWRT (logo_MULL $2 $1) $4 $0))))))
#(λ (logo_forLoop 7 (λ (λ (#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) 7
  ↳ logo_epsA $2 $0))))
#(λ (#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) logo_IFTY) (logo_SUBA logo-UA
  ↳ logo_epsA) (logo_MULL logo_epsL $0)))
#(λ (logo_forLoop logo_IFTY (λ (λ (logo_FWRT $2 logo_epsA (logo_FWRT logo_epsL logo_epsA $0))))))
#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) 4 # (logo_DIVA logo-UA 4))
#(logo_DIVA logo-UA)
#(λ (λ (λ (logo_forLoop $1 (λ (λ (logo_GETSET $2 (logo_FWRT logo_ZL $4 $0))))))
#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) logo_IFTY) (logo_DIVA logo_epsA
  ↳ 2) logo_epsL)
#(λ (λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) logo_IFTY) logo_epsA logo_epsL)
#(logo_forLoop 3 (λ (λ (#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0))))))
  ↳ logo_IFTY) (logo_DIVA logo_epsA 2) logo_epsL) (logo_FWRT logo_ZL # (logo_DIVA logo-UA 4)
  ↳ $0))))
#(λ (#(λ (λ (λ (logo_forLoop $1 (λ (λ (logo_GETSET $2 (logo_FWRT logo_ZL $4 $0))))))
  ↳ (#(logo_DIVA logo-UA) $0) $0))
#(λ (#(logo_PT (λ (logo_FWRT logo_UL logo_ZA $0))) (#(λ (logo_forLoop logo_IFTY (λ (λ (logo_FWRT
  ↳ $2 logo_epsA (logo_FWRT logo_epsL logo_epsA $0)))))) logo_epsL $0)))
#(λ (#(logo_PT (λ (logo_FWRT logo_UL logo_ZA $0))) (logo_FWRT logo_UL # (logo_DIVA logo-UA 4)
  ↳ $0)))
#(λ (#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) logo_IFTY) logo_epsA
  ↳ (logo_MULL logo_epsL $0)))
#(logo_FWRT logo_UL logo-UA)
#(λ (#(λ (logo_forLoop 7 (λ (λ (#(λ (λ (λ (logo_forLoop $2 (λ (λ (logo_FWRT $2 $3 $0)))))) 7
  ↳ logo_epsA $2 $0)))) (logo_MULL logo_epsL $0)))

```

A.10.4 Towers

```

#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0))))))
#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0))))))
#(λ (tower_loopM $0 (λ (λ (left 4 (#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) $0))))))
#(λ (tower_loopM $0 (λ (λ (left 8 (#(λ (1x3 (right 4 (1x3 (1x3 $0)))))) (left 4 (#(λ (1x3
  ↳ (1x3 (right 4 (1x3 (1x3 $0)))))) (#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) $0))))))
#(λ (λ (λ (tower_loopM $0 (λ (λ (left 3 (3x1 (right 3 (3x1 (left 4 $0)))))))))
#(λ (tower_loopM $0 (λ (λ (left 8 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (left 2 (3x1
  ↳ $0))))))

```

```

#(λ (1x3 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (1x3 (#(λ (1x3 (left 4 (1x3 (right 2 (3x1
  ↪ $0)))))) (right 4 $0))))))
#(λ (tower_loopM $0 (λ (λ (1x3 $0))))
#(tower_loopM 4 (λ (λ (right 6 (3x1 $0))))
#(λ (tower_loopM $0 (λ (λ (#(λ (1x3 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (1x3 (#(λ (1x3
  ↪ (left 4 (1x3 (right 2 (3x1 $0)))))) (right 4 $0)))))) $0))))
#(λ (right 4 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (#(λ (1x3 (left 4 (1x3 (right 2 (3x1
  ↪ $0)))))) $0))))
#(λ (right 2 (#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) $0)))
#(λ (λ (#(λ (tower_loopM $0 (λ (λ (#(λ (1x3 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (1x3 (#(λ
  ↪ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) (right 4 $0)))))) $0)))) $0 (right 2 (#(λ
  ↪ (tower_loopM $0 (λ (λ (left 4 (#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) $0)))))) $0
  ↪ $1))))))
#(tower_loopM 5 (λ (λ (#(λ (right 2 (#(λ (1x3 (left 4 (1x3 (right 2 (3x1 $0)))))) $0))) $0))))
#(λ (tower_loopM $0 (λ (λ (tower_embed (λ (#(λ (λ (λ (tower_loopM $0 (λ (λ (left $3 (3x1 (right
  ↪ 3 (3x1 (left $4 $0))))))))) 3 6 3 $0)) $0))))))
#(λ (λ (tower_loopM $0 (λ (λ (tower_embed (λ (#(λ (λ (λ (tower_loopM $0 (λ (λ (left $3 (3x1
  ↪ (right 3 (3x1 (left $4 $0))))))))) 4 5 $4 $0)) $0))))))
#(λ (λ (λ (tower_loopM $0 (λ (λ (left $3 (3x1 (right 3 (3x1 (left $4 $0))))))))) 4 5)
#(λ (λ (right 6 (tower_embed $0 $1))))
#(λ (tower_loopM $0 (λ (λ (#(λ (right 4 (#(λ (1x3 (1x3 (right 4 (1x3 (1x3 $0)))))) (#(λ (1x3
  ↪ (left 4 (1x3 (right 2 (3x1 $0)))))) $0)))) $0))))))

```

A.10.5 Symbolic regression

```

#(λ (+. $0 REAL))
#(λ (#(λ (+. $0 REAL)) (*. $0 REAL)))
#(/. REAL)
#(λ (#(λ (+. $0 REAL)) (*. $0 REAL))) (*. (#(λ (+. $0 REAL)) $0) $0)))
#(λ (/ . (#(/. REAL) $0) $0))
#(λ (#(λ (#(λ (+. $0 REAL)) (*. $0 REAL))) (*. $0 (#(λ (#(λ (#(λ (+. $0 REAL)) (*. $0 REAL)))
  ↪ (*. (#(λ (+. $0 REAL)) $0) $0)))) $0))))
#(λ (#(/. REAL) (#(λ (+. $0 REAL)) $0)))
#(λ (#(λ (+. $0 REAL)) (*. (#(λ (#(λ (#(λ (+. $0 REAL)) (*. $0 REAL))) (*. (#(λ (+. $0 REAL))
  ↪ $0) $0))) $0) $0)))

```

A.11 Hyperparameters and training details

Neural net architecture The recognition model for domains with sequential structure (list processing, text editing, regular expressions) is a recurrent neural network. We use a bidirectional GRU [6] with 64 hidden units that reads each input/output pair; we concatenate the input and output along with a special delimiter symbol between them. We use a 64-dimensional vectors to embed symbols in the input/output. We MaxPool the final hidden unit activations in the GRU along both passes of the bidirectional GRU.

The recognition model for domains with 2D visual structure (LOGO graphics, tower building, and symbolic regression) is a convolutional neural network. We take our convolutional architecture from [39].

We follow the RNN/CNN by an MLP with 128 hidden units and a ReLU activation which then outputs the Q_{ijk} matrix described in A.6.

Neural net training We train our recognition models using Adam [20] with a learning rate of 0.001.

DREAMCODER Hyperparameters Due to the high computational cost we performed only an informal coarse hyperparameter search. The most important parameter is the enumeration timeout during the wake phase; domains that present more challenging program synthesis problems require either longer timeouts, more CPUs, or both.

Domain	Timeout	CPUs	Batch size	λ (A.5)	α (A.5.4)	Max frontier (A.2)
Lists	7m	64	10	1	30	5
Text	7m	64	10	1	30	5
Graphics	1h	128	40	1.5	30	5
Symbolic regression	2m	40	10	1	30	5
Towers	2m	64	10	1	30	5
Regexes	30m	64	40	1	30	5

A.12 Results on held out text patterns

font

TrainTest	Human	MAP
JPCLN034.png JPCLN115.png JPCLN103.png JPCLN049.png JPCNN030.png hline JPCLN060.png JPCLN093.png JPCNN031.png JPCNN054.png JPCNN066.png	JPC\u\u\d+\.	((\u)*JPCLN\d\d\d\.)*(\u)*(\.)* FJPCLN829.F VARNLJPCLN980.R fu11 OTJPCLN466.HMNHESCJPCLN391.O GE3H JPCLN707.W6
WHS4_107 WHS7_103 WHS7_105 WHS6_116 WHS9_85 hline WHS3_43 WHS7_134 MDG_0000000033 WHS7_139 WHS4_117	WHS\d_\d+	WHS\d_(\d)* WHS4_ WHS7_ WHS1_4 WHS8_7 WHS5_6
WSW NNW W NNE E hline ESE SSE S SSW NE	\u+	(N)*(\u)* NNNNSS N NFTQ
CONJ ADP VERB ADV . hline DET ADJ NUM PRON NOUN	(\u)+ \.	.(\u)* vMLP] T z

.1301411		
.2883103		
.1275464		
.1376095		\. \d\d\d\d\d\d\d\d
.2999201		.9709716
hline	\. \d+	.6990069
9.225009E-02		.7696710
.1675194		.0109181
.1039294		.3759246
.1376776		
.128372		
u43		
u10		
u16		
u58		u\d\d
u02		u09
hline	u\d\d	u67
u46		u61
u22		u74
u42		u06
u30		
u36		
PM		
PF		
RF		
DM		\u\u
GS		CO
hline	\u\u	CI
PE		RS
PB		TE
DO		IW
SO		
CT		
SB1		
OND		
NH		
CH		(\u)*(\d)*
PSC		ZERT6
hline	(\u)+\u\d?	C318
NGH		037
USB		EH
NASCC		GE8
EB		
WHP		

SR504			
HR530			
HR512			
HR559		\uR5\d\d	
HR514		AR535	
hline	\uR5\d\d	JR595	
HR526		GR584	
HR544		DR506	
HR520		NR532	
HR562			
HR506			
<hr/>			
-4.26			
-1.69			
-3.91			
-1.622		-(\d)*(\d)*\.\d\d)*(\d)*	
-1		-7	
hline	-\d(\.(\d)+)?	-.94	
-2.316		-0.667.1266.375.57	
-3.19		-3627.51	
-0.933		-7	
-1.568			
-0.96			
<hr/>			
MN			
SD			
WY			
NE		\u\u	
MO		KF	
hline	\u\u	FD	
MS		OX	
TX		VB	
KY		RZ	
NM			
DC			
<hr/>			
hu48AC54			
hu264A0A			
hu294056			
hu26B551		(\d)*hu(\d)*(\u)*(\d)*d(\u)*d(\u)*	hu
hu3CDB6A		hu6LT9633NVB5	hu
hline	hu\d(\d \u)+	hu9621	hu
hu5CABA7		1huA397Y	hu
hu24A473		hu67439DR	hu
huFAA0CF		5hu7A1BYB5JGJBYFDOY	hu
hu154164			
hu7DE7FD			
<hr/>			

AJDJ6SQUDORP6					
A35C4777EXZZQH					
A2C9H5QRLXXICT					
A1JVYYO7G56DS				A ((\d) * (\u) * (\d) * \d (\u) * \d (\u) *) *	A (
A382RYVVZOJ8PX				A	A8
hline	A (\d \u) **			A	A3
A3GTTANAYUVSDY				A02W5BPYWKL3R23799KZ	A9
AZIN9ATSFV0ZG				A	A4
A2YO4SCWAWNYBI				A	A9
A2Q7SZISE1RXQQ					
AARIXIQYNH38W					
JJ					
RBR					
EX					
NNS				(\u) *	
VBN				N	
hline	\u+				
end				A	
NN				I	
PRP\$					
MD					
WDT					
CA					
VA					
FL					
CO				\u\u	
KY				MC	
hline	\u\u			MR	
MD				GL	
MI				BF	
CT				SD	
OH					
TX					
IH					
AP					
RS					
F				(\u) *	
MX				E	
hline	\u\u?			BFO	
XR				UT	
UP					
NM					
PZ					
QO					

GA		
PA		
LA		
IL		\u\u
VA		IS
hline	\u\u	GJ
WI		AL
FL		CU
CA		WR
NC		
OH		
PHI		
USA		
SCO		
RUS		\u\u\u
SVK		YJG
hline	\u\u\u	TRH
NZL		SKF
GER		RMP
CHN		LEM
NIR		
ENG		
NP		
AO		
CP		
SPP		\u(P)*(\u)*
FP		PQL
hline	\u+	CPPPTHBWYIOPD
TR		Z
CS		UXZJ
BD		NPPPK
MAR		
FIS		
(715) 967-2697		
(608) 819-2220		
(920) 988-2524		
(608) 442-0253		(.)*(\d)*\ \d\d\d-\d\d\d\d
(262) 723-4043) 342-1123
hline	\(\d\d\d\ \d\d\d-\d\d\d\d	79) 301-8637
(920) 887-4282		201) 398-4839
(262) 930-5176		70062) 942-1282
(920) 743-7943		!f) 085-7021
(715) 485-8764		
(608) 873-5217		

rec-93-org		
rec-51-org		
rec-321-org		
rec-281-dup-0		(\d)*rec-(\d)*-\1\1(.)*
rec-227-org		rec--ac
hline	rec-\d\d\d?-(org) (dup-0)	rec--thd
rec-372-org		rec-7-az
rec-39-org		rec--uo
rec-13-org		rec-27-bdM
rec-355-org		
rec-198-org		
<hr/>		
N22		
N33		
N34		
N13		N\d\d
N23		N28
hline	N\d\d	N06
N2		N51
N42		N38
N19		N07
N27		
N20		
<hr/>		
#217		
#222		
#070		
#007		#\d\d\d
#215		#972
hline	#\d\d\d	#223
#054		#071
#221		#320
#062		#142
#065		
#204		
<hr/>		
A-05334-58432		
A-05333-58214		
C-05271-37204		
C-05276-38389		\u-05\d\d\d-\d\d\d\d\d
A-05334-58579		N-05828-53285
hline	A C-\d+--\d+	H-05036-73233
A-05276-38355		E-05897-08233
A-05334-58331		Y-05034-95605
C-05276-38616		U-05337-85329
C-05334-58599		
A-05323-55341		

A1XB03X4J35ATE		
AXPOBOA6RMCV0		
A1K8QNLYYYYX21W		
A2QKM3JUFWBSMO		A((\d)*\d(\u)*\d)*(\u)*(\d)*\d(\u)*\d(\u)
A1T0NW527WZT7G		A7VB0Z
	A(\u \d)++	A957HS0JE
hline		A76387844N61
A24LC97AU3QC7G		AW9I9
A1MUEKEQQVROE7		ASG023O2
A3LC6M2EMDBBXP		
A1K9QGO39NG7DV		
A100Y89FZO4J0B		
<hr/>		
(281) 460-9574		
(801) 473-6431		
(713) 478-4292		
(713) 591-1559		(.)*(\d)*\)\ \d\d\d-\d\d\d\d\d
(281) 460-2390		7) 868-0287
		<) 807-8191
hline	\(\d\d\d\) \d\d\d-\d\d\d\d	') 242-9023
(314) 800-3559) 084-2796
(956) 778-2268		y) 973-8086
(662) 687-1620		
(281) 382-0084		
(757) 784-5526		
<hr/>		
us13C03		
us13C13		
us13C11		
us13A08		us13\u\d\d
us13B06		us13T43
		us13D30
hline	us13\u\d\d	us13Q43
us13B05		us13S41
us13C04		us13E89
us13A10		
us13B10		
us13C01		
<hr/>		
E07000193		
E07000170		
E07000127		
E07000027		E07000\d\d\d
E07000225		E07000380
		E07000213
hline	E07000\d\d\d	E07000430
E06000005		E07000874
E07000196		E07000079
E06000021		
E07000010		
E07000219		
<hr/>		

kt831qh3476		
vw545jd0450		
st391cf4769		
vx806dr5738		\l\l\d\d(\d)*\l\l\d\d\d\d
rp023xr5328		dl93627ix2215
hline	\l+\d+\l+\d+	qp26bz8609
vx793tc4890		xd6126oh3403
qy821tc8518		ng77cc0994
zw345zp0788		fm19iw7172
bb896gq6656		
qc288yw9773		
DE		
SE		
BD		
NA		\u\u
CH		MN
hline	\u\u	JJ
GB		WS
NL		QH
FR		KH
UA		
US		
-00:16:05.9		
-00:19:52.9		
-00:33:24.7		
-00:44:02.3		-00:\d\d: (\d)*\.(\d)*
-00:24:25.0		-00:00:.
hline	-00:\d\d:\d\d.\d	-00:39:.64
-00:17:28.6		-00:24:5927.
+00:03:04.8		-00:09:5656027.
-00:22:06.4		-00:48:3.8
-00:10:05.6		
-00:36:31.5		
SRX897027		
SRX897025		
SRX897032		
SRX894622		SRX89\d\d\d\d
SRX897016		SRX894550
hline	SRX89\d+	SRX894555
SRX897019		SRX897405
SRX894623		SRX892357
SRX897038		SRX899625
SRX897022		
SRX897035		

[0.0175]		
[0.035]		
[0.0475]		
[0.03]		(\d)*[0\.0(\d)*]
[0.015]		[0.0]
hline	[0\.0\d+]	[0.018]
[0.055]		52[0.00767]
[0.1275]		5[0.0]
[]		2[0.0]
[0.01]		
[0.005]		
'02:13:00'		
'03:27:00'		
'09:33:00'		
'15:40:00'		(.)*\d\d(.)*\d\d:00'
'22:24:00'		9707:00'
hline	'\d\d:\d\d:00'	6627:00'
'17:07:00'		03qD39:00'
'02:10:00'		HY10m8\dJf49:00'
'05:10:00'		24d07:00'
'05:55:00'		
'13:05:00'		
S5795		
S5552		
S6506		
H5529		\u\d\d\d\d
H5859		K4185
hline	(S H)\d+	A2648
H1099		V5688
H5580		A4248
H2261		Z8884
H8145		
H2879		
r6		
r1		
r2		
v1		\l(\d)*
r		d4
hline	(r v)\d?	o1
g1		f
g3		s
i6		t
g5		
x		

KW-0497		
KW-0494		
KW-1216		
KW-0528		KW-\d\d\d\d
KW-0568		KW-5188
hline	KW-\d+	KW-4639
KW-1203		KW-4354
KW-0959		KW-9880
KW-1071		KW-4387
KW-0505		
KW-0499		
<hr/>		
22G		
2		
22EXC		
9		(\d)*(\u)*
BL		
hline	\d*\u*	SWSX
31		101X
34		
3		
N_W04		
09X		
<hr/>		
N00028139		
N00009573		
N00027605		
N00024852		N000\d\d\d\d\d
N00035483		N00071101
hline	N000\d+	N00039120
N00005582		N00034062
N00033443		N00049088
N00031129		N00034153
N00005282		
N00033054		
<hr/>		
-11.50%		
14.20%		
0.70%		
33.80%		(.)*(\d)*\.(\\d)*0%
6.70%		7.0%
hline	-?\d?\d\.\d\d%	\V.0%
35.50%		1522.0%
26.20%		6.0%
25.20%		5.0%
46.60%		
7.20%		
<hr/>		

NY			
IL			
MS			
VA			\u\u
FL			PJ
			ZH
hline	\u\u		RJ
LA			RY
WI			IQ
AZ			
GA			
TX			
<hr/>			
-1.33 (0.465)			
-0.432 (0.0241)			
0.158 (0.0246)			
0.196 (0.0222)			(.)*(\d)*\.(\d)* \ (0\.\d\d\d(\d)*(.)*
0.206 (0.0233)			.6 (0.646
			9.56 (0.970199y
hline	-?\d\.\d+(\ (0\.\d+)\))?		q93.8 (0.255778"
NA			: s~. (0.7385uu
-0.186 (0.0215)			.927 (0.941I
1.11 (0)			
0.163 (0)			
-0.523 (0.0227)			
<hr/>			
(210)			
(220)			
(41)			
(635)			\ (\d\d(\d)*\)
(38)			(229)
			(18591)
hline	\ (\d+)\)		(7948)
(219)			(37)
(25)			(34)
(42)			
(6)			
(201)			
<hr/>			
2.5			
4			
--			
1.5			(\d)*(.)*
2			314E
			92
hline	(\d(\.\d)?) (--)		'9
1			33
3.5			
5			
3			
4.5			
<hr/>			

DS_258KADVNLH			
DS_25ROGVOIRY			
DS_252PR5PS19			
DS_25KIYDUYV3			DS_25 ((\d) * \u) * (\u) * (\d) *
DS_25AN1EA3PV			DS_25W709D9W2HQ754
hline	DS_25 (\u \d) +		DS_2581
DS_25VTXMIN8H			DS_2560A00L
DS_257J9CUMKG			DS_257PUQ7
DS_254F8A4M6R			DS_250
DS_251O26B4QW			
DS_258JLNIZ94			
<hr/>			
Q			
KTA			
JA			
TR			(\u) *
GB			QAUC
hline	\u+		NBD
QH			LN
KDC			QN
QM			OMS
DE			
DD			
<hr/>			
ms025846			
ms019454			
ms012586			
ms008925			ms0\d\d\d\d\d
ms002850			ms083084
hline	ms0\d+		ms093969
ms012952			ms053361
ms017651			ms088378
ms007498			ms059591
ms015186			
ms019618			
<hr/>			
HS2			
FCI			
LGD			
RWSW			(\u) * (\d) *
HS1			
hline	\u+ \d?		6
RSW			P
SHWT			O
WCI			YXP KP X7
RTWT			
D05G			
<hr/>			

ManH.010		
ManH.004		
ManH.009		
ManH.025		ManH\ .0\d\d
ManH.014		ManH.086
hline	ManH.0\d\d	ManH.012
ManH.016		ManH.048
ManH.023		ManH.030
ManH.027		ManH.025
ManH.021		
ManH.015		
<hr/>		
AIUTP		
GSE		
SPED-ATRC		
SPED		\u\u\u ((\u) *-) * (\u) *
GSE-DO		JDW
hline	\u+ (-\u+) ?	HIYINBP
MISL		XPC---A
OCCD		QCUW
HGCDC		MRE--
COUN		
ELP		
<hr/>		
Y2015/1093		
Y2013/1010		
Y2014/1017		
Y2015/1421		(\d) *Y201 (\d) * /\d\d\d\d
Y2017/1162		Y2018717/3895
hline	Y201\d/\d\d\d\d	Y2012944/4926
Y2011/1011		397Y201/2148
Y2015/1152		0Y201/5956
Y2018/1096		Y20196/0924
Y2017/1148		
Y2017/1206		
<hr/>		
q0005_0003		
q0009_0003		
q0002		
q0001		(\d) *q000 (\d) * (_\d\d\d\d) *
q0009_0008		7q000
hline	q000\d (_000\d) ?	44q00088905
q0011_0002		3q0002
q0010_0009		013334q0008_6631
q0011_0001		2q00098_7401_6849_8164
q0007_0008		
q0009_0014		
<hr/>		

3626cpr1748		
3626bpb0033		
56656		
x197		(\d)*(\l)*(\d)*\d\d
12550pr0002		73bg42228
		20920
hline	\d*\l*\d*	8151
x196		572837
3626cpr1687		98191
20598pat01		
3626cpr1206		
3626bpr0496		
<hr/>		
-122.31189		
-122.317017		
-122.342205		
-122.338913		-122\.3\d\d\d(\d)*
-122.312805		-122.3279
		-122.3810
hline	-122.3\d+	-122.34146
-122.338203		-122.34835
-122.330304		-122.3927
-122.331777		
-122.354093		
-122.325249		
<hr/>		
SFTCB		
RFTCB		
TETXB		
GETXB		\u\uT\uB
PATXB		FZTNB
		XZTJB
hline	\u\uT\uB	ITTNB
JNTCB		SDTNB
MGTXB		TRTXB
JFTXB		
KSTXB		
HYTCB		
<hr/>		
ENGL281		
ENGL300		
ENGL247		
ENGL280		((\u)*(\d)*) (ENGL\d\d\d)
ENG281		
		ENGL873
hline	ENGL?\d\d\d	ENGL768
ENGL119		ENGL586
281		ENGL579
ENGL442		
ENGL461		
GWS281		
<hr/>		

MF			
MNE			
MENRW			
MH			(\u)*M(\u)*
DPM			DM
	M?\u+		M
hline			MAEECBR
MSNE			FTIZSMD
MIC			M
MST			
MLG			
MRPWH			
<hr/>			
BUS M 277			
BUS M 440			
BUS M 498			
BUS M 490R TTh			BUS M \d\d\d(.)*
BUS M 490R F			BUS M 94909
	BUS M \d\d\d.*		BUS M 108G
hline			BUS M 238
BUS M 462			BUS M 548
BUS M 390 MW			BUS M 108
BUS M 478A			
BUS M 581			
BUS M 481			
<hr/>			
NIC			
MKD			
POL			
SWZ			\u\u\u
SUR			LSP
	\u\u\u		CNA
hline			JVP
AUS			FDW
ISL			QPU
GMB			
NZL			
COL			
<hr/>			
L - ??			\u - ((\d)*(\d)*\.(\\d)*)*(.)*
L - 31.0 lbs.			T - 424.72.4
L - 10.0 lbs.			V - ,
S - 8.6 lbs.			R - ,
L - 25.2 lbs.			Z - 0.B
	L \u - (\\?) (\d?\d\.\d lbs\.)		O - .,
hline			
L - 29.0 lbs.			
S - 9.0 lbs.			
L - 23.0 lbs.			
L - 22.8 lbs.			
S - ??			
<hr/>			

SF			
SD			
CHI			
HOU			(\u)*
KC			
hline	\u+		AEPCT
BAL			
LA			
DEN			
WSH			
MIA			
MAM.OSBS.2014.06			
MAM.OSBS.2013.07			
MAM.OSBS.2013.09			
MAM.OSBS.2014.05			MAM\ .OSBS\ .201(\d)*\ .\d\d
MAM.OSBS.2014.11			MAM.OSBS.20190235.90
hline	MAM\ .OSBS\ .201\d\ .\d\d		MAM.OSBS.2010.02
MAM.OSBS.2014.08			MAM.OSBS.201.60
MAM.OSBS.2014.09			MAM.OSBS.201.39
MAM.OSBS.2014.07			MAM.OSBS.201.13
MAM.OSBS.2014.10			
MAM.OSBS.2013.08			
EIRE			
DT4 9TG			
SW19 2HX			
SK15 3HN			(\u)*(\d)*...\u
SP6 3LR			5N%-D
hline	(\u \d)+((\u \d)+)*		2;+~P
HP14 4NE			9,~~V
NG1 1PU			Z75a/dI
ST16 1DW			PMUQIZmP
RG8 0HL			
LN6 0EJ			
-1,287659			
-0,890684			
-2			
-0,807021			-(\d)*,\d\d\d\d)*(\d)*
-2,485698			-71
hline	-\d(,\d+)?		- ,30503
-1,698316			- ,110639882333,9884
-2,672197			-9
-2,86825			-5
-1,905426			
-1,8			

cat. 16		
cat. 32		
cat. 19		
cat. 43		cat\.\d\d
cat. 13		cat. 42
hline	cat\.\d\d	cat. 73
cat. 29		cat. 09
cat. 22		cat. 40
cat. 3		cat. 57
cat. 52		
cat. 58		
<hr/>		
YlBv6gDob-Y		
iz151VUZx_c		
k63Hhf4zJpM		
iCiGjzEV7VI		(.)*.....
jk0kYP5djs0		oc1n7E.xq
hline	.	Od\9t%!fY
q6pGedpPjSQ	+	JI ^2N;f\rCyM
uUJaQWWMn6E		\3@QovC>a !83;
ov-mXl5s-yU		Kf}NfS`u"sV
ktWSfk5vRLI		
QmvJBvPz4Uk		
<hr/>		
D12		
C50		
E36		
C65		\u\d\d
C54		Y46
hline	\u\d\d	C53
C75		S18
F28		P43
C70		U46
D23		
G14		
<hr/>		
C32		
R15		
S07		
F23		\u\d\d
F11		H09
hline	\u\d\d	E48
L18		A62
C25		N92
L05		K47
P18		
R27		
<hr/>		

A.17		
A.12		
A.73		
A.63		A\.\d\d
A.35		A.97
hline	A\.\d\d	A.44
A.36		A.92
A.10		A.67
A.101		A.86
A.74		
A.72		
Resp19		
Resp44		
Resp17		
Resp20		Resp\d\d
Resp28		Resp51
hline	Resp\d\d	Resp17
Resp4		Resp43
Resp43		Resp95
Resp14		Resp33
Resp47		
Resp48		
DEG F		
LBMOL/HR		
MMBTU/HR		
PSI		\u\u\u(.)*((\u)*/*)*(\u)*
PSIA		RVCO/D//
hline	\u+((/)\u+)?	KHUD<wtHJ
PPM		GGUUI
PH		XYX/
LB/HR		IKRJD
PSIG		
MOL %		
US \$ 2.95		
US \$ 2.99		
US \$ 0.60		
US \$ 2.50		US \$ (\d)*\.\d\d
US \$ 1.95		US \$ 67.26
hline	US \$ \d\.\d\d	US \$.82
US \$ 3.50		US \$.28
US \$ 2.25		US \$ 03.90
US \$ 1.99		US \$.50
US \$ 1.00		
US \$ 0.75		

Z:-0.53		
Z:-0.61		
Z:0.17		
Z:0.32		Z:.(\\d)*\\.\\d\\d
Z:-0.16		Z:>.47
hline	Z:-?0\\.\\d\\d	Z:V6.90
Z:1.13		Z:>7.57
Z:0.14		Z:S6826.91
Z:0.95		Z:}.57
Z:-0.09		
Z:1.14		
<hr/>		
t1_cvk5ckb		
t1_cvjpw05		
t1_cvjq9vq		
t1_cvkds8m		t1_cv\\l....
t1_cvk5fd5		t1_cvs=N\\~
hline	t1_cv(\\l \\d)+	t1_cvp,WnD
t1_cvjpo3j		t1_cvhsO/C
t1_cvjpxk		t1_cvf[]'e
t1_cvjq15y		t1_cvp8,H3
t1_cvjrtvs		
t1_cvk7s7w		
<hr/>		
OldRC7		
HSP1		
RC16		
BARD11		\\u(.)*(\\u)*(\\d)*
YM2		AP1
hline	(\\u \\l)+\\d+	A
BRP11		L
KBH7		Fa-
BRP10		Q@
LVA3		
NSP7		
<hr/>		
\$6.5100		
\$12.7100		
\$6.6500		
\$39.5800		\$ (\\d)*\\.\\d\\d00
\$6.5000		\$60.2600
hline	\$\\d+\\.\\d+	\$.7200
\$13.1000		\$.6000
\$13.2400		\$.2500
\$34.6200		\$.6900
\$12.6900		
\$13.1400		
<hr/>		

OLE10026		
OLE2380894		
OLE1000174542		
1227		(\u)*(\d)*\d\d\d\d
1216		64833
hline	(OLE)?\d+	H3752
OLE1000092502		24900
OLE111		XY20106
OLE1000009196		66436899
OLE1000082147		
OLE1000064676		
<hr/>		
I		
F		
G		
C		\u
J		X
hline	\u	Q
B		N
E		B
K		W
D		
H		
<hr/>		
EFO_0001656		
EFO_0000572		
EFO_0001654		
EFO_0005135		EFO_000\d\d\d\d
EFO_0002897		EFO_0000780
hline	EFO_000\d+	EFO_0002254
EFO_0001266		EFO_0005533
EFO_0000826		EFO_0001096
EFO_0003168		EFO_0009258
EFO_0004017		
EFO_0001040		
<hr/>		
\$42,644.00		
\$65,602.00		
\$120,232.00		
\$49,474.00		\$(\d)*,\d\d\d\d.00
\$20,000.00		\$,414.00
hline	\$(\d+ (, \d\d\d\d)*\.\d\d	\$,222.00
\$56,767.00		\$,179.00
\$70,091.00		\$,747.00
\$58,995.00		\$1,504.00
\$78,487.00		
\$62,669.00		
<hr/>		

HARIJAN TOLI			
MAJHARIA			
MUSHWER TOLI			
HARIJAN TOLA			\u\u\u\u\u\u\u((\u)* \u\u\u\u)* (
PATKHOULI			PJCSGGJM
hline	\u+ (\u+)*		GJBQWWD JOE QGY
DIWAN TOLI			LHNCPC RENQF
MASAHl			LSRQDDA RCUK
HARPUR			LYUOBC CGV PFC FYQT
MUSHER TOLI			
AMOLWA			
<hr/>			
XEU			
DEM			
ITL			
CHF			\u\u\u
ESP			UUD
hline	\u\u\u		LPK
FRF			LML
ZAR			COV
CAD			UNK
USD			
AUD			
<hr/>			
UKC1			
UKL2			
UKG3			
UKL1			UK\u\d
UKM3			UKD2
hline	UK\u\d		UKJ1
UKE4			UKZ5
UKF1			UKH4
UKI1			UKO7
UKG2			
UKC2			
<hr/>			
sri			
gci			
mgl			
aut			\l\l\l
rsa			prq
hline	\l\l\l		vzu
slo			ruk
syr			euc
mar			geg
lux			
est			
<hr/>			

C15605276		
C35998515		
C11206768		
C18899260		C\d\d\d\d\d\d\d\d
C44659230		C75510881
hline	C\d+	C14542640
C13575104		C57453712
C19537329		C34779350
C99047219		C93417076
C45423422		
C43634821		
<hr/>		
Tx303		
81-1		
N192		
H100		(\u)*(\d)*(.)*\d
NC298		SQ3
hline	(\u \l \d -)+\d+	V14
H91		g8
C49A		@1
NC230		21
B97		
NC362		
<hr/>		
Q2-2019		
Q3-2018		
Q3-2019		
Q4-2018		Q(\d)*-201(\d)*
Q3-2017		Q-201
hline	Q\d-201\d	Q6-2018
Q1-2019		Q-201
Q1-2017		Q998-201
Q4-2017		Q-201
Q4-2019		
Q2-2017		
<hr/>		
SS0339		
FM0225		
FM2500		
FM0001		\u\u\d\d\d\d
SL0304		FV7666
hline	\u\u\d\d\d\d	AK8120
SH0094		RT7855
US0084		YM7795
US0259		GC8212
SH0204		
SH0147		
<hr/>		

IHYP			
YBEZ			
YBEU			
IHYS			\u\u\u\u
YBFV			UNUE
hline	\u\u\u\u		QYNB
L48H			XRGT
CGCE			EEVQ
KGM7			RLEC
KGL6			
YBEX			
IMPC_ACS_030_001			
IMPC_ACS_039_001			
IMPC_ACS_013_001			
IMPC_ACS_032_001			(.)*((\u)* (\d)*) (IMPC_ACS_0\d\d)_ (
IMPC_CBC_037_001			J@ 7_
hline	IMPC_\u\u\u_\d\d\d_\d\d\d		/Z8G_
IMPC_OFD_025_001			8_
IMPC_GRS_013_001			IMPC_ACS_079_V
IMPC_IPG_006_001			(H{ ' IMPC_ACS_074_
IMPC_ACS_031_001			
IMPC_CBC_043_001			
#79			
#101			
#94/2			
#4/2/95/2			(\d)*#(\d)* (/ (\d)*)*(.)*
#8/110/3-2			#1
hline	#\d+ ((/ -)\d+)*		#33626
#7/110/3-2			05#/*
#59			71#3
#3/82/3-2			#j,h\
#5/3/83/5-2			
#111			
GBP			
CNY			
EUR			
RON			\u\u\u
CAD			AIZ
hline	\u\u\u		EGY
AUD			JZW
INR			HBL
SEK			SHB
JPY			
USD			

akejr;ekr a;kfv;qkev alsf v;er rads sd		((.)*\1\1)*
hline jbq;ekrq; svdsgv 234 adv;etv afndv ;qkev q	(. \1)*	
PC.2300.040 PE.0500.010 PE.0300.010 PC.2000.190 PC.1500.010		P(\u)*(\d)*\.\d\d00\.\d\d
hline AC.0500.040 VE.0200.020 LY.0600.010 PC.2300.080 CP.1700.010	P\u\.\d\d\d\d\.\d\d\d	P.1300.020 PKM.8900.690 PA.3500.770 P.5100.740 P.0000.200
A00031832 A00038616 A00034074 A00037081 A00028150		((\u)*(\d)*) (A0003\d\d\d\d\
hline A00021039 A00039655 A00039257 A00031145 A00039607	A000\d+	IB1 A00034271 A00038893 C A00036764
EPS UGIS ASAMST RHETOR EDUC		(\u)*\u\u\u
hline AMERSTD L & S POL SCI ANTHRO ESPM	\u+	WDJY IPLI UUXI JVP LQB

V06873		
V06904		
V06914		
V06892		V06\d\d\d
V06921		V06026
hline	V06\d+	V06882
V06899		V06975
V06916		V06417
V06886		V06619
V06878		
V06890		
<hr/>		
F053830644		
F053858396		
4052189334		
F053799626		.05\d\d\d\d\d\d\d
F053841887		<052875194
hline	F?\d+	\$050434929
F053818771		q051077843
4052189176]050076233
F053852090		z051086391
F053859122		
F053823308		
<hr/>		
-79.5034323		
-79.29815396		
-79.50463054		
-79.24852584		-79\.\d\d\d(\d)*\d\d\d\d
-79.28102863		-79.70554673
hline	-79.\d+	-79.0002918
-79.40948687		-79.182928490
-79.21697343		-79.87319359
-79.29675842		-79.48449197
-79.5440321		
-79.34727006		
<hr/>		
als15		
Ssv15		
eb1613		
sg3415		(.)*l\l(\d)*1(\d)*
dt2315		eT"ve15
hline	(\u \l)+\d+	yi481
yg5615		9jf91
Dl3715		qx943417810
EM2815		fw6012
go114		
Hb915		
<hr/>		

POL-BRA		
BUL-POL		
POL-ARG		
RUS-IRI		\u\u\u-\u\u\u
BRA-FRA		YPM-NQF
hline	\u\u\u-\u\u\u	FFC-IQD
BUL-ARG		OKN-HGA
SRB-ITA		NKT-BOA
KOR-NED		KHS-OZZ
IRI-GER		
GER-RUS		
IL		
MO		
PA		
MI		\u\u
IA		TM
hline	\u\u	BX
AL		HT
IN		XP
CA		KC
DC		
MS		
0.5453040485		
1.1706019576		
0.1258506529		
0.3415773243		(.)*(\d)*\.(\d)*\d\d\d\d\d\d
-0.095009156		.290718619
hline	-\?d\.\d*	5.950453162
-0.455099117		G5.864143376
1.4648765223		1126.69897048776
0.2878761938		.88889553
1.2823425934		
6.1876327613		
Q65		
Q57		
Q71		
Q48		Q\d\d
Q49		Q05
hline	Q\d\d	Q02
Q68		Q59
Q64		Q40
Q47		Q07
Q58		
Q46		

4FF7DE80F4192276		
8F23E84FDD8DF9D0		
4807686BED3992FB		
53735C581B1DC0C0		((\u)*(\d)*\d(\u)*\d(\u)*,
74236837A4AFD042		
hline	(\d \u)+	6PRQKC1LD1EIUH4ABVEUREZB13
2257DD17485E53DF		Z78QC6R5U3PGPI
F9AED5486FD85258		5OCX1FRHOTG6877UT784X
EB9B653A49B15F67		
2D827450737DD778		
2C2DD9A93B415A54		
<hr/>		
6,820		
155,515		
117,127		
57,185		(\d)*,((\d)*,)*\d\d\d
6,617,103		6,,55,482
		,07,534
hline	\d+(,\d\d\d)+	294,038
36,984		,313
-		9,6,090
39,895		
124,168		
5,833,587		
<hr/>		
Oct 23 2013		
Jul 12 2010		
Sep 17 2012		
Jun 1 2001		\u\l\l (\d)* 20(\d)*
Nov 1 2004		Mvk 202
		Okf 20
hline	\u\l\l \d+ \d\d\d\d	Oae 0 20
Oct 10 2007		Vty 956 20
Mar 15 2016		Vyl 205
Nov 10 2011		
Oct 27 2005		
Jul 9 2015		
<hr/>		
FOS21001		
FOS20606		
FOS102		
FOS20803		FOS\d\d\d(\d)*
FOS20801		FOS203
		FOS7401
hline	FOS\d\d+	FOS8123144
FOS20904		FOS387
FOS1013		FOS8806
FOS108		
FOS140		
FOS20910		
<hr/>		

QWVvx22%{- Enrzla5^?@ FMBmr96#^[JTfvw6, ({/ Oh44>#?*%)			\u..... AlJ.au\$K3W NSdl].yag\$ Ov8+wks=*L RrQn_BDw [FUz jULVtWJ
hline Yuaz4>!_\$, FZfyve60#[CNWNdt076- JZrvo644_\$ HFqrq6 { }]	.+		
S0000214 S0000509 S0000411 S0000915 S0001042			((\u)*(\d)*) (S0000\d\d\d S0000207 SAAGX
hline S0001802 S0000685 S0000892 S0001472 S0000249	S000\d+		Z6451 D
S19000768 S19000884 S19000755 S19001103 S19000824			S1900\d\d\d\d S19005389 S19001985 S19007209 S19008380 S19006305
hline S19000776 S19000932 S19000981 S19001048 S19001110	S1900\d+		
\$31,800 \$35,400 \$37,200 \$24,000 \$31,200			\$(\d)*,(\d)*00(\d)* \$7,00 \$76,748008 \$22,00 \$8,26000 \$22,00
hline \$46,200 \$33,000 \$29,400 \$20,400 \$37,800	\$\d\d(,\d\d\d)+		

nIqsz Pwek1 bTFer 6h7Fs fg6ac		 QEbSb "xavP >rC-m L\#H> !\ra
hline sh6vJ hs2Yo kC8TK R0KY2 vE3us	(\u \l \d)+		
(66)28X8007 2816/180406 28X8007 4912/130100 28X8007.25			((.*)*\d\d)* 117 88
hline 28X8007.65 4910/110100 2812/130406 28X8007.14 28X8007.16	(\d \u \. \/ \(\))+		
R 340.18 R 995.25 R 383.50 R 530.10 R 194.94			R \d\d\d(\d)*\.\d\d R 15089.30 R 8358.10 R 007.68 R 369.89 R 8241.45
hline R 116.94 R 936.00 R 310.70 R 474.24 R 648.70	R \d\d\d.\d\d		
-0.7302 -0.7414 -0.7348 -0.7274 -0.729			-0\.\7\d\d(\d)* -0.77832 -0.7513 -0.75234 -0.7870 -0.7633
hline . -0.7632 -0.7286 -0.7256 -0.753	-0.7\d+		

TAS37818			
VIC30431			
TAS37825			
VIC30433			((\u)*(\d))* (\u\u\u3\d\d\d
GA4			ZNC31215
			117
hline	\u+?\d+		7
SA0062407			OUZ30460
WA100146691			DEV35548
TAS37819			
GA1			
NSW39502			
COMISARIA 52			
COMISARIA 11			
COMISARIA 51			
COMISARIA 40			COMISARIA \d\d
COMISARIA 14			COMISARIA 74
hline	COMISARIA \d\d		COMISARIA 58
COMISARIA 3			COMISARIA 78
COMISARIA 35			COMISARIA 95
COMISARIA 10			COMISARIA 80
COMISARIA 20			
COMISARIA 17			
K60			
N18			
R59			
D12			\u\d\d
I84			X97
hline	\u\d\d		A24
I20			S11
A04			H61
M13			B34
S82			
A90			
MPAS, PA-C			
DNP			
F.N.P.			
N.P.			((\u)*\.)*(\u)*P(.*
FP			.JP
hline	(.* \u\.?)+		SP2f
M.D.			I.ZPw`
CNP			W.BP
PA-C			V..R.P
FNP			
Dr.			

\$5.70		
\$3.40		
\$2.80		
\$5.40		$\$(\backslash d)^*\backslash.\backslash d\backslash d$
\$3.70		\$3.36
hline	$\$(\backslash d).\backslash d\backslash d$	\$42.74
\$3.60		\$.47
\$4.20		\$0.54
\$3.00		\$.01
\$5.60		
\$0.00		
R4703		
R3552		
R4708		
R3452		$R\backslash d\backslash d\backslash d\backslash d$
R4196		R0138
hline	$R\backslash d\backslash d\backslash d\backslash d$	R9715
R3650		R1064
R4707		R9971
R3547		R9908
R3574		
R3604		
P60004454		
P20005286		
P20005245		
P20005252		$P\backslash d000\backslash d\backslash d\backslash d\backslash d$
P20004669		P90009931
hline	$P\backslash d000\backslash d\backslash d\backslash d\backslash d$	P50005496
P80003353		P70000353
P80005572		P10002610
P20004263		P50006094
P20002903		
P20004479		
W 4-2		
W 4-3		
L 7-10		$\backslash u ((\backslash d)^-)*(\backslash u)^*(\backslash d)^*$
W 8-3		N OTZQ93
L 7-9		E 4-1-78
hline	$(W L) \backslash d-\backslash d^+$	K 76-691770---48
L 5-6		O
L 0-6		J 7--0-8
W 6-2		
L 2-7		
L 3-5		

MDEL18.1a		
MDEL1.2c		
MDEL18.1b		
MDEL6.2b		MDEL(\d)*\.(\\d)*\l
MDEL13.2b		MDEL.c
hline	MDEL\d\d?\.\d\l	MDEL7.0u
MDEL1.2d		MDEL.81m
MDEL13.2c		MDEL2.w
MDEL1.2e		MDEL4.30s
MDEL6.2d		
MDEL7.1a		
c04p0100a		
c04p01007		
c04p01009		
c04p0100c		c04p0100.
c04p01002		c04p01007
hline	c04p0100(\\l \\d)	c04p0100x
c04p01003		c04p0100\\
c04p0100o		c04p0100_
c04p0100e		c04p0100`
c04p0100m		
c04p0100f		
S6E-S-216M		
MC160		
I5-S-157M		
N4-S-38		(\\u)*(\\d)*((\\u)*-(\\d))**(\\u)
ZEN6-NL-SS-B		L
hline	(\\u \\d)+(-(\\u \\d))+*	5701YK-66786Z
S5-S-270		-R
N4-S-295		X
N920-NL-SS-B		689-8
S5-S-263M		
S6E-S-215M		
N1		
N7		
N8		
N5		N\\d
N2		N1
hline	N\\d	N1
N4		N4
N6		N0
N9		N5
N3		
N10		