
Learning Libraries of Subroutines for Neurally-Guided Bayesian Program Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Successful approaches to program induction require a hand-engineered domain-
2 specific language (DSL), constraining the space of allowed programs and imparting
3 prior knowledge of the domain. We contribute a program induction algorithm
4 called ECC that learns a DSL while jointly training a neural network to efficiently
5 search for programs in the learned DSL. We use our model to solve symbolic
6 regression problems, edit strings, and synthesize functions on lists, showing how
7 the model learns a domain-specific library of program components for expressing
8 solutions to problems in the domain.

9 1 Introduction

10 Automatically inducing programs from examples is a long-standing goal of artificial intelligence.
11 Recent work has successfully used symbolic search techniques (e.g., Metagol: [1], FlashFill: [2]),
12 neural networks trained from a corpus of examples (e.g., RobustFill: [3]), and hybrids of neural and
13 symbolic methods (e.g., Neural-guided deductive search: [4], DeepCoder: [5]) to synthesize programs
14 for task domains such as string transformations, list processing, and robot navigation and planning.
15 However, all these approaches – symbolic, neural and neural-symbolic – rely upon a hand-engineered
16 *Domain-Specific Language* (DSL). DSLs contain an inventory of restricted programming primitives,
17 encoding domain-specific knowledge about the space of programs. In practice we often have only a
18 few input/output examples for each program to be induced, and thus success often hinges on having a
19 good DSL that provides a crucial inductive bias for what would otherwise be an unconstrained search
20 through the space of all computable functions. Here we ask, to what extent can we dispense with
21 such highly hand-engineered domain-specific languages?

22 We propose *learning* the DSL by inducing a library of domain-specific subroutines. We consider
23 the setting where we have a collection of related programming tasks, each specified by a set of
24 input/output examples. Starting from a weaker or more general library of primitives, we give an
25 algorithm for constructing a richer, more powerful, and better-tuned DSL. Our algorithm is called
26 **Explore/Compress/Compile** (ECC), and iterates between three different steps: an **Explore** step
27 uses the DSL to explore the space of programs, searching for ones that solve the tasks; a **Compress**
28 step modifies the DSL by discovering regularities in the programs found by the previous Explore step;
29 and a **Compile** step, which improves the program search procedure by training a neural network to
30 write programs in the current DSL, in the spirit of “amortized” or “compiled” inference [6]. We call
31 the neural net a **recognition model** (c.f. Hinton 1995 [7]). The learned DSL distills commonalities
32 across programs that solve tasks, helping the agent solve related program induction problems. The
33 neural recognition model ensures that searching for programs remains tractable even as the DSL (and
34 hence the search space for programs) expands.

35 Because any model may be encoded as a (deterministic or probabilistic) program, we carefully
36 delineate the scope of program learning problems considered here. We think of ECC as learning to

solve the kinds of problems that humans can solve relatively quickly – once they acquire the relevant domain expertise. These correspond to short programs – once you have the right DSL. Even with the right DSL, program search may be intractable, so we amortize the cost of program search by training a neural network to assist the search procedure.

We apply ECC to three domains: symbolic regression; FlashFill-style [2] string processing; and Lisp-style functions on lists. For each of these we initially provide a generic set of programming primitives in a Lisp-like language. Our algorithm then discovers its own domain-specific vocabulary for expressing solutions in the domain (Tbl. 1).

2 The ECC Algorithm

Our goal is to induce a DSL while finding programs solving each of the tasks. We take inspiration primarily from the Exploration-Compression algorithm for bootstrap learning [8]. Exploration-Compression alternates between exploring the space of solutions to a set of tasks, and compressing those solutions to suggest new search primitives for the next exploration stage. We extend these ideas into an inference strategy that iterates through three steps: an **Explore** step uses the current DSL and recognition model to search for programs that solve the tasks. The **Compress** and **Compile** steps update the DSL and the recognition model, respectively. Crucially, these steps bootstrap off each other (Fig. 2):

Exploration: Searching for programs. Our program search is informed by both the DSL and the recognition model. When these improve, we find more programs solving the tasks.

Compression: Improving the DSL. We induce the DSL from the programs found in the exploration phase, aiming to maximally compress (or, raise the prior probability of) these programs. As we solve more tasks, we hone in on DSLs that more closely match the domain.

Compilation: Learning a neural recognition model. We update the recognition model by training on two data sources: samples from the DSL (as in the Helmholtz Machine’s “sleep” phase), and programs found by the search procedure during exploration. As the DSL improves and as search finds more programs, the recognition model gets more data to train on, and better data.

Domain	Part of the learned DSL
Regression	<code>(+ (* real x) real)</code> <i>(a linear function of x)</i>
Strings	<code>(map (lambda (x) (if (= x a) b x)))</code> <i>(replace occurrences of a w/ b)</i>
Lists	<code>(map (lambda (x) (+ x k)) l)</code> <i>(add k to every element of list l)</i>

Figure 1: Examples of structure found in DSLs learned by our algorithm. ECC builds a new DSL by discovering and reusing useful subroutines.




Figure 2: ECC solves for programs, the DSL, and a neural network (recognition model). Each of these steps iteratively bootstraps off of the others.

2.1 Hierarchical Bayesian Framing

ECC takes as input a set of *tasks*, written X , each of which is a program induction problem. It has at its disposal a *likelihood model*, written $\mathbb{P}[x|p]$, which scores the likelihood of a task $x \in X$ given a program p . Its goal is to solve each of the tasks by writing a program, and also to infer a DSL, written \mathcal{D} . We equip \mathcal{D} with a real-valued weight vector θ , and together (\mathcal{D}, θ) define a generative model over programs. We frame our goal as maximum a posteriori (MAP) inference of (\mathcal{D}, θ) given X .

85 Writing J for the joint probability of (\mathcal{D}, θ) and X , we want the \mathcal{D}^* and θ^* solving:

$$J(\mathcal{D}, \theta) \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_p \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \int J(\mathcal{D}, \theta) d\theta \quad \theta^* = \arg \max_{\theta} J(\mathcal{D}^*, \theta) \quad (1)$$

86 The above equations summarize the problem from the point of view of an ideal Bayesian learner.
 87 However, Eq. 1 is wildly intractable because evaluating $J(\mathcal{D}, \theta)$ involves summing over the infinite
 88 set of all programs. In practice we will only ever be able to sum over a finite set of programs. So, for
 89 each task, we define a finite set of programs, called a *frontier*, and only marginalize over the frontiers:
 90 **Definition.** A *frontier of task x* , written \mathcal{F}_x , is a finite set of programs s.t. $\mathbb{P}[x|p] > 0$ for all $p \in \mathcal{F}_x$.

91 Using the frontiers we define the following intuitive lower bound on the joint probability, called \mathcal{L} :

$$J \geq \mathcal{L} \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_{p \in \mathcal{F}_x} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] \quad (2)$$

92 ECC does approximate MAP inference by maximizing this lower bound on the joint probability,
 93 alternating maximization w.r.t. the frontiers (Exploration) and the DSL (Compression):

94 **Program Search: Maxing \mathcal{L} w.r.t. the frontiers.** Here (\mathcal{D}, θ) is fixed and we want to find new
 95 programs to add to the frontiers so that \mathcal{L} increases the most. \mathcal{L} most increases by finding programs
 96 where $\mathbb{P}[x, p|\mathcal{D}, \theta]$ is large.

97 **DSL Induction: Maxing $\int \mathcal{L} d\theta$ w.r.t. the DSL.** Here $\{\mathcal{F}_x\}_{x \in X}$ is held fixed, and so we can
 98 evaluate \mathcal{L} . Now the problem is that of searching the discrete space of DSLs and finding one
 99 maximizing $\int \mathcal{L} d\theta$. Once we have a DSL \mathcal{D} we can update θ to $\arg \max_{\theta} \mathcal{L}(\mathcal{D}, \theta, \{\mathcal{F}_x\})$.

100 Searching for programs is hard because of the large combinatorial search space. We ease this
 101 difficulty by training a neural recognition model, $q(\cdot|\cdot)$, during the compilation phase: q is trained to
 102 approximate the posterior over programs, $q(p|x) \propto \mathbb{P}[p|x, \mathcal{D}, \theta] = \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$, thus amortizing
 103 the cost of finding programs with high posterior probability.

104 **Neural recognition model: tractably maxing \mathcal{L} w.r.t. the frontiers.** Here we train a neural
 105 network, q , to predict a distribution over programs conditioned on a task. The objective of q is to
 106 assign high probability to programs p where $\mathbb{P}[x, p|\mathcal{D}, \theta]$ is large, because including those programs
 107 in the frontiers will most increase \mathcal{L} .

108 2.2 Exploration: Searching for Programs

109 Now our goal is to search for programs solving the tasks. We use the simple approach of enumerating
 110 programs from the DSL in decreasing order of their probability, and then checking if a program p
 111 assigns positive probability to a task ($\mathbb{P}[x|p] > 0$); if so, we incorporate p into the frontier \mathcal{F}_x .

112 To make this concrete we need to define what programs actually are and what form $\mathbb{P}[p|\mathcal{D}, \theta]$ takes.
 113 We represent programs as λ -calculus expressions. λ -calculus is a formalism for expressing functional
 114 programs that closely resembles the Lisp programming language. λ -calculus includes variables,
 115 function application, and the ability to create new functions. Throughout this paper we will write λ -
 116 calculus expressions in Lisp syntax. Our programs are all strongly typed. We use the Hindley-Milner
 117 polymorphic typing system [9] which is used in functional programming languages like OCaml. Type
 118 variables are always written using lowercase Greek letters and we write $\alpha \rightarrow \beta$ to mean a function
 119 that takes an input of type α and returns something of type β . We use the notation $p : \tau$ to mean that
 120 the λ -calculus expression p has the type τ . For example, to describe the type of the identity function
 121 we would say $(\text{lambda } (x) x) : \alpha \rightarrow \alpha$. We say a type α *unifies* with τ if every expression $p : \alpha$
 122 also satisfies $p : \tau$. Furthermore, the act of *unifying* a type α with τ is to introduce constraints on
 123 the type variables of α to ensure that α unifies with τ . See Supplement for more detail on program
 124 representation. With this notation in hand we now define DSLs:

125 **Definition:** (\mathcal{D}, θ) . A DSL \mathcal{D} is a set of typed λ -calculus expressions. A weight vector θ for a DSL
 126 \mathcal{D} is a vector of $|\mathcal{D}| + 1$ real numbers: one number for each DSL primitive $e \in \mathcal{D}$, written θ_e , and a
 127 weight controlling the probability of a variable occurring in a program, θ_{var} .

Algorithm 1 Generative model over programs

```

128 function sample( $\mathcal{D}, \theta, \mathcal{E}, \tau$ ):
129   Input: DSL ( $\mathcal{D}, \theta$ ), environment  $\mathcal{E}$ , type  $\tau$ 
130   Output: a program whose type unifies with  $\tau$ 
131   if  $\tau = \alpha \rightarrow \beta$  then
132     var  $\leftarrow$  an unused variable name
133     body  $\sim$  sample( $\mathcal{D}, \theta, \{\text{var} : \alpha\} \cup \mathcal{E}, \beta$ )
134     return (lambda (var) body)
135   end if
136   primitives  $\leftarrow \{p | p : \tau' \in \mathcal{D} \cup \mathcal{E}$ 
137     if  $\tau$  can unify with yield( $\tau'$ ) $\}$ 
138   Draw  $e \sim$  primitives, w.p.  $\propto \theta_e$  if  $e \in \mathcal{D}$ 
139     w.p.  $\propto \frac{\theta_{var}}{|\text{variables}|}$  if  $e \in \mathcal{E}$ 
140   Unify  $\tau$  with yield( $\tau'$ ).
141    $\{\alpha_k\}_{k=1}^K \leftarrow \text{args}(\tau')$ 
142   for  $k = 1$  to  $K$  do
143      $a_k \sim$  sample( $\mathcal{D}, \theta, \mathcal{E}, \alpha_k$ )
144   end for
145   return ( $e \ a_1 \ a_2 \ \dots \ a_K$ )
146   where:
147   yield( $\tau$ ) =  $\begin{cases} \text{yield}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ \tau & \text{otherwise.} \end{cases}$ 
148   args( $\tau$ ) =  $\begin{cases} [\alpha] + \text{args}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ [] & \text{otherwise.} \end{cases}$ 

```

The purpose of the recognition model is to amortize the cost of searching for programs. It does this by learning to predict programs which are probable under (\mathcal{D}, θ) while also assigning high likelihood to a task according to $\mathbb{P}[x|p]$. Concretely, the recognition model q is a neural network that predicts, for each task $x \in X$, a weight vector $q(x) = \theta^{(x)} \in \mathbb{R}^{|\mathcal{D}|+1}$. Together with the DSL, this defines a distribution over programs, $\mathbb{P}[p|\mathcal{D}, \theta = q(x)]$. We abbreviate this distribution as $q(p|x)$. The crucial aspect of this framing is that the neural network leverages the structure of the learned DSL, so it is *not* responsible for generating programs wholesale. We share this aspect with DeepCoder [5] and [15].

We want a recognition model that closely approximates the true posteriors over programs. We formulate this as minimizing the expected KL-divergence, $\mathbb{E}[\text{KL}(\mathbb{P}[p|x, \mathcal{D}, \theta] || q(p|x))]$, or equivalently maximizing $\mathbb{E}[\sum_p \mathbb{P}[p|x, \mathcal{D}, \theta] \log q(p|x)]$, where the expectation is taken over tasks. One could take this expectation over the empirical distribution of tasks, like how an autoencoder is trained [16]; or, one could take this expectation over samples from the generative model, like how a Helmholtz machine is trained [17]. We found it useful to maximize both an autoencoder-style objective \mathcal{L}_{AE} and a Helmholtz-style objective \mathcal{L}_{HM} , giving the objective for a recognition model, $\mathcal{L}_{\text{RM}} = \mathcal{L}_{\text{AE}} + \mathcal{L}_{\text{HM}}$:

$$\mathcal{L}_{\text{HM}} = \mathbb{E}_{(p,x) \sim (\mathcal{D}, \theta)} [\log q(p|x)] \quad \mathcal{L}_{\text{AE}} = \mathbb{E}_{x \sim X} \left[\sum_{p \in \mathcal{F}_x} \frac{\mathbb{P}[x, p|\mathcal{D}, \theta]}{\sum_{p' \in \mathcal{F}_x} \mathbb{P}[x, p'|\mathcal{D}, \theta]} \log q(p|x) \right]$$

The \mathcal{L}_{HM} objective is essential for data efficiency: all of our experiments train ECC on only a few hundred tasks, which is too little for a high-capacity neural network q . Once we bootstrap a (\mathcal{D}, θ) , we can draw unlimited samples from (\mathcal{D}, θ) and train q on those samples.

Evaluating \mathcal{L}_{HM} involves sampling programs from the current DSL, running them to get their outputs, and then training q to regress from the input/outputs to the program. Since these programs map inputs to outputs, we need to sample the inputs as well. Our solution is to sample the inputs from the empirical observed distribution of inputs in X .

2.4 Compression: Learning a Generative Model (a DSL)

The purpose of the DSL is to offer a set of abstractions that allow an agent to easily express solutions to the tasks at hand. Intuitively, we want the algorithm to look at the frontiers and generalize beyond

Alg. 1 is a procedure for drawing samples from the generative model (\mathcal{D}, θ) . In practice, we enumerate programs in order of their probability under Alg. 1 rather than sample them.

Why enumerate, when the program synthesis community has invented many sophisticated algorithms that search for programs? [10, 11, 12, 13, 14]. We have two reasons: (1) A key point of our work is that learning the DSL, along with a neural recognition model, can make program induction tractable, even if the search algorithm is very simple. (2) Enumeration is a general approach that can be applied to any program induction problem. Many of these more sophisticated approaches require special conditions on the space of programs.

A drawback of using an enumerative search algorithm is that we have no efficient means of solving for arbitrary constants that might occur in the program. In Sec. 3, we will show how to find programs with real-valued constants by automatically differentiating through the program and setting the constants using gradient descent.

2.3 Compilation: Learning a Neural Recognition Model

Example programs in frontiers	Proposed subexpression
<pre>(lambda (a b) (fold b (cons "," a) (lambda (x z) (cons x z)))) (lambda (a b) (fold a b (lambda (x z) (cons x z))))</pre>	<pre>(fold a b (lambda (x z) (cons x z)))</pre>

Figure 3: The DSL induction algorithm proposes subexpressions of programs to add to the DSL. These subexpressions are taken from programs in the frontiers (left column), and can introduce new variables (right column: a and b). Here, the proposed subexpression appends two lists.

177 them, both so the DSL can better express the current solutions, and also so that the DSL might expose
 178 new abstractions which will later be used to discover more programs. Formally, we want the DSL
 179 maximizing $\int \mathcal{L} d\theta$ (Sec. 2.1). We replace this marginal with an AIC approximation, giving the
 180 following objective for DSL induction:

$$\log \mathbb{P}[\mathcal{D}] + \arg \max_{\theta} \sum_{x \in X} \log \sum_{p \in \mathcal{F}_x} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] + \log \mathbb{P}[\theta|\mathcal{D}] - \|\theta\|_0 \quad (3)$$

181 We induce a DSL by searching locally through the space of DSLs, proposing small changes to \mathcal{D} until
 182 Eq. 3 fails to increase. The search moves work by introducing new λ -expressions into the DSL. We
 183 propose these new expressions by extracting subexpressions from programs already in the frontiers.
 184 These subexpressions are fragments of the original programs, and can introduce new variables
 185 (Fig. 3), which then become new functions in the DSL. The idea of storing and reusing fragments of
 186 expressions comes from Fragment Grammars [18] and Tree-Substitution Grammars [19].

187 We define a prior distribution over DSLs which penalizes the sizes of the λ -calculus expressions in
 188 the DSL, and put a Dirichlet prior over the weight vector:

$$\mathbb{P}[\mathcal{D}] \propto \exp \left(-\lambda \sum_{p \in \mathcal{D}} \text{size}(p) \right) \quad \mathbb{P}[\theta|\mathcal{D}] = \text{Dir}(\theta|\alpha) \quad (4)$$

189 where $\text{size}(p)$ measures the size of the syntax tree of program p , λ acts as a regularizer on the size of
 190 the DSL, and α is a concentration parameter controlling the smoothness of the prior over θ .

191 To appropriately score each proposed
 192 \mathcal{D} we must reestimate the weight vec-
 193 tor θ . Although this may seem very
 194 similar to estimating the parameters
 195 of a probabilistic context free gram-
 196 mar, for which we have effective ap-
 197 proaches like the Inside/Outside al-
 198 gorithm [20], our DSLs are context-
 199 sensitive due to the presence of vari-
 200 ables in the programs and also due
 201 to the polymorphic typing system. In
 202 the Supplement we derive a tractable
 203 MAP estimator for θ . Putting it all together, Alg. 2 describes how we combine program search,
 204 recognition model training, and DSL induction.

Algorithm 2 The ECC Algorithm

Input: Initial DSL \mathcal{D} , set of tasks X , iterations I
Hyperparameters: Enumeration timeout T
 Initialize $\theta \leftarrow \text{uniform}$
for $i = 1$ **to** I **do**
 $\mathcal{F}_x^\theta \leftarrow \{p|p \in \text{enum}(\mathcal{D}, \theta, T) \text{ if } \mathbb{P}[x|p] > 0\}$ (**Explore**)
 $q \leftarrow \text{train recognition model, maximizing } \mathcal{L}_{\text{RM}}$ (**Compile**)
 $\mathcal{F}_x^q \leftarrow \{p|p \in \text{enum}(\mathcal{D}, q(x), T) \text{ if } \mathbb{P}[x|p] > 0\}$ (**Explore**)
 $\mathcal{D}, \theta \leftarrow \text{induceDSL}(\{\mathcal{F}_x^\theta \cup \mathcal{F}_x^q\}_{x \in X})$ (**Compress**)
end for
return \mathcal{D}, θ, q

3 Symbolic Regression Experiments

205 As a simple first example, we apply ECC to symbolic regression
 206 problems. Here, the agent observes points along the curve of a func-
 207 tion, and must write a program that fits those points. We initially
 208 equip our learner with addition, multiplication, and division, and
 209 task it with solving 100 symbolic regression problems, each either
 210 a polynomial of degree 1–4 or a rational function. The recognition
 211 model is a convolutional network that observes an image of the target
 212 function’s graph (Fig. 4) – visually, different kinds of polynomials
 213




Figure 4: Recognition model input for symbolic regression. While the DSL learns subroutines for rational functions & polynomials, the recognition model jointly learns to look at a graph of the function (above) and predict which of those subroutines is appropriate for explaining the observation.

```

f0(a,b) = (fold a b (lambda (x y) (cons x y))))
(f0: Appends lists (of characters))
f1(s,c) = (fold s s (lambda (a x) (cdr (if (= c x) s a))))
(f1: Drop first characters from s until c reached)
f2(s) = (unfold s empty? car (lambda (z) (f1 z SPACE)))
(f2: Abbreviates a sequence of words)

```

Table 1: Some string editing learned subroutines

and rational functions produce different kinds of graphs, and so the recognition model can learn to look at a graph and predict what kind of function best explains it. A key difficulty, however, is that these problems are best solved with programs containing real numbers. Our solution to this difficulty is to allow the system to write programs with real-valued parameters, and then fit those parameters by automatically differentiating through the programs the system writes and use gradient descent to fit the parameters. We define the likelihood model, $\mathbb{P}[x|p]$, by assuming a Gaussian noise model for the input/output examples, and penalize the use of real-valued parameters using the BIC [21].

ECC learns a DSL containing templates for polynomials of different orders, as well as ratios of polynomials (Fig. 5). The algorithm also discovers programs that minimize the number of continuous degrees of freedom. For example, it learns to represent linear functions with the program `(* real (+ x real))`, which has two continuous degrees of freedom, and represents quartic functions using the invented DSL primitive f_6 in Tbl. 5 which has five continuous parameters. This phenomenon arises from our Bayesian framing – both the bias towards shorter programs and the likelihood model’s BIC penalty.

```

f0(x) = (* real (+ x real))
f1(x) = (f0 (* x (+ real x)))
f2(x) = (/ (/ x real) x)

```

4 Sequence manipulating programs

Figure 5: Some learned subroutines for symbolic regression. The system starts with addition, multiplication, division, and real numbers, and learns to build rational functions and polynomials up to 4th order.

We apply ECC to text editing (Section 4.1) and list processing (Section 4.2). For both these domains we use a bidirectional GRU [22] for the recognition model, and initially provide the system with a generic set of list processing primitives: `fold`, `unfold`, `if`, `map`, `length`, `index`, `=`, `+`, `-`, `0`, `1`, `cons`, `car`, `cdr`, `nil`, and `is-nil`.

4.1 String Editing

Synthesizing programs that manipulate strings is a classic problem in the programming languages and AI literatures [15, 23], and algorithms that learn string editing programs ship in Microsoft Excel [2]. This prior work presumes a ready-made DSL, expertly crafted to suit string editing. We show ECC can instead start out with generic Lisp primitives and recover many of the higher-level building blocks that have made these other system successful. We automatically generated 109 string editing tasks (Fig. 6) and model strings as lists of characters. At first, ECC cannot find any correct programs for most of the tasks. It assembles a DSL (Tbl. 1) that lets it rapidly explore the space of programs and find solutions to all of the tasks.

How well does the learned DSL generalized to real text-editing scenarios? We tested, but did not train, our system on problems from the SyGuS [24] program synthesis competition. Before any learning, ECC solves 32/108 of the problems with an average search time of 11 minutes. After learning, it solves 80/108, and does so much faster, solving them in an average of 29 seconds. As of the 2017 SyGuS competition, the best-performing algorithm solves 86/108 problems, and does so with a *different hand-engineered DSL for each problem*. Here we learned a single DSL that applied generically to all of the tasks, and perform comparably to the best prior work.

Input	Output	Input 1	Input 2	Output
Temple Annalisa Haven 185	TAH1	Launa	Withers	Launa Withers
Lara Gregori Bradford	LGB	Rudolf	Akiyama	Rudolf Akiyama
$f(s) = (f_2 \ s)$		$f(a, b) = (f_0 \ a \ (\text{cons } " \ " \ b))$		

Figure 6: Two string edit tasks (top) and the programs ECC writes for them (bottom). f_0 and f_2 are subroutines written by ECC, defined in Tbl. 1.

Name	Input	Output
append-4	[7 0 2]	[7 0 2 4]
len	[3 5 12 1]	4
has-2	[4 5 7 4]	false
repeat-2	[7 0]	[7 0 7 0]
drop-3	[0 3 8 6 4]	[6 4]
count-head-in-tail	[1 2 1 1 3]	2
rotate-2	[8 14 1 9]	[1 9 8 14]
pow-3	[10 4 13]	[1000 64 2197]
keep-mod-5	[5 9 14 6 3 0]	[5 0]

Table 2: Sample tasks from our list function domain

4.2 List Functions

Synthesizing programs that manipulate data structures is a widely studied problem in the programming languages community [12]. We consider this problem within the context of learning functions that manipulate lists. We created 225 Lisp-style list manipulation tasks, each with 15 input/output examples (Tbl. 2). Our data set is challenging along two dimensions: many of the functions are very complicated, and the agent must learn to solve these complicated problems from only 225 tasks. Our data set primarily consists of arithmetic operations upon sequences, and so, in addition to the Lisp primitives provided for the text editing experiments, we additionally start the system out with the following primitives: `mod`, `*`, `>`, `is-square`, `is-prime`.

A complete repertoire of higher-order functions is a staple of functional programming standard libraries. Although we provided ECC with some of the standard higher-order functions, like `fold` and `unfold`, we did not include others. When trained on these list functions, our system rediscovers and then reuses the higher-order function `filter` **Lucas: not sure if it actually does this! It would be cool if we can write something like this though.**

4.3 Quantitative Results

We compare with four baselines on held-out tasks:

Ours (no NN), which lesions the recognition model.

RF/DC, which holds the generative model (\mathcal{D}, θ) fixed and learns a recognition model only from samples from the fixed generative model. This is equivalent to our algorithm with $\lambda = \infty$ (Sec. 2.4) and $\mathcal{L}_{\text{RM}} = \mathcal{L}_{\text{HM}}$ (Sec. 2.3). We call this baseline RF/DC because this setup is closest to how RobustFill [3] and DeepCoder [5] are trained. We can not compare directly with these systems, because they are engineered for one specific domain, and do not have publicly available code and datasets.

PCFG, which lesions the recognition model, learns θ , and fixes \mathcal{D} . This is equivalent to our algorithm with $q(x) = \theta$ and $\lambda = \infty$, and is like learning the parameters of a PCFG while not learning any of the structure.

Enum, which does no learning and just enumerates a frontier. This is equivalent to our first wake cycle.

For each domain, we are interested both in how many tasks the agent can solve and how quickly it can find those solutions. For symbolic regression, we also care about the quality of the solution, as measured by the likelihood model $\mathbb{P}[x|p]$, e.g. did the agent correctly explain a linear function

```

f0(i, ℓ) = (singleton (index i ℓ))
(f0: Put the i-th index into a list)
f1(i, ℓ) = (++ ℓ (f0 i ℓ))
(f1: Append the i-th index)
f2(n, ℓ) = (any (lambda (x) (= n x)) ℓ)
(f2: Whether n appears in the list)
f3(i, ℓ) = (index (negate i) (sort ℓ))
(f3: Get the i-th largest number)
f4(n, ℓ) = (mapi (lambda (i x) (mod x n)) ℓ)
f5(k, ℓ) = (mapi (lambda (i x) (+ x k)) ℓ)
f6(k, n, ℓ) = (f4 n (f5 k ℓ))
(f6: Caesar shift of k in integers modulo n)

```

Table 3: Some learned DSL primitives for list functions

	Ours	Ours (no NN)	RF/DC	PCFG	Enum
<i>Symbolic Regression</i>					
% solved	98%	94%	0.7%	0%	2%
Solve time	2.7s	2.8s	3.6s	—	2.1s
MDL (nats)	9.3	11.0	2.3	—	4.6
<i>String Editing</i>					
% solved	99%	82%	9%	24%	18%
Solve time	0.9s	2.2s	2.7s	2.3s	2.4s
<i>List functions</i>					
% solved	86%	84%	60%	74%	70%
Solve time	0.8s	0.7s	1.0s	1.0s	1.1s

Table 4: % solved w/ 5 sec timeout. Solve time: averaged over solved tasks. RF/DC: trained like RobustFill/DeepCoder. PCFG: model w/o structure learning. Enum: model w/o any learning. MDL: $-\mathbb{E} [\mathbb{P}[x|p]]$. For domains other than symbolic regression MDL is 0 nats. For symbolic regression MDL is a proxy for # of continuous parameters.

287 using two numbers, or did it introduce extraneous parameters? Tbl. 4 compares our model against
288 these baselines. Our full model consistently improves on the baselines, sometimes dramatically
289 (string editing and symbolic regression). The recognition model consistently increases the number of
290 solved held-out tasks, and lesioning it also slows down the convergence of the algorithm, taking more
291 wake/sleep cycles to reach a given number of tasks solved (Fig. 7). This supports the view of the
292 recognition model as a way of accelerating the search over programs.

293 5 Related Work

294 Our work is far from the first for learning to learn programs, an idea that goes back to Solomonoff [25]:




Figure 7: Learning curves for ECC both with (blue) and without (red) the recognition model as the frontier size is varied (solid/dashed/dotted lines).

295 *Deep learning*: Much recent work in the ML community has focused on creating neural networks that
 296 regress from input/output examples to programs [3, 26, 15, 5]. These neural networks are typically
 297 trained with strong supervision (i.e., with annotated ground-truth programs) on massive data sets
 298 (i.e., hundreds of millions [3]). Our work considers a weakly-supervised regime where ground truth
 299 programs are not provided and the agent must learn from a few hundred tasks.

300 *Inventing new subroutines for program induction*: Several program induction algorithms, most
 301 prominently the EC algorithm [8], take as their goal to learn new, reusable subroutines that are
 302 shared in a multitask setting. We find this work inspiring and motivating, and extend it along two
 303 dimensions: (1) we propose a new algorithm for inducing reusable subroutines, based on Fragment
 304 Grammars [18]; and (2) we show how to combine these techniques with bottom-up neural recognition
 305 models. Other instances of this related idea are [27], Schmidhuber’s OOPS model [28], and predicate
 306 invention in ILP [29].

307 Our work is an instance of Bayesian Program Learning (BPL; see [30, 8, 31, 27]). Previous BPL
 308 systems have largely assumed a fixed DSL (but see [27]), and our contribution here is a general way
 309 of doing BPL with less hand-engineering of the DSL.

310 6 Contribution and Outlook

311 We contribute an algorithm, ECC, that learns to program by bootstrapping a DSL with new domain-
 312 specific primitives that the algorithm itself discovers, together with a neural recognition model
 313 that learns how to efficiently deploy the DSL on new tasks. We believe this integration of top-
 314 down symbolic representations and bottom-up neural networks – both of them learned – could help
 315 make program induction systems more generally useful for AI. Many directions remain open. Two
 316 immediate goals are to integrate more sophisticated neural recognition models [3] and program
 317 synthesizers [10], which may improve performance in some domains over the generic methods used
 318 here. Another direction is to explore DSL meta-learning: Can we find a *single* universal primitive set
 319 that could effectively bootstrap DSLs for new domains, including the four domains considered, but
 320 also many others?

321 References

- 322 [1] Stephen H Muggleton, Dianhuan Lin, and Alireza Tamaddoni-Nezhad. Meta-interpretive learning of
 323 higher-order dyadic datalog: Predicate invention revisited. *Machine Learning*, 100(1):49–73, 2015.
- 324 [2] Sumit Gulwani. Automating string processing in spreadsheets using input-output examples. In *ACM*
 325 *SIGPLAN Notices*, volume 46, pages 317–330. ACM, 2011.
- 326 [3] Jacob Devlin, Jonathan Uesato, Surya Bhupatiraju, Rishabh Singh, Abdel-rahman Mohamed, and Pushmeet
 327 Kohli. Robustfill: Neural program learning under noisy i/o. *arXiv preprint arXiv:1703.07469*, 2017.
- 328 [4] Ashwin Kalyan, Abhishek Mohta, Oleksandr Polozov, Dhruv Batra, Prateek Jain, and Sumit Gulwani.
 329 Neural-guided deductive search for real-time program synthesis from examples. *ICLR*, 2018.
- 330 [5] Matej Balog, Alexander L Gaunt, Marc Brockschmidt, Sebastian Nowozin, and Daniel Tarlow. Deepcoder:
 331 Learning to write programs. *ICLR*, 2016.
- 332 [6] Tuan Anh Le, Atılım Güneş Baydin, and Frank Wood. Inference Compilation and Universal Probabilistic
 333 Programming. In *AISTATS*, 2017.
- 334 [7] Geoffrey E Hinton, Peter Dayan, Brendan J Frey, and Radford M Neal. The "wake-sleep" algorithm for
 335 unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.
- 336 [8] Eyal Dechter, Jon Malmaud, Ryan P. Adams, and Joshua B. Tenenbaum. Bootstrap learning via modular
 337 concept discovery. In *IJCAI*, 2013.
- 338 [9] Benjamin C. Pierce. *Types and programming languages*. MIT Press, 2002.
- 339 [10] Armando Solar Lezama. *Program Synthesis By Sketching*. PhD thesis, 2008.
- 340 [11] Eric Schkufza, Rahul Sharma, and Alex Aiken. Stochastic superoptimization. In *ACM SIGARCH Computer*
 341 *Architecture News*, volume 41, pages 305–316. ACM, 2013.

- 342 [12] John K Feser, Swarat Chaudhuri, and Isil Dillig. Synthesizing data structure transformations from input-
343 output examples. In *PLDI*, 2015.
- 344 [13] Peter-Michael Osera and Steve Zdancewic. Type-and-example-directed program synthesis. In *ACM*
345 *SIGPLAN Notices*, volume 50, pages 619–630. ACM, 2015.
- 346 [14] Oleksandr Polozov and Sumit Gulwani. Flashmeta: A framework for inductive program synthesis. *ACM*
347 *SIGPLAN Notices*, 50(10):107–126, 2015.
- 348 [15] Aditya Menon, Omer Tamuz, Sumit Gulwani, Butler Lampson, and Adam Kalai. A machine learning
349 framework for programming by example. In *ICML*, pages 187–195, 2013.
- 350 [16] Geoffrey E Hinton and Ruslan R Salakhutdinov. Reducing the dimensionality of data with neural networks.
351 *Science*, 2006.
- 352 [17] Peter Dayan, Geoffrey E Hinton, Radford M Neal, and Richard S Zemel. The helmholtz machine. *Neural*
353 *computation*, 7(5):889–904, 1995.
- 354 [18] Timothy J. O’Donnell. *Productivity and Reuse in Language: A Theory of Linguistic Computation and*
355 *Storage*. The MIT Press, 2015.
- 356 [19] Trevor Cohn, Phil Blunsom, and Sharon Goldwater. Inducing tree-substitution grammars. *Journal of*
357 *Machine Learning Research*, 11(Nov):3053–3096, 2010.
- 358 [20] J.D. Lafferty. *A Derivation of the Inside-outside Algorithm from the EM Algorithm*. Research report.
- 359 [21] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*.
360 2006.
- 361 [22] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger
362 Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical
363 machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- 364 [23] Tessa Lau. *Programming by demonstration: a machine learning approach*. PhD thesis, University of
365 Washington, 2001.
- 366 [24] Rajeev Alur, Dana Fisman, Rishabh Singh, and Armando Solar-Lezama. Sygus-comp 2016: results and
367 analysis. *arXiv preprint arXiv:1611.07627*, 2016.
- 368 [25] Ray J Solomonoff. A system for incremental learning based on algorithmic probability. Sixth Israeli
369 Conference on Artificial Intelligence, Computer Vision and Pattern Recognition, 1989.
- 370 [26] Jacob Devlin, Rudy R Bunel, Rishabh Singh, Matthew Hausknecht, and Pushmeet Kohli. Neural program
371 meta-induction. In *NIPS*, 2017.
- 372 [27] Percy Liang, Michael I. Jordan, and Dan Klein. Learning programs: A hierarchical bayesian approach. In
373 *ICML*, 2010.
- 374 [28] Jürgen Schmidhuber. Optimal ordered problem solver. *Machine Learning*, 54(3):211–254, 2004.
- 375 [29] Dianhuan Lin, Eyal Dechter, Kevin Ellis, Joshua B. Tenenbaum, and Stephen Muggleton. Bias reformulation
376 for one-shot function induction. In *ECAI 2014*, 2014.
- 377 [30] Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. Human-level concept learning through
378 probabilistic program induction. *Science*, 350(6266):1332–1338, 2015.
- 379 [31] Kevin Ellis, Armando Solar-Lezama, and Josh Tenenbaum. Sampling for bayesian program learning. In
380 *Advances in Neural Information Processing Systems*, 2016.