
Learning Libraries of Subroutines for Neurally-Guided Bayesian Program Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Successful approaches to program induction require a hand-engineered domain-
2 specific language (DSL), constraining the space of allowed programs and imparting
3 prior knowledge of the domain. We contribute a program induction algorithm
4 called ECC that learns a DSL while jointly training a neural network to efficiently
5 search for programs in the learned DSL. We use our model to solve symbolic
6 regression problems, edit strings, and synthesize functions on lists, showing how
7 the model learns a domain-specific library of program components for expressing
8 solutions to problems in the domain.

9 1 Introduction

10 Automatically inducing programs from examples is a long-standing goal of artificial intelligence.
11 Recent work has successfully used symbolic search techniques (e.g., Metagol: [1], FlashFill: [2]),
12 neural networks trained from a corpus of examples (e.g., RobustFill: [3]), and hybrids of neural and
13 symbolic methods (e.g., Neural-guided deductive search: [4], DeepCoder: [5]) to synthesize programs
14 for task domains such as string transformations, list processing, and robot navigation and planning.
15 However, all these approaches – symbolic, neural and neural-symbolic – rely upon a hand-engineered
16 *Domain-Specific Language* (DSL). DSLs contain an inventory of restricted programming primitives,
17 encoding domain-specific knowledge about the space of programs. In practice we often have only a
18 few input/output examples for each program to be induced, and thus success often hinges on having a
19 good DSL that provides a crucial inductive bias for what would otherwise be an unconstrained search
20 through the space of all computable functions. Here we ask, to what extent can we dispense with
21 such highly hand-engineered domain-specific languages?

22 We propose *learning* the DSL by inducing a library of domain-specific subroutines. We consider
23 the setting where we have a collection of related programming tasks, each specified by a set of
24 input/output examples. Starting from a weaker or more general library of primitives, we give an
25 algorithm for constructing a richer, more powerful, and better-tuned DSL. Our algorithm is called
26 **Explore/Compress/Compile** (ECC), and iterates between three different steps: an **Explore** step
27 uses the DSL to explore the space of programs, searching for ones that solve the tasks; a **Compress**
28 step modifies the DSL by discovering regularities in the programs found by the previous Explore step;
29 and a **Compile** step, which improves the program search procedure by training a neural network to
30 write programs in the current DSL, in the spirit of “amortized” or “compiled” inference [6]. We call
31 the neural net a **recognition model** (c.f. Hinton 1995 [7]). The learned DSL distills commonalities
32 across programs that solve tasks, helping the agent solve related program induction problems. The
33 neural recognition model ensures that searching for programs remains tractable even as the DSL (and
34 hence the search space for programs) expands.

35 Because any model may be encoded as a (deterministic or probabilistic) program, we carefully
36 delineate the scope of program learning problems considered here. We think of ECC as learning to

37 solve the kinds of problems that humans can solve relatively quickly – once they acquire the relevant
 38 domain expertise. These correspond to short programs – once you have the right DSL. Even with the
 39 right DSL, program search may be intractable, so we amortize the cost of program search by training
 40 a neural network to assist the search procedure.

41 We apply ECC to three domains: symbolic regression; FlashFill-style [2] string pro-
 42 cessing; and Lisp-style functions on lists. For each of these we initially provide a
 43 generic set of programming primitives in a Lisp-like language. Our algorithm then dis-
 44 covers its own domain-specific vocabulary for expressing solutions in the domain (Tbl. 1).
 45

46 2 The ECC Algorithm

47 Our goal is to induce a DSL while
 48 finding programs solving each of the
 49 tasks. We take inspiration primar-
 50 ily from the Exploration-Compression
 51 algorithm for bootstrap learning [8].
 52 Exploration-Compression alternates
 53 between exploring the space of solu-
 54 tions to a set of tasks, and compress-
 55 ing those solutions to suggest new
 56 search primitives for the next explo-
 57 ration stage. We extend these ideas into an inference strategy that iterates through three steps: an
 58 **Explore** step uses the current DSL and recognition model to search for programs that solve the tasks.
 59 The **Compress** and **Compile** steps update the DSL and the recognition model, respectively. Crucially,
 60 these steps bootstrap off each other (Fig. 2):

61 **Exploration: Searching for programs.** Our
 62 program search is informed by both the DSL
 63 and the recognition model. When these improve,
 64 we find more programs solving the tasks.

65 **Compression: Improving the DSL.** We in-
 66 duce the DSL from the programs found in the
 67 exploration phase, aiming to maximally compress
 68 (or, raise the prior probability of) these
 69 programs. As we solve more tasks, we hone in
 70 on DSLs that more closely match the domain.

71 **Compilation: Learning a neural recognition**
 72 **model.** We update the recognition model by
 73 training on two data sources: samples from the
 74 DSL (as in the Helmholtz Machine’s “sleep”
 75 phase), and programs found by the search proce-
 76 dure during exploration. As the DSL improves
 77 and as search finds more programs, the recogni-
 78 tion model gets more data to train on, and better data.

Domain	Part of the learned DSL
Regression	<code>(+ (* real x) real)</code> <i>(a linear function of x)</i>
Strings	<code>(map (lambda (x) (if (= x a) b x)))</code> <i>(replace occurrences of a w/ b)</i>
Lists	<code>(fold nil (lambda (a b) (cons a b)))</code> <i>(reverses a list)</i>

Figure 1: Examples of structure found in DSLs learned by our algorithm. ECC builds a new DSL by discovering and reusing useful subroutines.




Figure 2: ECC solves for programs, the DSL, and a neural network (recognition model). Each of these steps iteratively bootstraps off of the others.

79 2.1 Hierarchical Bayesian Framing

80 ECC takes as input a set of *tasks*, written X , each of which is a program induction problem. It has at
 81 its disposal a *likelihood model*, written $\mathbb{P}[x|p]$, which scores the likelihood of a task $x \in X$ given a
 82 program p . Its goal is to solve each of the tasks by writing a program, and also to infer a DSL, written
 83 \mathcal{D} . We equip \mathcal{D} with a real-valued weight vector θ , and together (\mathcal{D}, θ) define a generative model
 84 over programs. We frame our goal as maximum a posteriori (MAP) inference of (\mathcal{D}, θ) given X .

85 Writing J for the joint probability of (\mathcal{D}, θ) and X , we want the \mathcal{D}^* and θ^* solving:

$$J(\mathcal{D}, \theta) \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_p \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \int J(\mathcal{D}, \theta) d\theta \quad \theta^* = \arg \max_{\theta} J(\mathcal{D}^*, \theta) \quad (1)$$

86 The above equations summarize the problem from the point of view of an ideal Bayesian learner.
 87 However, Eq. 1 is wildly intractable because evaluating $J(\mathcal{D}, \theta)$ involves summing over the infinite
 88 set of all programs. In practice we will only ever be able to sum over a finite set of programs. So, for
 89 each task, we define a finite set of programs, called a *frontier*, and only marginalize over the frontiers:
 90 **Definition.** A *frontier of task x* , written \mathcal{F}_x , is a finite set of programs s.t. $\mathbb{P}[x|p] > 0$ for all $p \in \mathcal{F}_x$.

91 Using the frontiers we define the following intuitive lower bound on the joint probability, called \mathcal{L} :

$$J \geq \mathcal{L} \triangleq \mathbb{P}[\mathcal{D}, \theta] \prod_{x \in X} \sum_{p \in \mathcal{F}_x} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] \quad (2)$$

92 ECC does approximate MAP inference by maximizing this lower bound on the joint probability,
 93 alternating maximization w.r.t. the frontiers (Exploration) and the DSL (Compression):

94 **Program Search: Maxing \mathcal{L} w.r.t. the frontiers.** Here (\mathcal{D}, θ) is fixed and we want to find new
 95 programs to add to the frontiers so that \mathcal{L} increases the most. \mathcal{L} most increases by finding programs
 96 where $\mathbb{P}[x, p|\mathcal{D}, \theta]$ is large.

97 **DSL Induction: Maxing $\int \mathcal{L} d\theta$ w.r.t. the DSL.** Here $\{\mathcal{F}_x\}_{x \in X}$ is held fixed, and so we can
 98 evaluate \mathcal{L} . Now the problem is that of searching the discrete space of DSLs and finding one
 99 maximizing $\int \mathcal{L} d\theta$. Once we have a DSL \mathcal{D} we can update θ to $\arg \max_{\theta} \mathcal{L}(\mathcal{D}, \theta, \{\mathcal{F}_x\})$.

100 Searching for programs is hard because of the large combinatorial search space. We ease this
 101 difficulty by training a neural recognition model, $q(\cdot|\cdot)$, during the compilation phase: q is trained to
 102 approximate the posterior over programs, $q(p|x) \propto \mathbb{P}[p|x, \mathcal{D}, \theta] = \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta]$, thus amortizing
 103 the cost of finding programs with high posterior probability.

104 **Neural recognition model: tractably maxing \mathcal{L} w.r.t. the frontiers.** Here we train a neural
 105 network, q , to predict a distribution over programs conditioned on a task. The objective of q is to
 106 assign high probability to programs p where $\mathbb{P}[x, p|\mathcal{D}, \theta]$ is large, because including those programs
 107 in the frontiers will most increase \mathcal{L} .

108 2.2 Exploration: Searching for Programs

109 Now our goal is to search for programs solving the tasks. We use the simple approach of enumerating
 110 programs from the DSL in decreasing order of their probability, and then checking if a program p
 111 assigns positive probability to a task ($\mathbb{P}[x|p] > 0$); if so, we incorporate p into the frontier \mathcal{F}_x .

112 To make this concrete we need to define what programs actually are and what form $\mathbb{P}[p|\mathcal{D}, \theta]$ takes.
 113 We represent programs as λ -calculus expressions. λ -calculus is a formalism for expressing functional
 114 programs that closely resembles the Lisp programming language. λ -calculus includes variables,
 115 function application, and the ability to create new functions. Throughout this paper we will write λ -
 116 calculus expressions in Lisp syntax. Our programs are all strongly typed. We use the Hindley-Milner
 117 polymorphic typing system [9] which is used in functional programming languages like OCaml. Type
 118 variables are always written using lowercase Greek letters and we write $\alpha \rightarrow \beta$ to mean a function
 119 that takes an input of type α and returns something of type β . We use the notation $p : \tau$ to mean that
 120 the λ -calculus expression p has the type τ . For example, to describe the type of the identity function
 121 we would say $(\text{lambda } (x) x) : \alpha \rightarrow \alpha$. We say a type α *unifies* with τ if every expression $p : \alpha$
 122 also satisfies $p : \tau$. Furthermore, the act of *unifying* a type α with τ is to introduce constraints on
 123 the type variables of α to ensure that α unifies with τ . See Supplement for more detail on program
 124 representation. With this notation in hand we now define DSLs:

125 **Definition:** (\mathcal{D}, θ) . A DSL \mathcal{D} is a set of typed λ -calculus expressions. A weight vector θ for a DSL
 126 \mathcal{D} is a vector of $|\mathcal{D}| + 1$ real numbers: one number for each DSL primitive $e \in \mathcal{D}$, written θ_e , and a
 127 weight controlling the probability of a variable occurring in a program, θ_{var} .

Algorithm 1 Generative model over programs

```

128 function sample( $\mathcal{D}, \theta, \mathcal{E}, \tau$ ):
129   Input: DSL ( $\mathcal{D}, \theta$ ), environment  $\mathcal{E}$ , type  $\tau$ 
130   Output: a program whose type unifies with  $\tau$ 
131   if  $\tau = \alpha \rightarrow \beta$  then
132     var  $\leftarrow$  an unused variable name
133     body  $\sim$  sample( $\mathcal{D}, \theta, \{\text{var} : \alpha\} \cup \mathcal{E}, \beta$ )
134     return (lambda (var) body)
135   end if
136   primitives  $\leftarrow \{p \mid p : \tau' \in \mathcal{D} \cup \mathcal{E}$ 
137     if  $\tau$  can unify with  $\text{yield}(\tau')$ 
138   Draw  $e \sim$  primitives, w.p.  $\propto \theta_e$  if  $e \in \mathcal{D}$ 
139     w.p.  $\propto \frac{\theta_{\text{var}}}{|\text{variables}|}$  if  $e \in \mathcal{E}$ 
140   Unify  $\tau$  with  $\text{yield}(\tau')$ .
141    $\{\alpha_k\}_{k=1}^K \leftarrow \text{args}(\tau')$ 
142   for  $k = 1$  to  $K$  do
143      $a_k \sim$  sample( $\mathcal{D}, \theta, \mathcal{E}, \alpha_k$ )
144   end for
145   return ( $e \ a_1 \ a_2 \ \dots \ a_K$ )
146   where:
147   yield( $\tau$ ) =  $\begin{cases} \text{yield}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ \tau & \text{otherwise.} \end{cases}$ 
148   args( $\tau$ ) =  $\begin{cases} [\alpha] + \text{args}(\beta) & \text{if } \tau = \alpha \rightarrow \beta \\ [] & \text{otherwise.} \end{cases}$ 

```

The purpose of the recognition model is to amortize the cost of searching for programs. It does this by learning to predict programs which are probable under (\mathcal{D}, θ) while also assigning high likelihood to a task according to $\mathbb{P}[x|p]$. Concretely, the recognition model q is a neural network that predicts, for each task $x \in X$, a weight vector $q(x) = \theta^{(x)} \in \mathbb{R}^{|\mathcal{D}|+1}$. Together with the DSL, this defines a distribution over programs, $\mathbb{P}[p|\mathcal{D}, \theta = q(x)]$. We abbreviate this distribution as $q(p|x)$. The crucial aspect of this framing is that the neural network leverages the structure of the learned DSL, so it is *not* responsible for generating programs wholesale. We share this aspect with DeepCoder [5] and [15].

We want a recognition model that closely approximates the true posteriors over programs. We formulate this as minimizing the expected KL-divergence, $\mathbb{E}[\text{KL}(\mathbb{P}[p|x, \mathcal{D}, \theta] \| q(p|x))]$, or equivalently maximizing $\mathbb{E}[\sum_p \mathbb{P}[p|x, \mathcal{D}, \theta] \log q(p|x)]$, where the expectation is taken over tasks. One could take this expectation over the empirical distribution of tasks, like how an autoencoder is trained [16]; or, one could take this expectation over samples from the generative model, like how a Helmholtz machine is trained [17]. We found it useful to maximize both an autoencoder-style objective \mathcal{L}_{AE} and a Helmholtz-style objective \mathcal{L}_{HM} , giving the objective for a recognition model, $\mathcal{L}_{\text{RM}} = \mathcal{L}_{\text{AE}} + \mathcal{L}_{\text{HM}}$:

$$\mathcal{L}_{\text{HM}} = \mathbb{E}_{(p,x) \sim (\mathcal{D}, \theta)} [\log q(p|x)] \quad \mathcal{L}_{\text{AE}} = \mathbb{E}_{x \sim X} \left[\sum_{p \in \mathcal{F}_x} \frac{\mathbb{P}[x, p|\mathcal{D}, \theta]}{\sum_{p' \in \mathcal{F}_x} \mathbb{P}[x, p'|\mathcal{D}, \theta]} \log q(p|x) \right]$$

The \mathcal{L}_{HM} objective is essential for data efficiency: all of our experiments train ECC on only a few hundred tasks, which is too little for a high-capacity neural network q . Once we bootstrap a (\mathcal{D}, θ) , we can draw unlimited samples from (\mathcal{D}, θ) and train q on those samples.

Evaluating \mathcal{L}_{HM} involves sampling programs from the current DSL, running them to get their outputs, and then training q to regress from the input/outputs to the program. Since these programs map inputs to outputs, we need to sample the inputs as well. Our solution is to sample the inputs from the empirical observed distribution of inputs in X .

2.4 Compression: Learning a Generative Model (a DSL)

The purpose of the DSL is to offer a set of abstractions that allow an agent to easily express solutions to the tasks at hand. Intuitively, we want the algorithm to look at the frontiers and generalize beyond

Alg. 1 is a procedure for drawing samples from the generative model (\mathcal{D}, θ) . In practice, we enumerate programs in order of their probability under Alg. 1 rather than sample them.

Why enumerate, when the program synthesis community has invented many sophisticated algorithms that search for programs? [10, 11, 12, 13, 14]. We have two reasons: (1) A key point of our work is that learning the DSL, along with a neural recognition model, can make program induction tractable, even if the search algorithm is very simple. (2) Enumeration is a general approach that can be applied to any program induction problem. Many of these more sophisticated approaches require special conditions on the space of programs.

A drawback of using an enumerative search algorithm is that we have no efficient means of solving for arbitrary constants that might occur in the program. In Sec. 4, we will show how to find programs with real-valued constants by automatically differentiating through the program and setting the constants using gradient descent.

2.3 Compilation: Learning a Neural Recognition Model

Example programs in frontiers	Proposed subexpression
<pre>(lambda (a b) (fold b (cons ", " a) (lambda (x z) (cons x z)))) (lambda (a b) (fold a b (lambda (x z) (cons x z))))</pre>	<pre>(fold a b (lambda (x z) (cons x z)))</pre>

Figure 3: The DSL induction algorithm proposes subexpressions of programs to add to the DSL. These subexpressions are taken from programs in the frontiers (left column), and can introduce new variables (right column: a and b). Here, the proposed subexpression appends two lists.

177 them, both so the DSL can better express the current solutions, and also so that the DSL might expose
 178 new abstractions which will later be used to discover more programs. Formally, we want the DSL
 179 maximizing $\int \mathcal{L} d\theta$ (Sec. 2.1). We replace this marginal with an AIC approximation, giving the
 180 following objective for DSL induction:

$$\log \mathbb{P}[\mathcal{D}] + \arg \max_{\theta} \sum_{x \in X} \log \sum_{p \in \mathcal{F}_x} \mathbb{P}[x|p] \mathbb{P}[p|\mathcal{D}, \theta] + \log \mathbb{P}[\theta|\mathcal{D}] - \|\theta\|_0 \quad (3)$$

181 We induce a DSL by searching locally through the space of DSLs, proposing small changes to \mathcal{D} until
 182 Eq. 3 fails to increase. The search moves work by introducing new λ -expressions into the DSL. We
 183 propose these new expressions by extracting subexpressions from programs already in the frontiers.
 184 These subexpressions are fragments of the original programs, and can introduce new variables
 185 (Fig. 3), which then become new functions in the DSL. The idea of storing and reusing fragments of
 186 expressions comes from Fragment Grammars [18] and Tree-Substitution Grammars [19].

187 To define the prior distribution over (\mathcal{D}, θ) , we penalize the syntactic complexity of the λ -calculus
 188 expressions in the DSL, defining $\mathbb{P}[\mathcal{D}] \propto \exp\left(-\lambda \sum_{p \in \mathcal{D}} \text{size}(p)\right)$ where $\text{size}(p)$ measures the size
 189 of the syntax tree of program p , and place a symmetric Dirichlet prior over the weight vector θ .

190 To appropriately score each proposed
 191 \mathcal{D} we must reestimate the weight vec-
 192 tor θ . Although this may seem very
 193 similar to estimating the parameters
 194 of a probabilistic context free gram-
 195 mar, for which we have effective ap-
 196 proaches like the Inside/Outside al-
 197 gorithm [20], our DSLs are context-
 198 sensitive due to the presence of vari-
 199 ables in the programs and also due
 200 to the polymorphic typing system. In
 201 the Supplement we derive a tractable
 202 MAP estimator for θ . Putting all these ingredients together, Alg. 2 describes how we combine
 203 program search, recognition model training, and DSL induction.

Algorithm 2 The ECC Algorithm

Input: Initial DSL \mathcal{D} , set of tasks X , iterations I
Hyperparameters: Enumeration timeout T
 Initialize $\theta \leftarrow$ uniform
for $i = 1$ **to** I **do**
 $\mathcal{F}_x^\theta \leftarrow \{p | p \in \text{enum}(\mathcal{D}, \theta, T) \text{ if } \mathbb{P}[x|p] > 0\}$ (**Explore**)
 $q \leftarrow$ train recognition model, maximizing \mathcal{L}_{RM} (**Compile**)
 $\mathcal{F}_x^q \leftarrow \{p | p \in \text{enum}(\mathcal{D}, q(x), T) \text{ if } \mathbb{P}[x|p] > 0\}$ (**Explore**)
 $\mathcal{D}, \theta \leftarrow \text{induceDSL}(\{\mathcal{F}_x^\theta \cup \mathcal{F}_x^q\}_{x \in X})$ (**Compress**)
end for
return \mathcal{D}, θ, q

204 3 Sequence manipulating programs

205 We apply ECC to list processing (Section 3.2) and text editing (Section 3.1). For both these domains
 206 we use a bidirectional GRU [21] for the recognition model, and initially provide the system with a
 207 generic set of list processing primitives: fold, unfold, if, map, length, index, =, +, -, 0, 1, cons, car,
 208 cdr, nil, and is-nil.

209 3.1 List Functions

210 Synthesizing programs that manipulate data structures is a widely studied problem in the programming
 211 languages community [12]. We consider this problem within the context of learning functions that
 212 manipulate lists. We created 225 Lisp-style list manipulation tasks, each with 15 input/output
 213 examples (Tbl. ??). Our data set is challenging along two dimensions: many of the functions are very
 214 complicated, and the agent must learn to solve these complicated problems from only 225 tasks. Our
 215 data set primarily consists of arithmetic operations upon sequences, and so, in addition to the Lisp

```

f0(a,b) = (fold a b (lambda (x y) (cons x y))))
(f0: Appends lists (of characters))
f1(s,c) = (fold s s (lambda (a x) (cdr (if (= c x) s a))))
(f1: Drop first characters from s until c reached)
f2(s) = (unfold s empty? car (lambda (z) (f1 z SPACE)))
(f2: Abbreviates a sequence of words)

```

Table 1: Some string editing learned subroutines

Input	Output	Input 1	Input 2	Output
Temple Annalisa Haven 185	TAH1	Launa	Withers	Launa Withers
Lara Gregori Bradford	LGB	Rudolf	Akiyama	Rudolf Akiyama
$f(s) = (f_2 \ s)$		$f(a,b) = (f_0 \ a \ (cons \ " \ " \ b))$		

Figure 4: Two string edit tasks (top) and the programs ECC writes for them (bottom). f_0 and f_2 are subroutines written by ECC, defined in Tbl. 1.

primitives provided for the text editing experiments, we additionally start the system out with the following primitives: `mod`, `*`, `>`, `is-square`, `is-prime`.

A complete repertoire of higher-order functions is a staple of functional programming standard libraries. Although we provided ECC with some of the standard higher-order functions, like `fold` and `unfold`, we did not include others. When trained on these list functions, our system rediscovers and then reuses the higher-order function `filter` **Lucas: not sure if it actually does this! It would be cool if we can write something like this though.**

3.2 String Editing

Synthesizing programs that manipulate strings is a classic problem in the programming languages and AI literatures [15, 22], and algorithms that learn string editing programs ship in Microsoft Excel [2]. This prior work presumes a ready-made DSL, expertly crafted to suit string editing. We show ECC can instead start out with generic Lisp primitives and recover many of the higher-level building blocks that have made these other system successful. We automatically generated 109 string editing tasks (Fig. 4) and model strings as lists of characters. At first, ECC cannot find any correct programs for most of the tasks. It assembles a DSL (Tbl. 1) that lets it rapidly explore the space of programs and find solutions to all of the tasks.

How well does the learned DSL generalized to real text-editing scenarios? We tested, but did not train, our system on problems from the SyGuS [23] program synthesis competition. Before any learning, ECC solves 32/108 of the problems with an average search time of 11 minutes. After learning, it solves 80/108, and does so much faster, solving them in an average of 29 seconds. As of the 2017 SyGuS competition, the best-performing algorithm solves 86/108 problems, and does so with a *different hand-engineered DSL for each problem*. Here we learned a single DSL that applied generically to all of the tasks, and perform comparably to the best prior work.

4 Symbolic Regression Experiments

We apply ECC to symbolic regression problems. Here, the agent observes points along the curve of a function, and must write a program that fits those points. We initially equip our learner with addition, multiplication, and division, and task it with solving 100 symbolic regression problems, each either a polynomial of degree 1–4 or a rational function. The recognition model is a convolutional network that observes an image of the target function’s graph (Fig. 5) – visually, different kinds of polynomials and rational functions produce different kinds of graphs, and so the recognition model can learn to look at a graph and predict what kind of function best explains it. A



Figure 5: Recognition model input for symbolic regression. While the DSL learns subroutines for rational functions & polynomials, the recognition model jointly learns to look at a graph of the function (above) and predict which of those subroutines is appropriate for explaining the observation.




Figure 7: Learning curves for ECC both with (blue) and without (red) the recognition model as the frontier size is varied (solid/dashed/dotted lines).

key difficulty, however, is that these problems are best solved with programs containing real numbers. Our solution to this difficulty is to allow the system to write programs with real-valued parameters, and then fit those parameters by automatically differentiating through the programs the system writes and use gradient descent to fit the parameters. We define the likelihood model, $P[x|p]$, by assuming a Gaussian noise model for the input/output examples, and penalize the use of real-valued parameters using the BIC [24].

ECC learns a DSL containing templates for polynomials of different orders, as well as ratios of polynomials (Fig. 6). The algorithm also discovers programs that minimize the number of continuous degrees of freedom. For example, it learns to represent linear functions with the program `(* real (+ x real))`, which has two continuous degrees of freedom, and represents quartic functions using the invented DSL primitive f_6 in Tbl. 6 which has five continuous parameters. This phenomenon arises from our Bayesian framing – both the bias towards shorter programs and the likelihood model’s BIC penalty.

```

f0(x) = (* real (+ x real))
f1(x) = (f0 (* x (+ real x)))
f2(x) = (/ (/ x real) x)

```

Figure 6: Some learned subroutines for symbolic regression. The system starts with addition, multiplication, division, and real numbers, and learns to build rational functions and polynomials up to 4th order.

4.1 Quantitative Results

We compare with four baselines on held-out tasks:

Ours (no NN), which lesions the recognition model.

RF/DC, which holds the generative model (\mathcal{D}, θ) fixed and learns a recognition model only from samples from the fixed generative model. This is equivalent to our algorithm with $\lambda = \infty$ (Sec. 2.4) and $\mathcal{L}_{\text{RM}} = \mathcal{L}_{\text{HM}}$ (Sec. 2.3). We call this baseline RF/DC because this setup is closest to how RobustFill [3] and DeepCoder [5] are trained. We can not compare directly with these systems, because they are engineered for one specific domain, and do not have publicly available code and datasets.

PCFG, which lesions the recognition model, learns θ , and fixes \mathcal{D} . This is equivalent to ECC with $q(x) = \theta$ and $\lambda = \infty$, and is like learning the parameters of a PCFG while not learning its structure.

Enum, which enumerates a frontier without any learning – equivalently, our first exploration cycle.

For each domain, we are interested both in how many tasks the agent can solve and how quickly it can find those solutions. Tbl. 2 compares our model against these baselines. Our full model consistently improves on the baselines, sometimes dramatically (string editing and symbolic regression). The recognition model consistently increases the number of solved held-out tasks, and lesioning it also slows down the convergence of the algorithm, taking more iterations to reach a given number of tasks solved (Fig. 7). This supports a view of the recognition model as a way of amortizing the cost of searching for programs.

5 Related Work

Our work is far from the first for learning to learn programs, an idea that goes back to Solomonoff [25]:

Deep learning: Much recent work in the ML community has focused on creating neural networks that regress from input/output examples to programs [3, 26, 15, 5]. These neural networks are typically trained with strong supervision (i.e., with annotated ground-truth programs) on massive data sets (i.e., hundreds of millions [3]). Our work considers a weakly-supervised regime where ground truth programs are not provided and the agent must learn from a few hundred tasks.

Inventing new subroutines for program induction: Several program induction algorithms, most prominently the EC algorithm [8], take as their goal to learn new, reusable subroutines that are shared in a multitask setting. We find this work inspiring and motivating, and extend it along two dimensions: (1) we propose a new algorithm for inducing reusable subroutines, based on Fragment Grammars [18]; and (2) we show how to combine these techniques with bottom-up neural recognition models. Other instances of this related idea are [27], Schmidhuber’s OOPS model [28], and predicate invention in ILP [29].

Our work is an instance of Bayesian Program Learning (BPL; see [30, 8, 31, 27]). Previous BPL systems have largely assumed a fixed DSL (but see [27]), and our contribution here is a general way of doing BPL with less hand-engineering of the DSL.

	Ours	Ours (no NN)	RF/DC	PCFG	Enum
<i>Symbolic Regression</i>					
% solved	98%	94%	0.7%	0%	2%
Solve time	2.7s	2.8s	3.6s	–	2.1s
<i>String Editing</i>					
% solved	99%	82%	9%	24%	18%
Solve time	0.9s	2.2s	2.7s	2.3s	2.4s
<i>List functions</i>					
% solved	86%	84%	60%	74%	70%
Solve time	0.8s	0.7s	1.0s	1.0s	1.1s

Table 2: % solved w/ 5 sec timeout. Solve time: averaged over solved tasks. RF/DC: trained like RobustFill/DeepCoder. PCFG: model w/o structure learning. Enum: model w/o any learning.

6 Contribution and Outlook

We contribute an algorithm, ECC, that learns to program by bootstrapping a DSL with new domain-specific primitives that the algorithm itself discovers, together with a neural recognition model that learns how to efficiently deploy the DSL on new tasks. We believe this integration of top-down symbolic representations and bottom-up neural networks – both of them learned – could help make program induction systems more generally useful for AI. Many directions remain open. Two immediate goals are to integrate more sophisticated neural recognition models [3] and program synthesizers [10], which may improve performance in some domains over the generic methods used here. Another direction is to explore DSL meta-learning: Can we find a *single* universal primitive set that could effectively bootstrap DSLs for new domains, including the four domains considered, but also many others?

References

- [1] Stephen H Muggleton, Dianhuan Lin, and Alireza Tamaddon-Nezhad. Meta-interpretive learning of higher-order dyadic datalog: Predicate invention revisited. *Machine Learning*, 100(1):49–73, 2015.
- [2] Sumit Gulwani. Automating string processing in spreadsheets using input-output examples. In *ACM SIGPLAN Notices*, volume 46, pages 317–330. ACM, 2011.
- [3] Jacob Devlin, Jonathan Uesato, Surya Bhupatiraju, Rishabh Singh, Abdel-rahman Mohamed, and Pushmeet Kohli. Robustfill: Neural program learning under noisy i/o. *arXiv preprint arXiv:1703.07469*, 2017.
- [4] Ashwin Kalyan, Abhishek Mohta, Oleksandr Polozov, Dhruv Batra, Prateek Jain, and Sumit Gulwani. Neural-guided deductive search for real-time program synthesis from examples. *ICLR*, 2018.
- [5] Matej Balog, Alexander L Gaunt, Marc Brockschmidt, Sebastian Nowozin, and Daniel Tarlow. Deepcoder: Learning to write programs. *ICLR*, 2016.
- [6] Tuan Anh Le, Atılım Güneş Baydin, and Frank Wood. Inference Compilation and Universal Probabilistic Programming. In *AISTATS*, 2017.

- [7] Geoffrey E Hinton, Peter Dayan, Brendan J Frey, and Radford M Neal. The "wake-sleep" algorithm for unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.
- [8] Eyal Dechter, Jon Malmaud, Ryan P. Adams, and Joshua B. Tenenbaum. Bootstrap learning via modular concept discovery. In *IJCAI*, 2013.
- [9] Benjamin C. Pierce. *Types and programming languages*. MIT Press, 2002.
- [10] Armando Solar Lezama. *Program Synthesis By Sketching*. PhD thesis, 2008.
- [11] Eric Schkufza, Rahul Sharma, and Alex Aiken. Stochastic superoptimization. In *ACM SIGARCH Computer Architecture News*, volume 41, pages 305–316. ACM, 2013.
- [12] John K Feser, Swarat Chaudhuri, and Isil Dillig. Synthesizing data structure transformations from input-output examples. In *PLDI*, 2015.
- [13] Peter-Michael Osera and Steve Zdancewic. Type-and-example-directed program synthesis. In *ACM SIGPLAN Notices*, volume 50, pages 619–630. ACM, 2015.
- [14] Oleksandr Polozov and Sumit Gulwani. Flashmeta: A framework for inductive program synthesis. *ACM SIGPLAN Notices*, 50(10):107–126, 2015.
- [15] Aditya Menon, Omer Tamuz, Sumit Gulwani, Butler Lampson, and Adam Kalai. A machine learning framework for programming by example. In *ICML*, pages 187–195, 2013.
- [16] Geoffrey E Hinton and Ruslan R Salakhutdinov. Reducing the dimensionality of data with neural networks. *Science*, 2006.
- [17] Peter Dayan, Geoffrey E Hinton, Radford M Neal, and Richard S Zemel. The helmholtz machine. *Neural computation*, 7(5):889–904, 1995.
- [18] Timothy J. O’Donnell. *Productivity and Reuse in Language: A Theory of Linguistic Computation and Storage*. The MIT Press, 2015.
- [19] Trevor Cohn, Phil Blunsom, and Sharon Goldwater. Inducing tree-substitution grammars. *JMLR*.
- [20] J.D. Lafferty. *A Derivation of the Inside-outside Algorithm from the EM Algorithm*. Research report.
- [21] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [22] Tessa Lau. *Programming by demonstration: a machine learning approach*. PhD thesis, 2001.
- [23] Rajeev Alur, Dana Fisman, Rishabh Singh, and Armando Solar-Lezama. Sygus-comp 2016: results and analysis. *arXiv preprint arXiv:1611.07627*, 2016.
- [24] Christopher M. Bishop. *Pattern Recognition and Machine Learning*. 2006.
- [25] Ray J Solomonoff. A system for incremental learning based on algorithmic probability. Sixth Israeli Conference on Artificial Intelligence, Computer Vision and Pattern Recognition, 1989.
- [26] Jacob Devlin, Rudy R Bunel, Rishabh Singh, Matthew Hausknecht, and Pushmeet Kohli. Neural program meta-induction. In *NIPS*, 2017.
- [27] Percy Liang, Michael I. Jordan, and Dan Klein. Learning programs: A hierarchical bayesian approach. In *ICML*, 2010.
- [28] Jürgen Schmidhuber. Optimal ordered problem solver. *Machine Learning*, 54(3):211–254, 2004.
- [29] Dianhuan Lin, Eyal Dechter, Kevin Ellis, Joshua B. Tenenbaum, and Stephen Muggleton. Bias reformulation for one-shot function induction. In *ECAI 2014*, 2014.
- [30] Brenden M Lake, Ruslan Salakhutdinov, and Joshua B Tenenbaum. Human-level concept learning through probabilistic program induction. *Science*, 350(6266):1332–1338, 2015.
- [31] Kevin Ellis, Armando Solar-Lezama, and Josh Tenenbaum. Sampling for bayesian program learning. In *Advances in Neural Information Processing Systems*, 2016.