

# DreamCoder: Bootstrapping Domain-Specific Languages for Neurally-Guided Bayesian Program Learning

Kevin Ellis, Lucas Morales, Mathias Sablé Meyer, Maxwell Nye, Luke Hewitt,  
Armando Solar-Lezama, Joshua B. Tenenbaum

Massachusetts Institute of Technology & MIT-IBM Watson AI lab

## Wake/Sleep DSL Induction

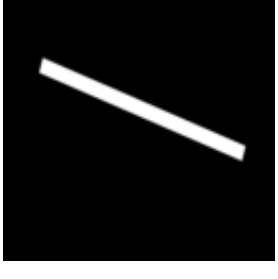



**Domain Specific Language (DSL):** A finely-tuned program representation, specialized to a domain of programming tasks. Prior work in program learning largely uses hand-engineered DSLs.

**Approach:** DREAMCODER algorithm, which bootstraps a learned DSL while jointly training a neural net to search for programs in the learned DSL. Given a few hundred programming tasks, alternately:

- **Wake:** synthesize programs
- **Sleep-R:** train neural net (Recognition model)
- **Sleep-G:** improve DSL (Generative model)


**Program representation:**  
 $\approx$ Lisp; conditionals, variables,  $\lambda$  abstraction

## Model outputs for three different task domains

	List Functions	Text Editing	Symbolic Regression
Programs & Tasks	$[7\ 2\ 3] \rightarrow [7\ 3]$ $[1\ 2\ 3\ 4] \rightarrow [3\ 4]$ $[4\ 3\ 2\ 1] \rightarrow [4\ 3]$ $f(\ell) = (f_1\ \ell\ (\lambda\ (x)\ (>\ x\ 2)))$	$+106\ 769\text{-}438 \rightarrow 106.769.438$ $+83\ 973\text{-}831 \rightarrow 83.973.831$ $f(s) = (f_0\ \text{"."}\ \text{"-"}\ (f_0\ \text{"."}\ \text{" "}\ (\text{cdr}\ s)))$	  $f(x) = (f_1\ x)$ $f(x) = (f_6\ x)$
	$[2\ 7\ 8\ 1] \rightarrow 8$ $[3\ 19\ 14] \rightarrow 19$ $f(\ell) = (f_2\ \ell)$	Temple Anna H $\rightarrow$ TAH Lara Gregori $\rightarrow$ LG $f(s) = (f_2\ s)$	  $f(x) = (f_4\ x)$ $f(x) = (f_3\ x)$
DSL	$f_1(\ell, p) = (\text{foldr}\ \ell\ \text{nil}\ (\lambda\ (x\ a)\ (\text{if}\ (p\ x)\ (\text{cons}\ x\ a)\ a))))$ <i>(f<sub>1</sub>: Higher-order filter function)</i>	$f_0(s, a, b) = (\text{map}\ (\lambda\ (x)\ (\text{if}\ (= x\ a)\ b\ x))\ s)$ <i>(f<sub>0</sub>: substitutes characters)</i>	$f_0(x) = (+\ x\ \text{real})$ $f_1(x) = (f_0\ (*\ \text{real}\ x))$ $f_2(x) = (f_1\ (*\ x\ (f_0\ x)))$ $f_3(x) = (f_0\ (*\ x\ (f_2\ x)))$ $f_4(x) = (f_0\ (*\ x\ (f_3\ x)))$ <i>(f<sub>4</sub>: 4th order polynomial)</i>
	$f_2(\ell) = (\text{foldr}\ \ell\ 0\ (\lambda\ (x\ a)\ (\text{if}\ (>\ a\ x)\ a\ x)))$ <i>(f<sub>2</sub>: Maximum element in list <math>\ell</math>)</i>	$f_1(s, c) = (\text{foldr}\ s\ s\ (\lambda\ (x\ a)\ (\text{cdr}\ (\text{if}\ (= c\ x)\ s\ a))))$ <i>(f<sub>1</sub>: Drop characters from s until c reached)</i>	$f_5(x) = (/ \text{real}\ x)$ $f_6(x) = (f_5\ (f_0\ x))$ <i>(f<sub>6</sub>: rational function)</i>
	$f_3(\ell, k) = (\text{foldr}\ \ell\ (\text{is-nil}\ \ell)\ (\lambda\ (x\ a)\ (\text{if}\ a\ a\ (= k\ x))))$ <i>(f<sub>3</sub>: Whether <math>\ell</math> contains k)</i>	$f_2(s) = (\text{unfold}\ s\ \text{is-nil}\ \text{car}\ (\lambda\ (z)\ (f_1\ z\ \text{" "})))$ <i>(f<sub>2</sub>: Abbreviates words)</i>	

Top: Tasks from three domains we apply our algorithm to, each followed by the programs DREAMCODER discovers for them. Bottom: Several examples from learned DSL. Notice that learned DSL primitives can call each other, and that DREAMCODER rediscovers higher-order functions like **filter** ( $f_1$  under List Functions)

## Bayesian framing



Observe  $N$  **tasks**, written  $\{x_n\}_{n=1}^N$ , each a program synthesis problem.

Solve task  $x_n$  with latent program  $p_n$

**Likelihood model**  $\mathbb{P}[x_n|p_n]$  scores program  $p_n$  on task  $x_n$

Latent **DSL**  $\mathcal{D}$  acts as generative model over programs:  $\mathbb{P}[p_n|\mathcal{D}]$

$$p_n^* = \arg \max_{p_n} \mathbb{P}[x_n|p_n] \mathbb{P}[p_n|\mathcal{D}^*]$$

Wake

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \mathbb{P}[\mathcal{D}] \prod_{n=1}^N \mathbb{P}[x_n|p_n] \mathbb{P}[p_n|\mathcal{D}]$$

Sleep-G

## Neural recognition model

Neural network  $Q(p|x)$  predicts distribution over programs conditioned on tasks. Simple  $Q$ : just predicts probabilities of DSL productions. Goal: learn to invert generative model

$$\min_Q \text{KL}(\mathbb{P}[p|x, \mathcal{D}] || Q(p|x))$$

Sleep-R

Train on two sources of data:

- **Samples (“Dreams”) from DSL:** Unlimited data, but only high-quality if generative model  $\mathcal{D}$  is good. Like Helmholtz Machine’s recognition model training. Loss:

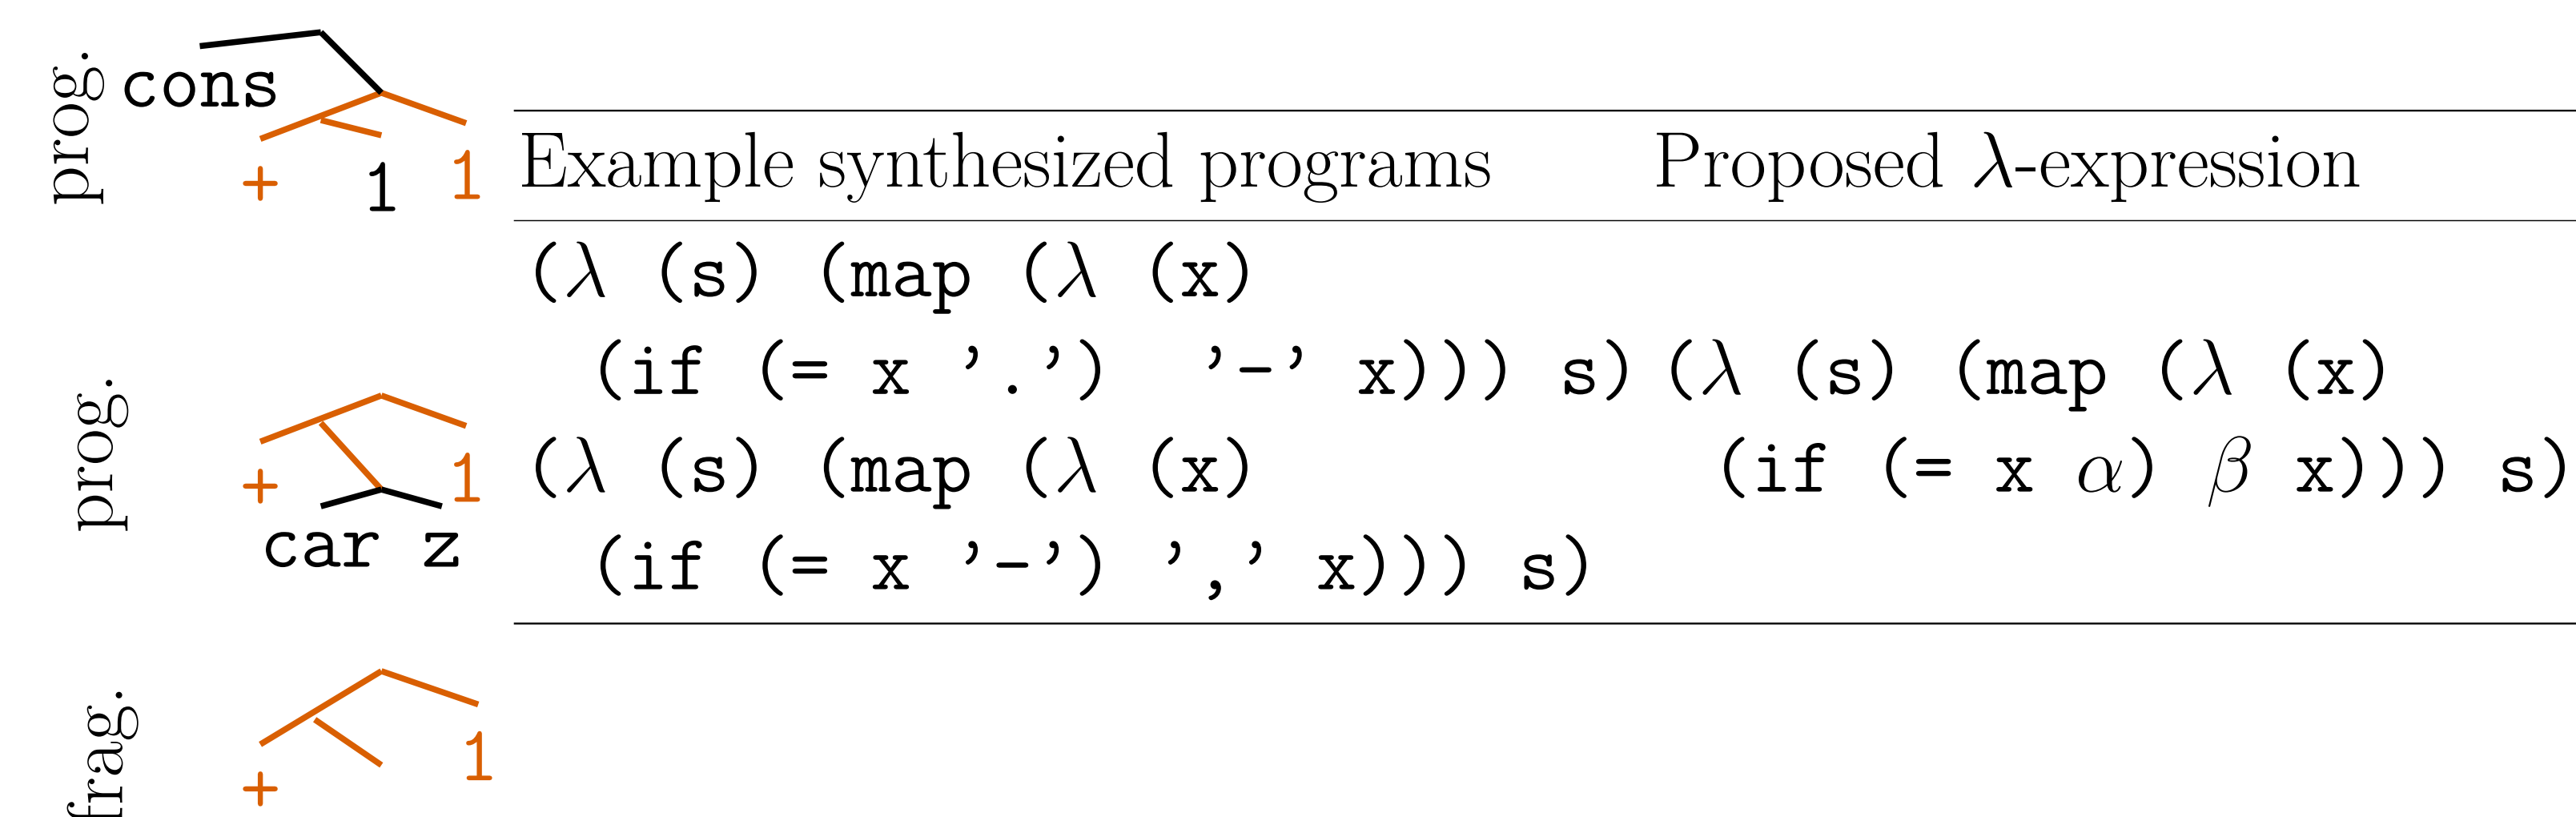
$$\mathbb{E}_{(p,x) \sim \mathcal{D}} [\log Q(p|x)]$$

- **Self-Supervised:**  $(x_n, p_n)$  pairs discovered during waking. Loss:

$$\frac{\mathbb{P}[x_n, p_n|\mathcal{D}]}{\sum_{(x_n, p'_n)} \mathbb{P}[x_n, p'_n|\mathcal{D}]} \log Q(p_n|x_n)$$

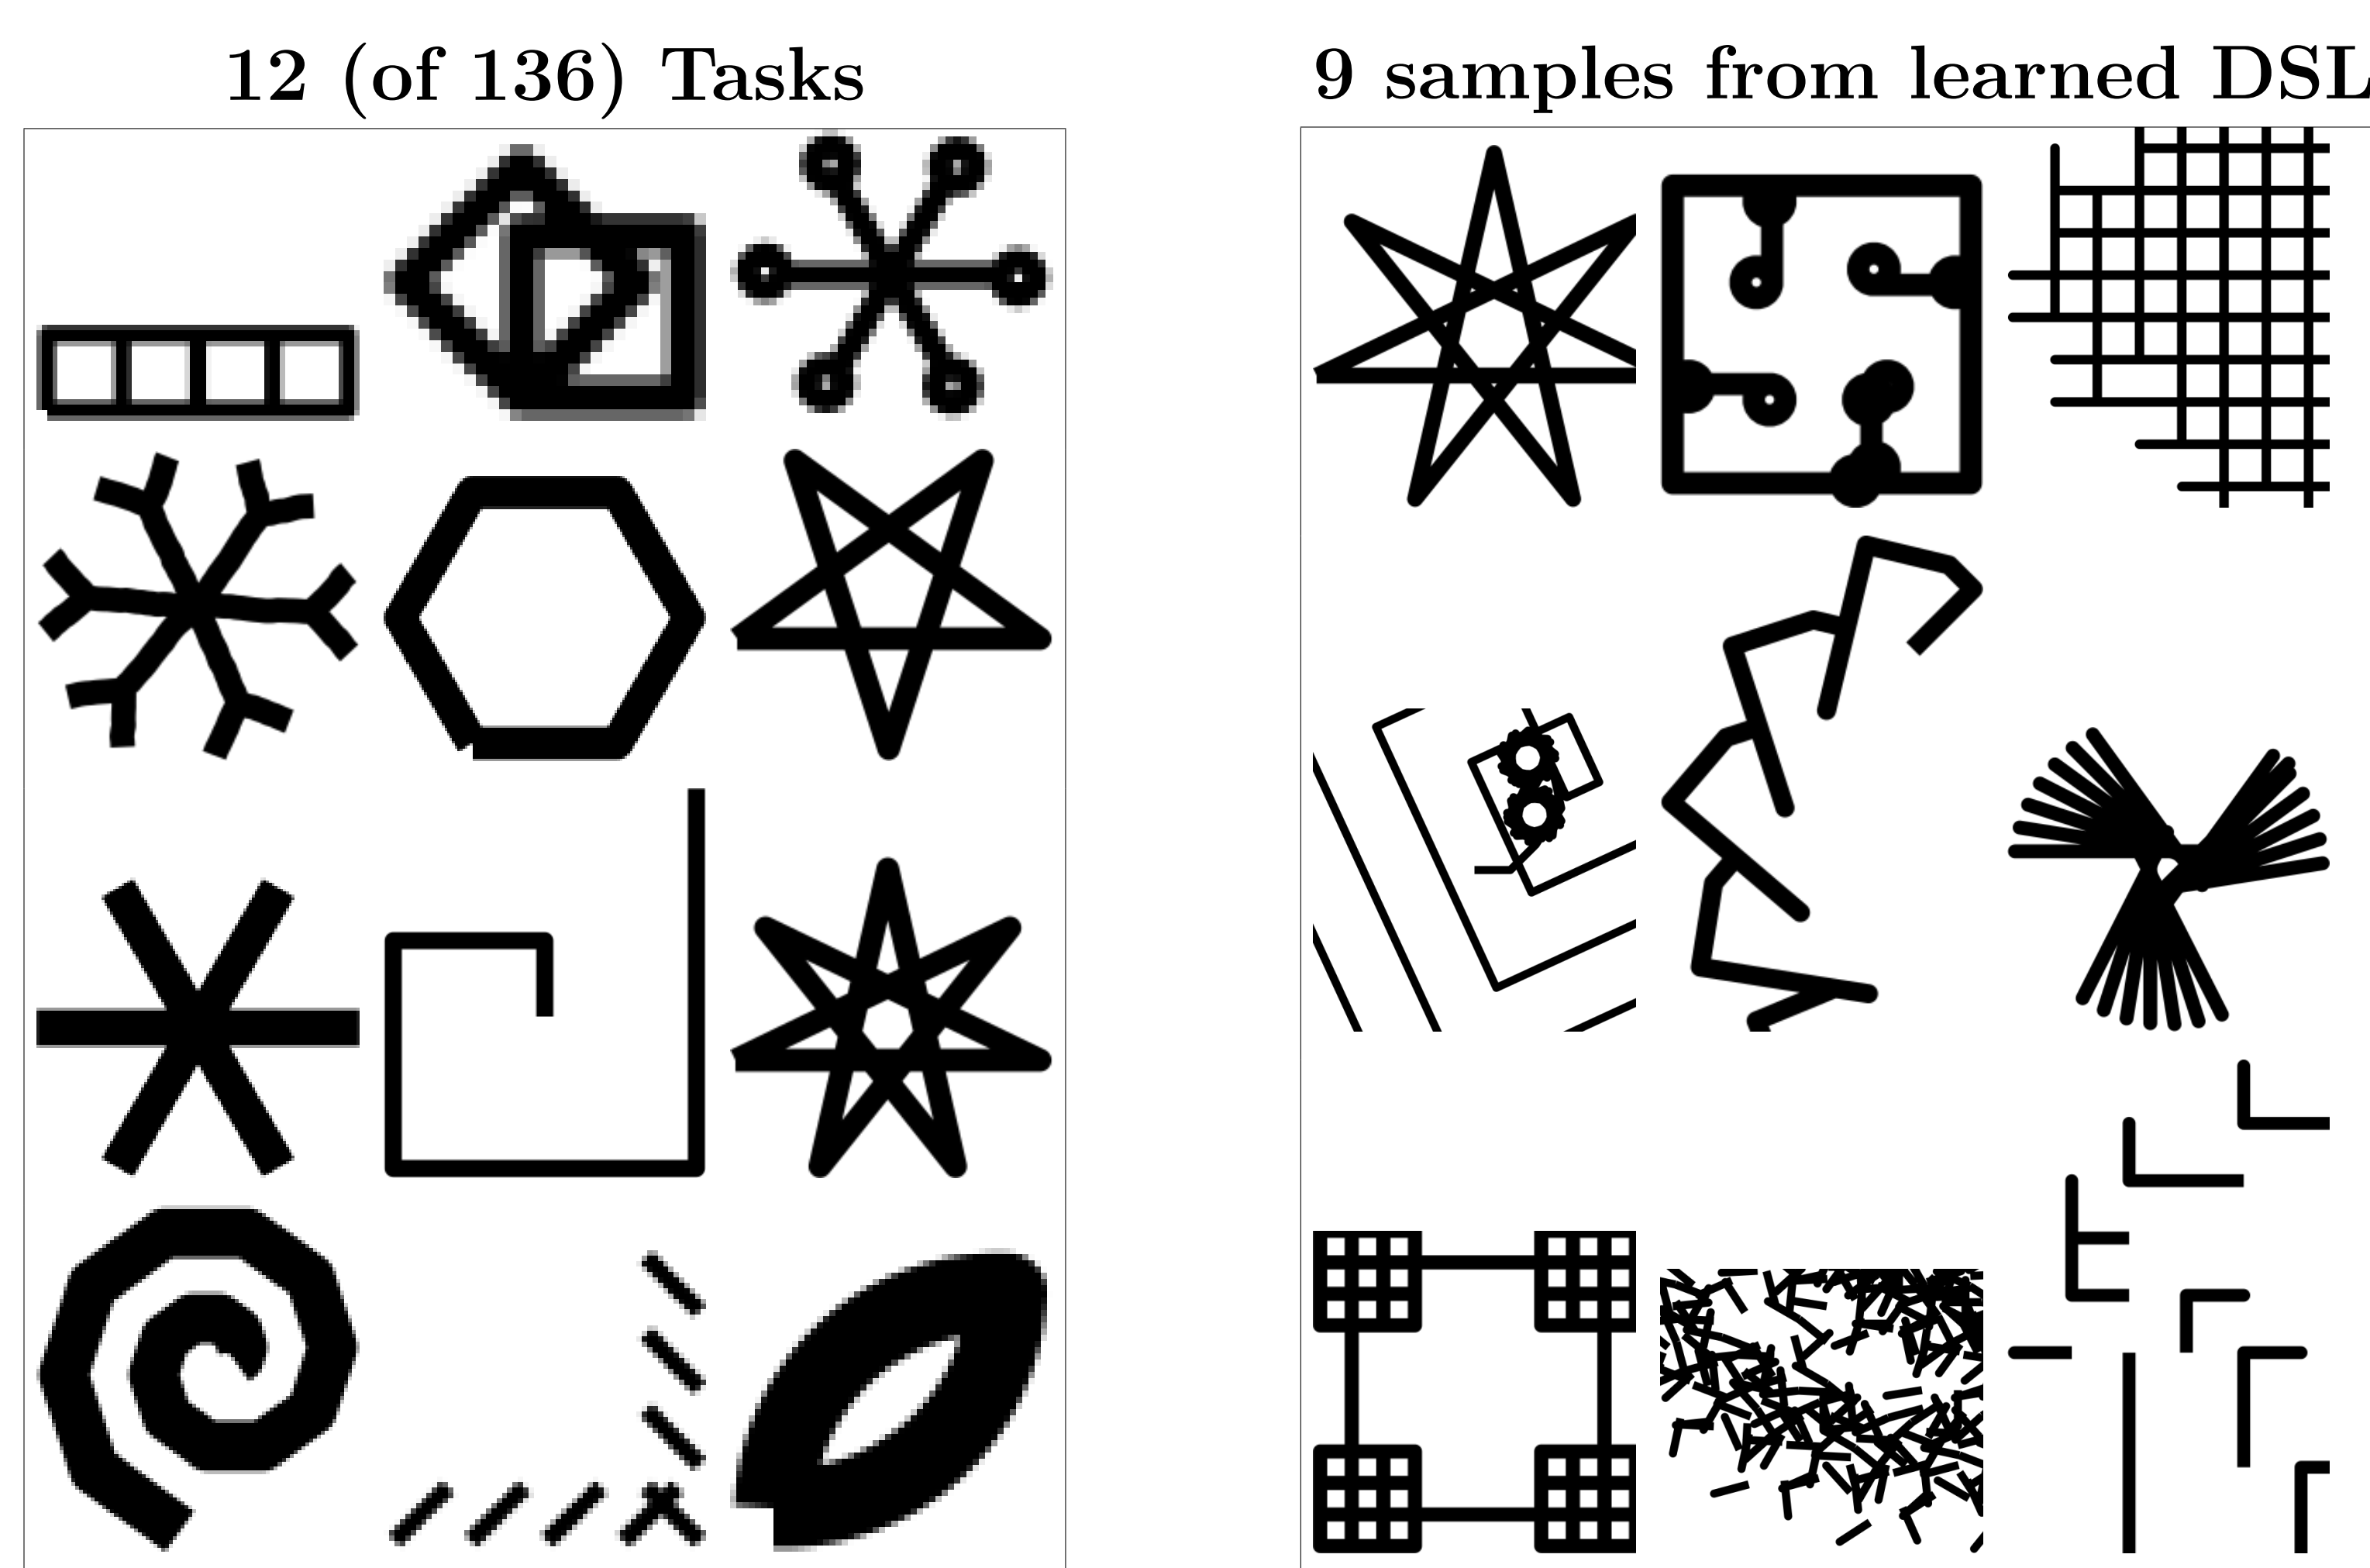
## Fragment Grammars: Inducing a DSL

Fragment grammars: introduced in computational linguistics (O’Donnell 2015)



**Figure 1: Left:** syntax trees of two programs sharing common structure, highlighted in orange, from which we extract a fragment and add it to the DSL (bottom). **Right:** actual programs, from which we extract fragments that perform character substitutions.

## Generative models of images: Turtle/LOGO Graphics

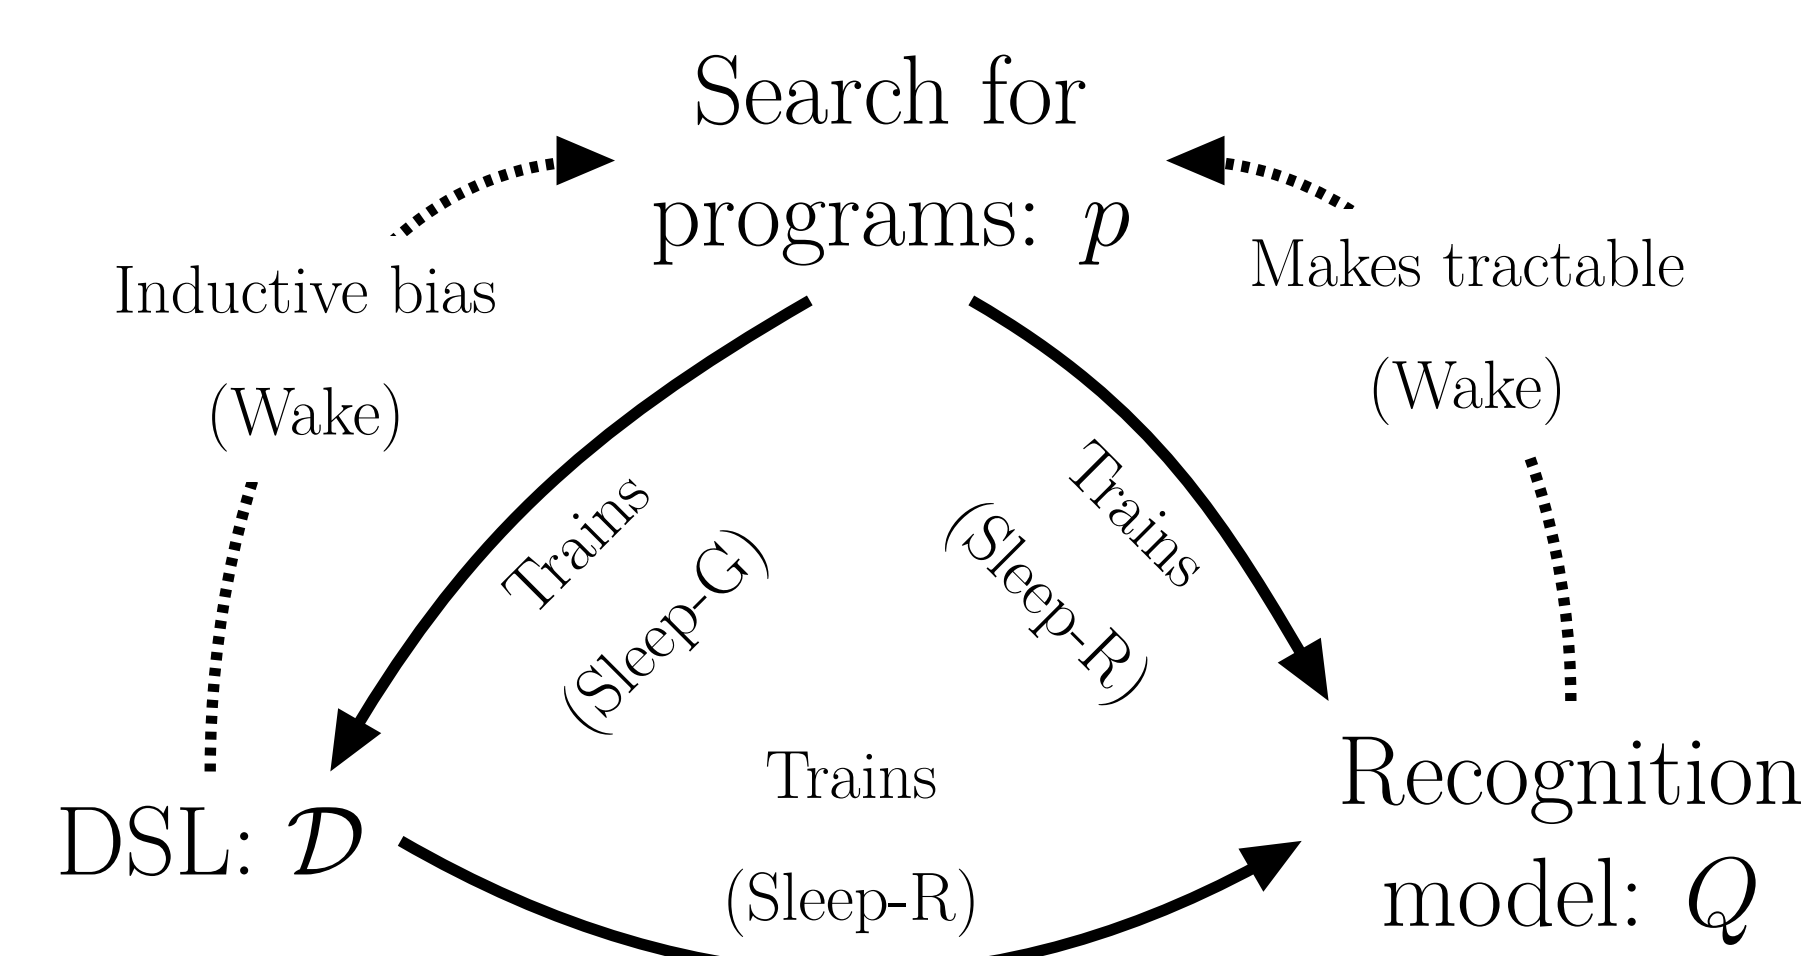


**Left:** Agent controls a ‘pen’ – tasked with drawing pictures. **Right:** During Sleep-R, ‘dream’ by sampling programs from learned DSL and rendering them.

## Acknowledgements

We gratefully acknowledge collaboration with Eyal Dechter, whose EC algorithm (Dechter et al, IJCAI 2013) provided the inspiration for DreamCoder, and Luke Hewitt, who graciously provided us with a regex learning data set.

## Why this works: Bootstrapping



## Generative models of text

Three Tasks		
1.14531	F	110.9
?	CL	163.2
1.29857	F	207.3
?	PCFL	143.3

Three learned generative models  
 $?|(1.\backslash d+)((\backslash u\backslash u)*)|F\backslash d\backslash d\backslash d.\backslash d$

Samples from generative models		
1.61	DQDF	343.8
?	F	241.2
?	F	647.5
1.2	KI	246.8

Learned DSL  
 $f_1 = \backslash u\backslash w*$   $f_2(x) = (x|f_1)*$   
 $f_3 = f_2(\text{" "}) = (\backslash u\backslash w*)*$   
*(f<sub>3</sub>: whitespace delimited words)*  
 $f_4(x) = (x*x)$  *(f<sub>4</sub>: regex ‘plus’)*

## Learning from Scratch

Start w/ McCarthy 1959 Lisp: recursion, conditionals, lists. Train on 22 programming exercises. After 93 hours on 64 CPUs, rediscovers: **map**, **fold**, **zip**, **unfold**, **index**, **length**, **range**, **incr**, **decr**.