

DreamCoder: Bootstrapping Domain-Specific Languages for Neurally-Guided Bayesian Program Learning

Kevin Ellis, Lucas Morales, Mathias Sablé Meyer, Maxwell Nye, Armando Solar-Lezama, Joshua B. Tenenbaum
Massachusetts Institute of Technology & MIT-IBM Watson AI lab

Wake/Sleep DSL Induction


Domain Specific Language (DSL): A finely-tuned program representation, specialized to a domain of programming tasks. Prior work in program learning largely uses hand-engineered DSLs.

Approach: DREAMCODER algorithm, which bootstraps a learned DSL while jointly training a neural net to search for programs in the learned DSL. Given a few hundred programming tasks, alternately:

- **Wake**: synthesize programs
- **Sleep-R**: train neural net (Recognition model)
- **Sleep-G**: improve DSL (Generative model)

Program representation:
 \approx Lisp; conditionals, variables,
 λ abstraction

Bayesian framing



Observe N **tasks**, written $\{x_n\}_{n=1}^N$, each a program synthesis problem.

Solve task x_n with latent program p_n

Likelihood model $\mathbb{P}[x_n|p_n]$ scores
program p_n on task x_n

Latent **DSL** \mathcal{D} acts as generative model over programs: $\mathbb{P}[x|\mathcal{D}]$

$$p_n^* = \arg \max_{p_n} \mathbb{P}[x_n | p_n] \mathbb{P}[p_n | \mathcal{D}^*]$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \underbrace{\mathbb{P}[\mathcal{D}] \prod_n \sum_{p_n} \mathbb{P}[x_n | p_n] \mathbb{P}[p_n | \mathcal{D}]}_{\text{Sleep-G}} \underbrace{\text{Wake}}_{\text{Wake}}$$

Neural recognition model

Neural network $Q(p|x)$ predicts distribution over programs conditioned on tasks. Simple Q : just predicts probabilities of DSL productions. Goal: learn to invert generative model

$$\underbrace{\min_Q \text{KL}(\mathbb{P}[p|x, \mathcal{D}] || Q(p|x))}_{\text{Sleep-R}}$$

Train on two sources of data:


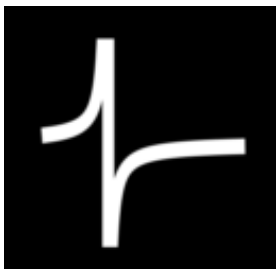

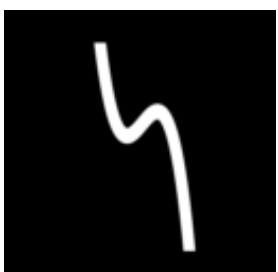
- **Samples (“Dreams”) from DSL:** Unlimited data, but only high-quality if generative model \mathcal{D} is good. Like Helmholtz Machine’s recognition model training. Loss:

$$\mathbb{E}_{(p,x) \sim \mathcal{D}} [\log Q(p|x)]$$

- **Self-Supervised:** (x_n, p_n) pairs discovered during waking. Loss:

$$\frac{\mathbb{P}[x_n, p_n | \mathcal{D}]}{\sum_{(x_n, p'_n)} \mathbb{P}[x_n, p_n | \mathcal{D}]} \log Q(p_n | x_n)$$

Model outputs for three different task domains

	List Functions	Text Editing	Symbolic Regression
Programs & Tasks	$[7\ 2\ 3] \rightarrow [7\ 3]$ $[1\ 2\ 3\ 4] \rightarrow [3\ 4]$ $[4\ 3\ 2\ 1] \rightarrow [4\ 3]$ $f(\ell) = (f_1\ \ell\ (\lambda\ (x)\ (>\ x\ 2)))$ $[2\ 7\ 8\ 1] \rightarrow 8$ $[3\ 19\ 14] \rightarrow 19$ $f(\ell) = (f_2\ \ell)$	$+106\ 769\text{-}438 \rightarrow 106.769.438$ $+83\ 973\text{-}831 \rightarrow 83.973.831$ $f(s) = (f_0\ \text{"." " "-}$ $(f_0\ \text{"." " " "$ $(\text{cdr } s)))$ Temple Anna H \rightarrow TAH Lara Gregori \rightarrow LG $f(s) = (f_2\ s)$	  $f(x) = (f_1\ x)$ $f(x) = (f_6\ x)$   $f(x) = (f_4\ x)$ $f(x) = (f_3\ x)$
	$f_1(\ell, p) = (\text{foldr } \ell\ \text{nil } (\lambda\ (x\ a)\ (\text{if } (p\ x)\ (\text{cons } x\ a)\ a)))$ $(f_1: \text{Higher-order filter function})$	$f_0(s, a, b) = (\text{map } (\lambda\ (x)\ (\text{if } (= x\ a)\ b\ x))\ s)$ $(f_0: \text{substitutes characters})$	$f_0(x) = (+\ x\ \text{real})$ $f_1(x) = (f_0\ (*\ \text{real } x))$ $f_2(x) = (f_1\ (*\ x\ (f_0\ x)))$ $f_3(x) = (f_0\ (*\ x\ (f_2\ x)))$ $f_4(x) = (f_0\ (*\ x\ (f_3\ x)))$ $(f_4: 4\text{th order polynomial})$
	$f_2(\ell) = (\text{foldr } \ell\ 0\ (\lambda\ (x\ a)\ (\text{if } (>\ a\ x)\ a\ x)))$ $(f_2: \text{Maximum element in list } \ell)$	$f_1(s, c) = (\text{foldr } s\ s\ (\lambda\ (x\ a)\ (\text{cdr } (\text{if } (= c\ x)\ s\ a))))$ $(f_1: \text{Drop characters from } s \text{ until } c \text{ reached})$	$f_5(x) = (/ \text{ real } x)$ $f_6(x) = (f_5\ (f_0\ x))$ $(f_6: \text{rational function})$
	$f_3(\ell, k) = (\text{foldr } \ell\ (\text{is-nil } \ell)\ (\lambda\ (x\ a)\ (\text{if } a\ a\ (= k\ x))))$ $(f_3: \text{Whether } \ell \text{ contains } k)$	$f_2(s) = (\text{unfold } s\ \text{is-nil } \text{car } (\lambda\ (z)\ (f_1\ z\ \text{" "})))$ $(f_2: \text{Abbreviates words})$	

Top: Tasks from three domains we apply our algorithm to, each followed by the programs `DREAMCODER` discovers for them. Bottom: Several examples from learned DSL. Notice that learned DSL primitives can call each other, and that `DREAMCODER` rediscovers higher-order functions like `filter` (f_1 under List Functions)

Fragment Grammars: Inducing a DSL

Fragment grammars: introduced in computational linguistics (O'Donnell 2015)




	cons	Example synthesized programs	Proposed λ -expression
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'.'}) \text{'-' } x))) s)$	$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \alpha) \beta x))) s)$
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'-'}) \text{' ,' } x))) s)$	
frag.			

Figure 1: **Left:** syntax trees of two programs sharing common structure, highlighted in orange, from which we extract a fragment and add it to the DSL (bottom). **Right:** actual programs, from which we extract fragments that perform character substitutions.

Ongoing work: Generative models

Learn probabilistic program (a regex) p_n from K strings $x_n = \{y_n^k\}_{k=1}^K$.

Likelihood model:

$$\mathbb{P}[x_n|p_n] = \prod_{k=1}^K \mathbb{P}[y_n^k|p_n]$$

Tasks:		
cut	F	Moss Side
control	CL	Burnage
control	F	City Centre
cut	PCFL	Brooklands

Learned generative models:	
$\backslash 1 * \backslash 1$	$((\backslash u \backslash u) *) F$
$(\backslash \backslash u \backslash w *) *$	

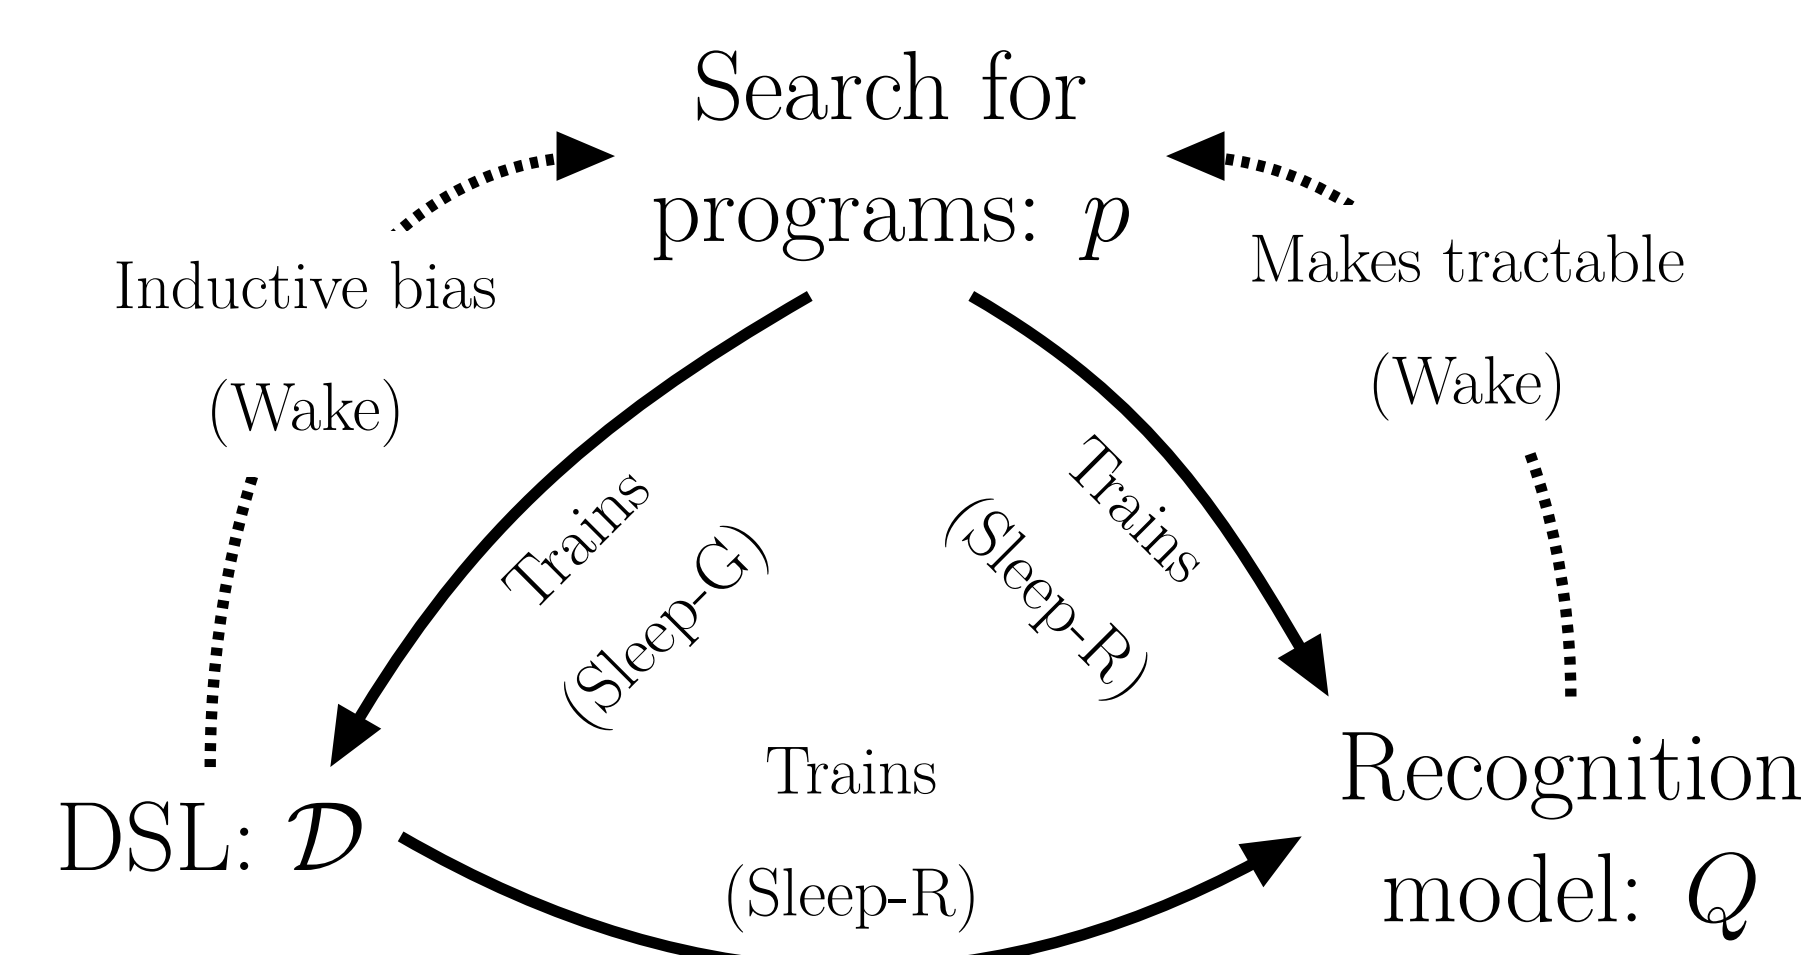
Samples from synthesized generative models:

ya	DQDF	Vr DR
glrwfdenc	F	BeF lKQ
mgs	F	W
piljnl	KI	kqBfZ 0
kj	F	ON
zci	GL	Bttc
sxpm	F	S

Learned DSL:	
$f_1()$	$\backslash u \backslash w *$
$f_2(x)$	$(x f_1) * = (x \backslash u \backslash w *) *$
$f_3(x)$	$f_2(\text{space}) = (\backslash \backslash u \backslash w *) *$
$f_4(x)$	$(x * x)$
<i>(equivalent to regex 'plus')</i>	
$f_5()$	$f_4(\backslash 1) = \backslash 1 * \backslash 1$

Iteration	% Tasks Solved
0	~30
1	~88
2	100
3	100

Why this works: Bootstrapping



- Search finds new programs \implies DSL+recognition model get more data
- DSL improves \implies easier search, recognition model gets better data
- Recognition model improves \implies easier search

Learning from Scratch

Start w/ McCarthy 1959 Lisp: recursion, conditionals, lists. Train on 22 programming exercises. After 93 hours on 64 CPUs, rediscovers 9 functional programming staples: **map**, **fold**, **zip**, **unfold**, **index**, **length**, **range**, **incr**, **decr**.

Acknowledgements

We gratefully acknowledge collaboration with Eyal Dechter, whose EC algorithm (Dechter et al, IJCAI 2013) provided the inspiration for DreamCoder, and Luke Hewitt, who graciously provided us with a regex learning data set.

