

DreamCoder: Bootstrapping Domain-Specific Languages for Neurally-Guided Bayesian Program Learning

Kevin Ellis, Lucas Morales, Mathias Sablé Meyer, Maxwell Nye, Armando Solar-Lezama, Joshua B. Tenenbaum
Massachusetts Institute of Technology

Wake/Sleep DSL Induction


Domain Specific Language (DSL): A finely-tuned program representation, specialized to a domain of programming tasks. Prior work in program learning largely uses hand-engineered DSLs.

Approach: DREAMCODER algorithm, which bootstraps a learned DSL while jointly training a neural net to search for programs in the learned DSL. Given a few hundred programming tasks, alternately:

- **Wake**: synthesize programs
- **Sleep-R**: train neural net (Recognition model)
- **Sleep-G**: improve DSL (Generative model)

Program representation:
 \approx Lisp; conditionals, variables,
 λ abstraction

Bayesian framing



Observe N **tasks**, written $\{x_n\}_{n=1}^N$, each a program synthesis problem.

Solve task x_n with latent program p_n

Likelihood model $\mathbb{P}[x_n|p_n]$ scores
program p_n on task x_n

Latent **DSL** \mathcal{D} acts as generative model over programs: $\mathbb{P}[x|\mathcal{D}]$

$$p_n^* = \arg \max_{p_n} \underbrace{\mathbb{P}[x_n|p_n]}_{\text{likelihood}} \underbrace{\mathbb{P}[p_n|\mathcal{D}^*]}_{\text{prior}}$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \underbrace{\mathbb{P}[\mathcal{D}]}_{\text{Wake}} \prod_n \sum_{p_n} \underbrace{\mathbb{P}[x_n|p_n]\mathbb{P}[p_n|\mathcal{D}]}_{\text{Sleep-G}}$$

Neural recognition model

Neural network $Q(p|x)$ predicts distribution over programs conditioned on tasks. Simple Q : just predicts probabilities of DSL productions. Goal: learn to invert generative model

$$\underbrace{\min_Q \text{KL}(\mathbb{P}[p|x, \mathcal{D}] || Q(p|x))}_{\text{Sleep-R}}$$

Train on two sources of data:

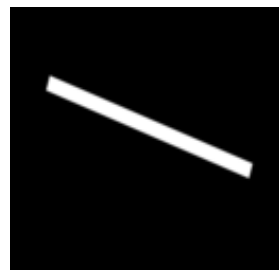
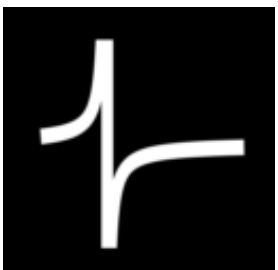
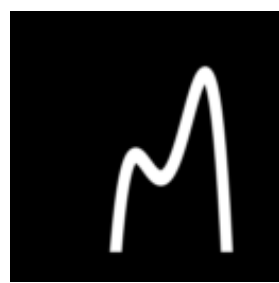
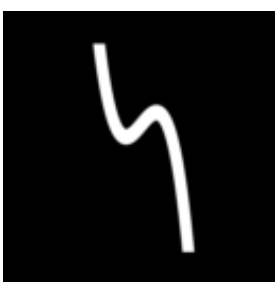
- **Samples (“Dreams”) from DSL:** Unlimited data, but only high-quality if generative model \mathcal{D} is good. Like Helmholtz Machine’s recognition model training. Loss:

$$\mathbb{E}_{(p,x) \sim \mathcal{D}} [\log Q(p|x)]$$

- **Self-Supervised:** (x_n, p_n) pairs discovered during waking. Loss:

$$\frac{\mathbb{P}[x_n, p_n | \mathcal{D}]}{\sum_{(x_n, p'_n)} \mathbb{P}[x_n, p_n | \mathcal{D}]} \log Q(p_n | x_n)$$

Model outputs for three different task domains

	List Functions	Text Editing	Symbolic Regression
Programs & Tasks	$[7\ 2\ 3] \rightarrow [7\ 3]$ $[1\ 2\ 3\ 4] \rightarrow [3\ 4]$ $[4\ 3\ 2\ 1] \rightarrow [4\ 3]$ $f(\ell) = (f_1\ \ell\ (\lambda\ (x)\ (>\ x\ 2)))$	$+106\ 769\text{-}438 \rightarrow 106.769.438$ $+83\ 973\text{-}831 \rightarrow 83.973.831$ $f(s) = (f_0\ \text{"."}\ \text{"-"}\ (f_0\ \text{"."}\ \text{" "}\ (\text{cdr}\ s)))$	  $f(x) = (f_1\ x)\ f(x) = (f_6\ x)$
	$[2\ 7\ 8\ 1] \rightarrow 8$ $[3\ 19\ 14] \rightarrow 19$ $f(\ell) = (f_2\ \ell)$	$[7\ 3] \rightarrow \text{False}$ $[3] \rightarrow \text{False}$ $[9\ 0\ 0] \rightarrow \text{True}$ $[0] \rightarrow \text{True}$ $[0\ 7\ 3] \rightarrow \text{True}$ $f(\ell) = (f_3\ \ell\ 0)$	  $f(x) = (f_4\ x)\ f(x) = (f_3\ x)$
		Temple Anna H \rightarrow TAH Lara Gregori \rightarrow LG $f(s) = (f_2\ s)$	
DSL	$f_1(\ell, p) = (\text{foldr}\ \ell\ \text{nil}\ (\lambda\ (x\ a)\ (\text{if}\ (p\ x)\ (\text{cons}\ x\ a)\ a)))$ <i>(f_1: Higher-order filter function)</i>	$f_0(s, a, b) = (\text{map}\ (\lambda\ (x)\ (\text{if}\ (= x\ a)\ b\ x))\ s)$ <i>(f_0: substitutes characters)</i>	$f_0(x) = (+\ x\ \text{real})$ $f_1(x) = (f_0\ (*\ \text{real}\ x))$ $f_2(x) = (f_1\ (*\ x\ (f_0\ x)))$ $f_3(x) = (f_0\ (*\ x\ (f_2\ x)))$ $f_4(x) = (f_0\ (*\ x\ (f_3\ x)))$ <i>(f_4: 4th order polynomial)</i>
	$f_2(\ell) = (\text{foldr}\ \ell\ 0\ (\lambda\ (x\ a)\ (\text{if}\ (>\ a\ x)\ a\ x)))$ <i>(f_2: Maximum element in list ℓ)</i>	$f_1(s, c) = (\text{foldr}\ s\ s\ (\lambda\ (x\ a)\ (\text{cdr}\ (\text{if}\ (= c\ x)\ s\ a))))$ <i>(f_1: Drop characters from s until c reached)</i>	
	$f_3(\ell, k) = (\text{foldr}\ \ell\ (\text{is-nil}\ \ell)\ (\lambda\ (x\ a)\ (\text{if}\ a\ a\ (= k\ x))))$ <i>(f_3: Whether ℓ contains k)</i>	$f_2(s) = (\text{unfold}\ s\ \text{is-nil}\ \text{car}\ (\lambda\ (z)\ (f_1\ z\ \text{" "})))$ <i>(f_2: Abbreviates words)</i>	$f_5(x) = (/ \text{real}\ x)$ $f_6(x) = (f_5\ (f_0\ x))$ <i>(f_6: rational function)</i>

Top: Tasks from three domains we apply our algorithm to, each followed by the programs `DREAMCODER` discovers for them. Bottom: Several examples from learned DSL. Notice that learned DSL primitives can call each other, and that `DREAMCODER` rediscovers higher-order functions like `filter` (f_1 under List Functions)

Fragment Grammars: Inducing a DSL

Fragment grammars: introduced in computational linguistics (O'Donnell 2015)

	cons	Example synthesized programs	Proposed λ -expression
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'.'}) \text{'-' } x))) s)$	$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \alpha) \beta x))) s)$
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'-'}) \text{' , ' } x))) s)$	
frag.			

Figure 1: **Left:** syntax trees of two programs sharing common structure, highlighted in orange, from which we extract a fragment and add it to the DSL (bottom). **Right:** actual programs, from which we extract fragments that perform character substitutions.

Ongoing work: Generative models

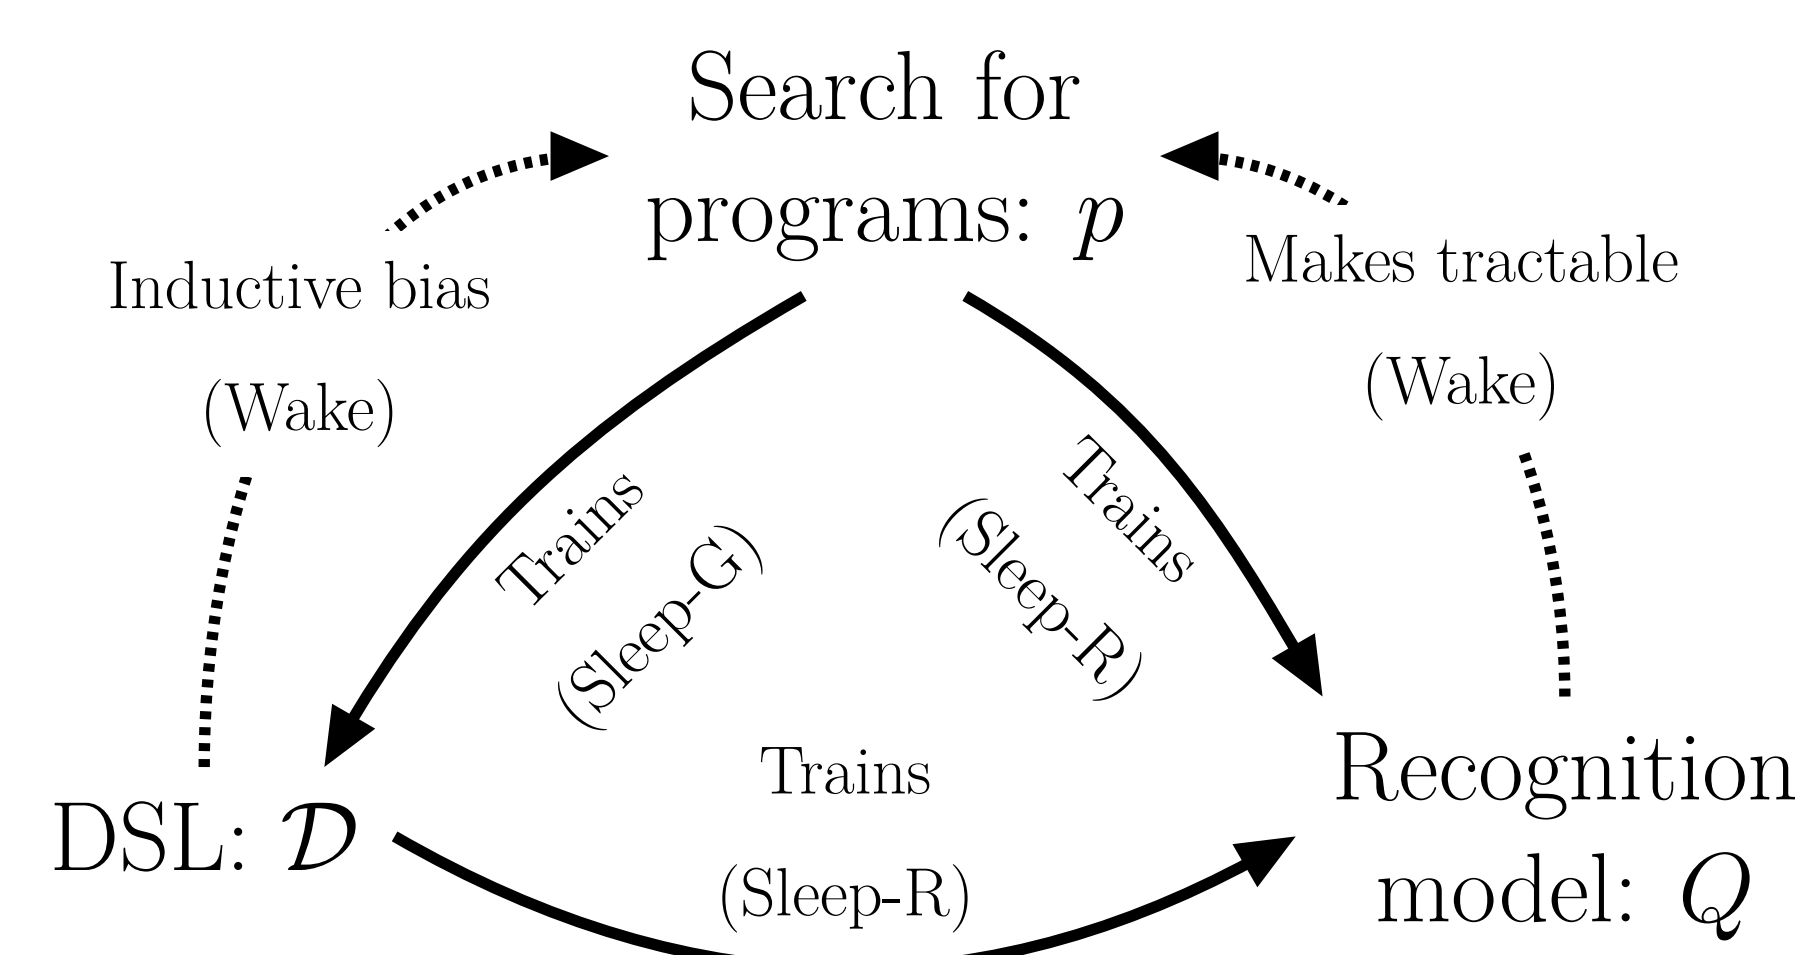
Learn probabilistic program (a regex) p_n from K strings $x_n = \{y_n^k\}_{k=1}^K$.

Likelihood model:

$$\mathbb{P}[x_n|p_n] = \prod_{k=1}^K \mathbb{P}[y_n^k|p_n]$$

Tasks:			Learned DSL:											
cut	F	Moss Side	$f_1() = \backslash u \backslash w *$											
control	CL	Burnage	$f_2(x) = (x f_1) * = (x \backslash u \backslash w *) *$											
control	F	City Centre	$f_3(x) = f_2(\text{space}) = (\quad \backslash u \backslash w *) *$											
cut	PCFL	Brooklands	$f_4(x) = (x * x)$											
Learned generative models:			<i>(equivalent to regex ‘plus’)</i>											
$\backslash 1 * \backslash 1$	$((\backslash u \backslash u) *) F$	$(\quad \backslash u \backslash w *) *$	$f_5() = f_4(\backslash 1) = \backslash 1 * \backslash 1$											
Samples from synthesized generative models:														
ya	DQDF	Vr DR	<table> <caption>Data for % Tasks Solved vs Iteration</caption> <tr> <th>Iteration</th> <th>% Tasks Solved</th> </tr> <tr> <td>0</td> <td>~25</td> </tr> <tr> <td>1</td> <td>~90</td> </tr> <tr> <td>2</td> <td>100</td> </tr> <tr> <td>3</td> <td>100</td> </tr> </table>		Iteration	% Tasks Solved	0	~25	1	~90	2	100	3	100
Iteration	% Tasks Solved													
0	~25													
1	~90													
2	100													
3	100													
glrwfdenc	F	BeF IKQ												
mgs	F	W												
piljnl	KI	kqBfZ 0												
kj	F	ON												
zci	GL	Bttc												
sxpm	F	S												

Why this works: Bootstrapping



- Search finds new programs \implies DSL+recognition model get more data
- DSL improves \implies easier search, recognition model gets better data
- Recognition model improves \implies easier search

Learning from Scratch

Start w/ McCarthy 1959 Lisp: recursion, conditionals, lists. Train on 22 programming exercises. After 93 hours on 64 CPUs, rediscovers 9 functional programming staples: **map**, **fold**, **zip**, **unfold**, **index**, **length**, **range**, **incr**, **decr**.

Acknowledgements

We gratefully acknowledge collaboration with Eyal Dechter, whose EC algorithm (Dechter et al, IJCAI 2013) provided the inspiration for DreamCoder, and Luke Hewitt, who graciously provided us with a regex learning data set.

