

DreamCoder: Bootstrapping Domain-Specific Languages for Neurally-Guided Bayesian Program Learning

Kevin Ellis, Lucas Morales, Mathias Sablé Meyer, Maxwell Nye, Armando Solar-Lezama, Joshua B. Tenenbaum
Massachusetts Institute of Technology

Wake/Sleep DSL Induction


Domain Specific Language (DSL): A finely-tuned program representation, specialized to a domain of programming tasks. Prior work in program learning largely uses hand-engineered DSLs.

Approach: DREAMCODER algorithm, which bootstraps a learned DSL while jointly training a neural net to search for programs in the learned DSL. Given a few hundred programming tasks, alternately:

- **Wake**: synthesize programs
- **Sleep-R**: train neural net (Recognition model)
- **Sleep-G**: improve DSL (Generative model)

Program representation:
 \approx Lisp; conditionals, variables,
 λ abstraction

Bayesian framing



Observe N **tasks**, written $\{x_n\}_{n=1}^N$, each a program synthesis problem.

Solve task x_n with latent program p_n

Likelihood model $\mathbb{P}[x_n|p_n]$ scores
program p_n on task x_n

Latent **DSL** \mathcal{D} acts as generative model over programs: $\mathbb{P}[x|\mathcal{D}]$

$$\underbrace{p_n^* = \arg \max_{p_n} \mathbb{P}[x_n | p_n] \mathbb{P}[p_n | \mathcal{D}^*]}$$

$$\mathcal{D}^* = \arg \max_{\mathcal{D}} \underbrace{\mathbb{P}[\mathcal{D}] \prod_n \sum_{p_n} \mathbb{P}[x_n | p_n] \mathbb{P}[p_n | \mathcal{D}]}_{\text{Sleep-G}} \underbrace{\mathbb{P}[\mathcal{D}]}_{\text{Wake}}$$

Neural recognition model

Neural network $Q(p|x)$ predicts distribution over programs conditioned on tasks. Simple Q : just predicts probabilities of DSL productions. Goal: learn to invert generative model

$$\underbrace{\min_Q \text{KL}(\mathbb{P}[p|x, \mathcal{D}] || Q(p|x))}_{\text{Sleep-R}}$$

Train on two sources of data:

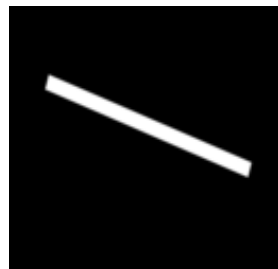
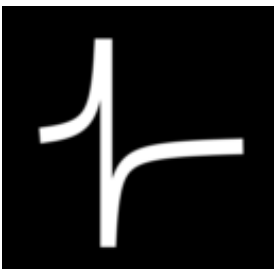

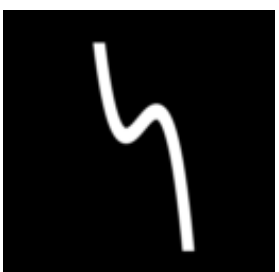
- **Samples (“Dreams”) from DSL:** Unlimited data, but only high-quality if generative model \mathcal{D} is good. Like Helmholtz Machine’s recognition model training. Loss:

$$\mathbb{E}_{(p,x) \sim \mathcal{D}} [\log Q(p|x)]$$

- **Self-Supervised:** (x_n, p_n) pairs discovered during waking. Loss:

$$\frac{\mathbb{P}[x_n, p_n | \mathcal{D}]}{\sum_{(x_n, p'_n)} \mathbb{P}[x_n, p_n | \mathcal{D}]} \log Q(p_n | x_n)$$

Model outputs for three different task domains

	List Functions	Text Editing	Symbolic Regression
Programs & Tasks	$[7\ 2\ 3] \rightarrow [7\ 3]$ $[1\ 2\ 3\ 4] \rightarrow [3\ 4]$ $[4\ 3\ 2\ 1] \rightarrow [4\ 3]$ $f(\ell) = (f_1\ \ell\ (\lambda\ (x)\ (>\ x\ 2)))$ $[2\ 7\ 8\ 1] \rightarrow 8$ $[3\ 19\ 14] \rightarrow 19$ $f(\ell) = (f_2\ \ell)$	$+106\ 769-438 \rightarrow 106.769.438$ $+83\ 973-831 \rightarrow 83.973.831$ $f(s) = (f_0\ \text{"." " "-}$ $(f_0\ \text{"." " " "$ $(cdr\ s)))$ Temple Anna H \rightarrow TAH Lara Gregori \rightarrow LG $f(s) = (f_2\ s)$	  $f(x) = (f_1\ x)$ $f(x) = (f_6\ x)$   $f(x) = (f_4\ x)$ $f(x) = (f_3\ x)$
	$f_1(\ell, p) = (\text{foldr}\ \ell\ \text{nil}\ (\lambda\ (x\ a)\ (\text{if}\ (p\ x)\ (\text{cons}\ x\ a)\ a)))$ <i>(f₁: Higher-order filter function)</i>	$f_0(s, a, b) = (\text{map}\ (\lambda\ (x)\ (\text{if}\ (= x\ a)\ b\ x))\ s)$ <i>(f₀: substitutes characters)</i>	$f_0(x) = (+\ x\ \text{real})$ $f_1(x) = (f_0\ (*\ \text{real}\ x))$ $f_2(x) = (f_1\ (*\ x\ (f_0\ x)))$ $f_3(x) = (f_0\ (*\ x\ (f_2\ x)))$ $f_4(x) = (f_0\ (*\ x\ (f_3\ x)))$ <i>(f₄: 4th order polynomial)</i>
	$f_2(\ell) = (\text{foldr}\ \ell\ 0\ (\lambda\ (x\ a)\ (\text{if}\ (>\ a\ x)\ a\ x)))$ <i>(f₂: Maximum element in list ℓ)</i>	$f_1(s, c) = (\text{foldr}\ s\ s\ (\lambda\ (x\ a)\ (\text{cdr}\ (\text{if}\ (= c\ x)\ s\ a))))$ <i>(f₁: Drop characters from s until c reached)</i>	$f_5(x) = (/ \text{real}\ x)$ $f_6(x) = (f_5\ (f_0\ x))$ <i>(f₆: rational function)</i>
	$f_3(\ell, k) = (\text{foldr}\ \ell\ (\text{is-nil}\ \ell)\ (\lambda\ (x\ a)\ (\text{if}\ a\ a\ (= k\ x))))$ <i>(f₃: Whether ℓ contains k)</i>	$f_2(s) = (\text{unfold}\ s\ \text{is-nil}\ \text{car}\ (\lambda\ (z)\ (f_1\ z\ \text{" "})))$ <i>(f₂: Abbreviates words)</i>	

Top: Tasks from three domains we apply our algorithm to, each followed by the programs `DREAMCODER` discovers for them. Bottom: Several examples from learned DSL. Notice that learned DSL primitives can call each other, and that `DREAMCODER` rediscovers higher-order functions like `filter` (f_1 under List Functions)

Fragment Grammars: Inducing a DSL

Fragment grammars: introduced in computational linguistics (O'Donnell 2015)




	cons	Example synthesized programs	Proposed λ -expression
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'.'}) \text{'-' } x))) s)$	$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \alpha) \beta x))) s)$
prog.		$(\lambda (s) (\text{map } (\lambda (x)$ $(\text{if } (= x \text{'-'}) \text{' ,' } x))) s)$	
frag.			

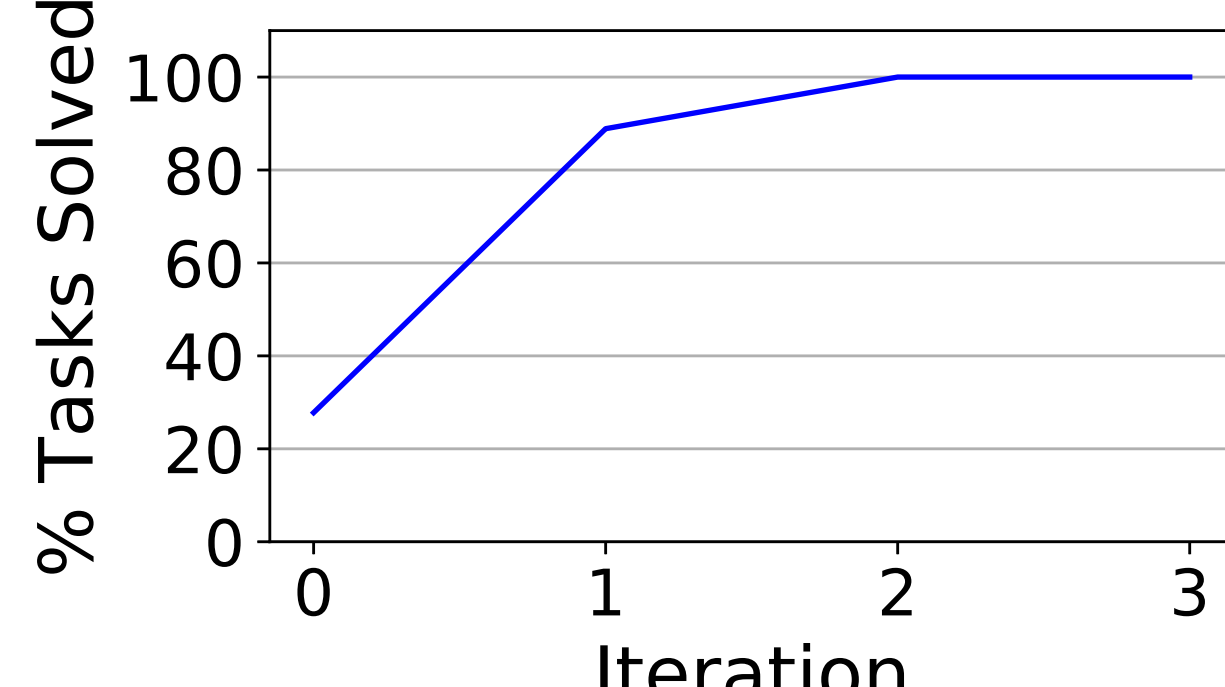
Figure 1: **Left:** syntax trees of two programs sharing common structure, highlighted in orange, from which we extract a fragment and add it to the DSL (bottom). **Right:** actual programs, from which we extract fragments that perform character substitutions.

Ongoing work: Generative models

Learn probabilistic program (a regex) p_n from K strings $x_n = \{y_n^k\}_{k=1}^k$.

Likelihood model:


$$\mathbb{P}[x_n|p_n] = \prod_{k=1}^K \mathbb{P}[y_n^k|p_n]$$

Tasks:			Learned DSL:	
cut	F	Moss Side	$f_1() = \backslash \mathbf{u} \backslash \mathbf{w}^*$	
control	CL	Burnage	$f_2(x) = (x \mid f_1)^* = (x \mid \backslash \mathbf{u} \backslash \mathbf{w}^*)^*$	
control	F	City Centre	$f_3(x) = f_2(\text{space}) = (\mid \backslash \mathbf{u} \backslash \mathbf{w}^*)^*$	
cut	PCFL	Brooklands	$f_4(x) = (x * x)$	
Learned generative models:			<i>(equivalent to regex 'plus')</i>	
$\backslash \mathbf{l} * \backslash \mathbf{l} \ ((\backslash \mathbf{u} \backslash \mathbf{u})^*) \mid \mathbf{F} \ (\mid \backslash \mathbf{u} \backslash \mathbf{w}^*)^*$			$f_5() = f_4(\backslash \mathbf{l}) = \backslash \mathbf{l} * \backslash \mathbf{l}$	
Samples from synthesized generative models:				
ya	DQDF	Vr DR		
glrwfdenc	F	BeF IKQ		
mgs	F	W		
piljnl	KI	kqBfZ 0		
kj	F	ON		
zci	GL	Bttc		
sxpm	F	S		

A line graph showing the percentage of tasks solved over three iterations for the 'Solved' series. The x-axis is labeled 'Iteration' and ranges from 0 to 3. The y-axis is labeled '% Tasks Solved' and ranges from 0 to 100. The data points are approximately: (0, 28), (1, 90), (2, 100), and (3, 100). The line is blue.

Iteration	% Tasks Solved
0	28
1	90
2	100
3	100

Why this works: Bootstrapping



- Search finds new programs \implies DSL+recognition model get more data
- DSL improves \implies easier search, recognition model gets better data
- Recognition model improves \implies easier search

Learning from Scratch

Start w/ McCarthy 1959 Lisp: recursion, conditionals, lists. Train on 22 programming exercises. After 93 hours on 64 CPUs, rediscovers 9 functional programming staples: **map**, **fold**, **zip**, **unfold**, **index**, **length**, **range**, **incr**, **decr**.

Acknowledgements

We gratefully acknowledge collaboration with Eyal Dechter, whose EC algorithm (Dechter et al, IJCAI 2013) provided the inspiration for DreamCoder, and Luke Hewitt, who graciously provided us with a regex learning data set.