

Growing libraries of concepts with wake-sleep program induction

Kevin Ellis

Joint with: Lucas Morales, Mathias Sablé Meyer, Armando Solar-Lezama,
Joshua B. Tenenbaum

Heavy inspiration from: Eyal Dechter

October 2018

MIT

The Language of Thought

Copyrighted Material

The Language of Thought

JERRY A. FODOR

The Language and Thought Series

Jerrold J. Katz

D. Terence Langendoen

George A. Miller

A FORMAL THEORY OF INDUCTIVE INFERENCE, Part I*†

Ray J. Solomonoff

Visiting Professor, Computer Learning Research Center
Royal Holloway, University of London
Mailing Address: P.O.B. 400404, Cambridge, Ma. 02140, U.S.A.

Information and Control, Volume 7, No. 1, Pp. 1-22, March 1964 Copyright
by Academic Press Inc.

Engineering the language of thought



Growing a domain-specific language of thought

Goal: acquire domain-specific knowledge needed to induce a class of programs

Growing a domain-specific language of thought

Goal: acquire domain-specific knowledge needed to induce a class of programs

- Library of concepts (declarative knowledge)
- Inference strategy (procedural knowledge)

DSL: Library of concepts

Tasks and Programs

```
[7 2 3] → [7 3]
[1 2 3 4] → [3 4]
[4 3 2 1] → [4 3]    [7 3] → False
f(ℓ) = (f1 ℓ (λ (x)    [3] → False
    (> x 2)))          [9 0 0] → True
                        [0] → True
                        [0 7 3] → True
                        f(ℓ) = (f3 ℓ 0)
[2 7 8 1] → 8
[3 19 14] → 19
f(ℓ) = (f2 ℓ)
```


DSL

```
f0(ℓ, r) = (foldr r ℓ cons)
           (f0: Append lists r and ℓ)
f1(ℓ, p) = (foldr ℓ nil (λ (x a)
    (if (p x) (cons x a) a)))
           (f1: Higher-order filter function)
f2(ℓ) = (foldr ℓ 0 (λ (x a)
    (if (> a x) a x)))
           (f2: Maximum element in list ℓ)
f3(ℓ, k) = (foldr ℓ (is-nil ℓ)
    (λ (x a) (if a a (= k x))))
           (f3: Whether ℓ contains k)
```

- **Wake:** Solve problems by writing programs
- **Sleep:** Improve DSL and neural recognition model:
 - **Sleep-G:** Improve DSL (**G**enerative model)
 - **Sleep-R:** Improve **R**ecognition model


Combines ideas from Wake-Sleep & Exploration-Compression algorithm by Eyal Dechter

DSL learning as Bayesian inference




[Dechter et al., 2013] [Liang et al, 2010]; [Lake et al, 2015]

DSL learning as **amortized** Bayesian inference

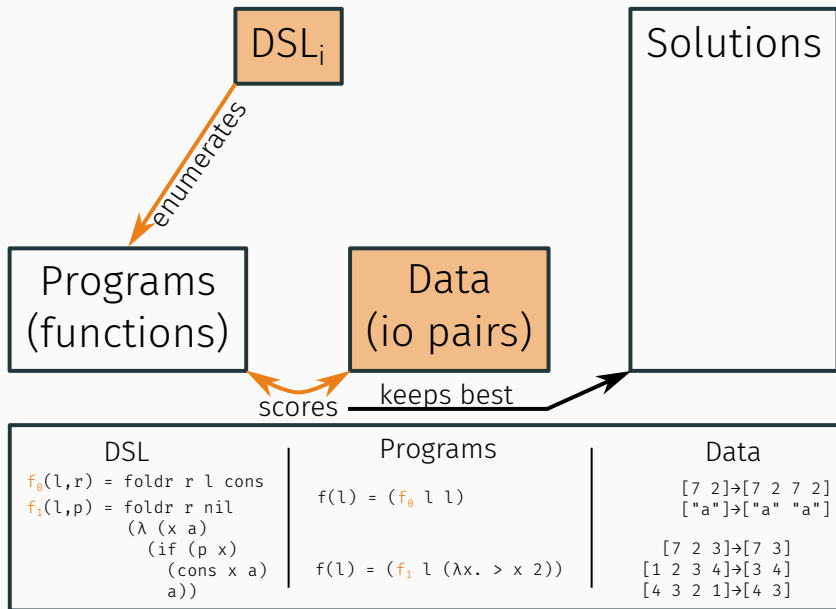


DSL learning as **amortized** Bayesian inference




New: amortized inference +
better program representation (Lisp) +
better DSL inference


Wake — as in Exploration/Compression Algorithm



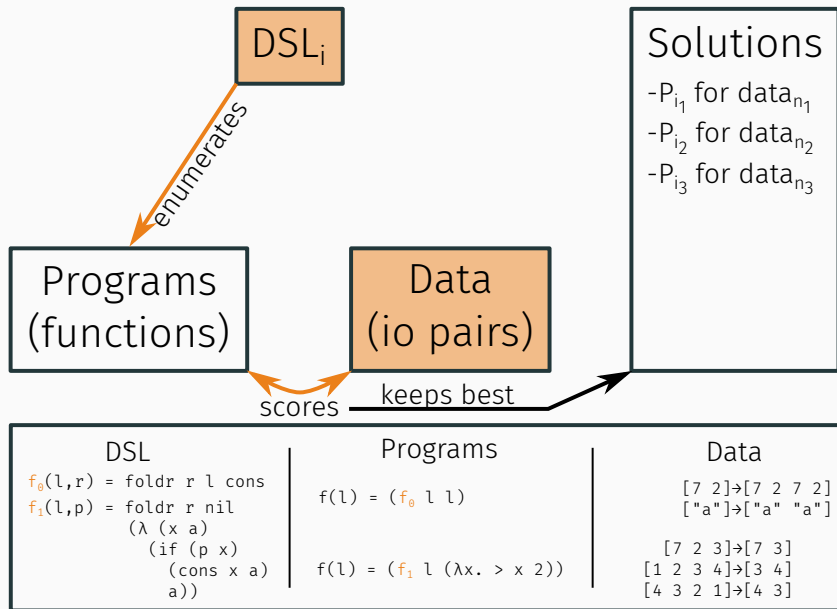
Wake — as in Exploration/Compression Algorithm



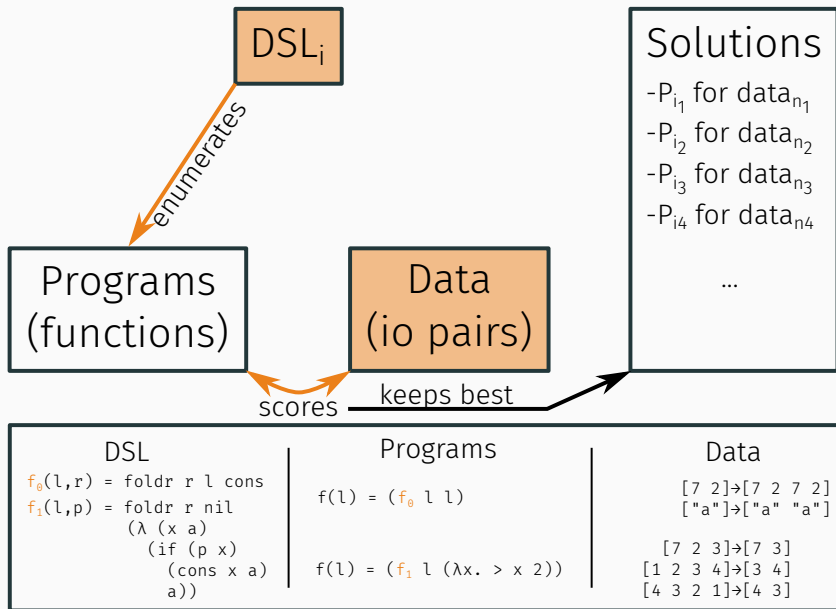
Wake — as in Exploration/Compression Algorithm




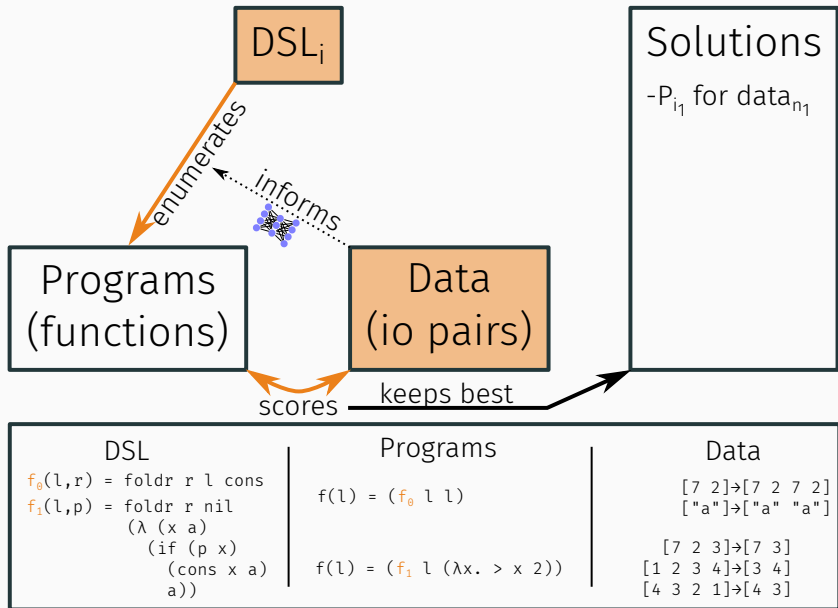
Wake — as in Exploration/Compression Algorithm

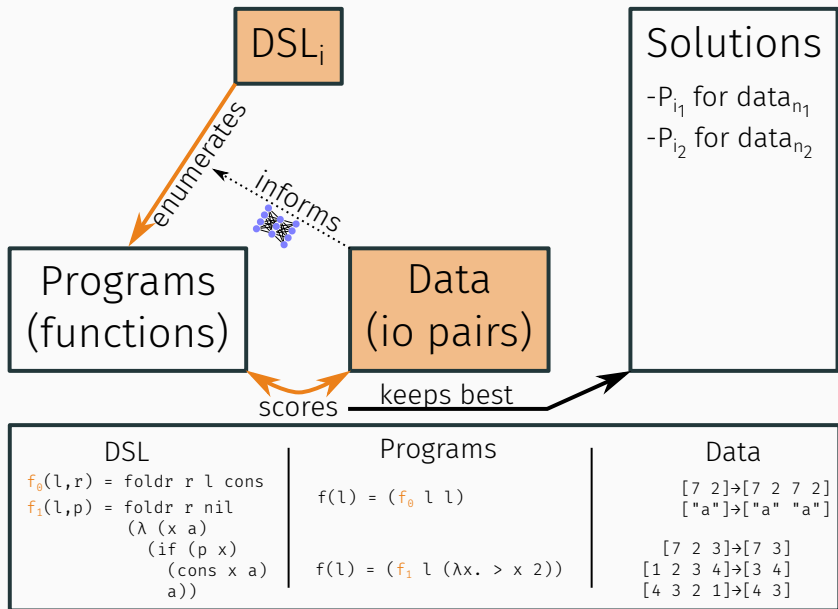


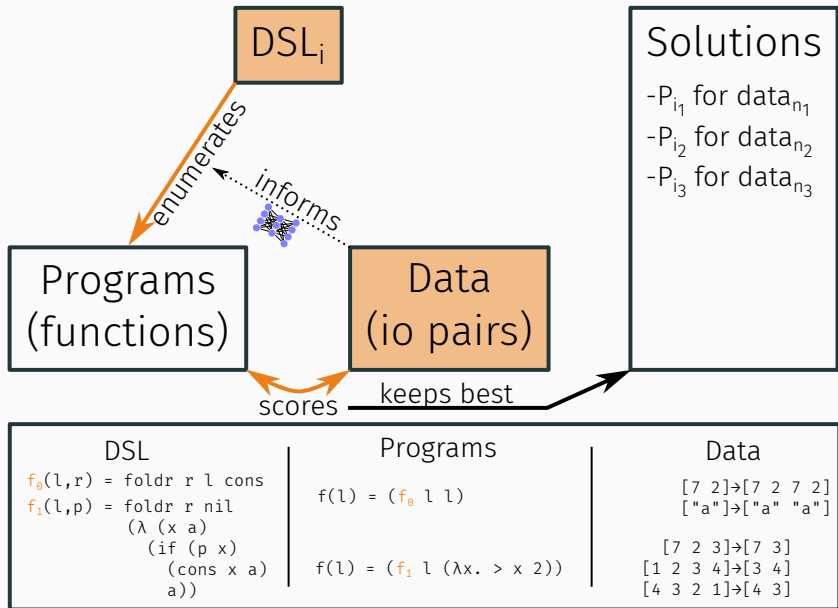
Wake — as in Exploration/Compression Algorithm






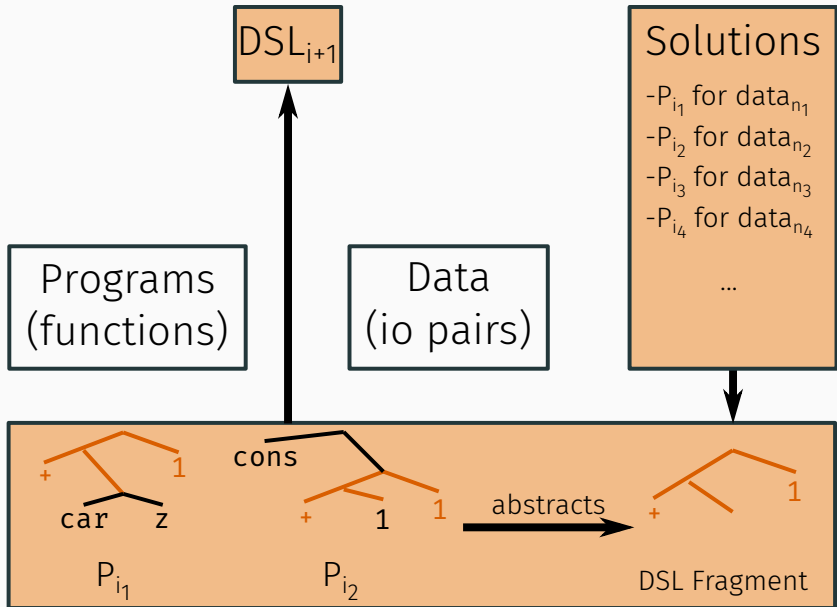








DreamCoder — Sleep-G



Learning higher-order map function

| Task | Program |
|---|---|
| $(1\ 2\ 3) \rightarrow (2\ 4\ 6)$ $(1\ 9\ 2) \rightarrow (2\ 18\ 4)$ | <pre>(Y (λ (r l) (if (nil? l) nil (cons (+ (car l) (car l)) (r (cdr l))))))</pre> |
| $(1\ 2\ 3) \rightarrow (2\ 3\ 4)$ $(1\ 9\ 2) \rightarrow (2\ 10\ 3)$ | <pre>(Y (λ (r l) (if (nil? l) nil (cons (+ (car l) 1) (r (cdr l))))))</pre> |

Learning higher-order `map` function

| Task | Program |
|---|---|
| $(1\ 2\ 3) \rightarrow (2\ 4\ 6)$ $(1\ 9\ 2) \rightarrow (2\ 18\ 4)$ | <pre>(Y (λ (r l) (if (nil? l) nil (cons (+ (car l) (car l)) (r (cdr l))))))</pre> |
| $(1\ 2\ 3) \rightarrow (2\ 3\ 4)$ $(1\ 9\ 2) \rightarrow (2\ 10\ 3)$ | <pre>(Y (λ (r l) (if (nil? l) nil (cons (+ (car l) 1) (r (cdr l))))))</pre> |

```
map = (λ (f) (Y (λ (r l) (if (nil? l) nil
  (cons (f (car l))
    (r (cdr l))))))
```

DreamCoder — Sleep-G (Refactoring)

Learning higher-order map function

| Task | Program |
|---|--|
| $(1\ 2) \rightarrow (2\ 4)$ $(1\ 9) \rightarrow (2\ 18)$ | <pre>((λ (f) (Y (λ (r l) (if (nil? l) nil (cons (f (car l)) (r (cdr l))))))) (λ (z) (+ z z)))</pre> |
| $(1\ 2) \rightarrow (2\ 3)$ $(1\ 9) \rightarrow (2\ 10)$ | <pre>((λ (f) (Y (λ (r l) (if (nil? l) nil (cons (f (car l)) (r (cdr l))))))) (λ (z) (+ z 1)))</pre> |
| map | <pre>(λ (f) (Y (λ (r l) (if (nil? l) nil (cons (f (car l)) (r (cdr l)))))))</pre> |


```
(Y (λ (r l) (if (nil? l) nil  
  (cons (+ (car l) (car l))  
        (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)  
  nil  
  (cons (f (car l))  
        (r (cdr l)))))))  
(λ (z) (+ z z)))
```

```
(Y (λ (r l) (if (nil? l) nil  
  (cons (+ (car l) 1)  
        (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)  
  nil  
  (cons (f (car l))  
        (r (cdr l)))))))  
(λ (z) (+ z 1)))
```

Compress (MDL/Bayes objective)

```
map = (λ (f) (Y (λ (r l) (if (nil? l) nil  
  (cons (f (car l))  
        (r (cdr l)))))))
```

```
(Y (λ (r l) (if (nil? l) nil
  (cons (+ (car l) (car l))
    (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)
  (cons (+ (car l) (car l))
    (r (cdr l))))))
  (λ (z) (+ z z))))
```

10¹⁴ refactorings

```
(Y (λ (r l) (if (nil? l) nil
  (cons (+ (car l) 1)
    (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)
  (cons (+ (car l) 1)
    (r (cdr l))))))
  (λ (z) (+ z 1))))
```

10¹³ refactorings

Compress (MDL/Bayes objective)

```
map = (λ (f) (Y (λ (r l) (if (nil? l) nil
  (cons (f (car l))
    (r (cdr l))))))
```

version space: set of programs
Lau 2003; Gulwani 2012

```
(cons (f (car 1) (car 1))  
      (r (cdr 1))))))
```

refactor

```
((λ (f) (Y (λ (r 1) (if (nil? 1)  
                        1))  
            (λ (z) (+ z z))))
```

10^{14}
refactorings

```
(cons (f (car 1) 1)  
      (r (cdr 1))))))
```

refactor

```
((λ (f) (Y (λ (r 1) (if (nil? 1)  
                        1))  
            (λ (z) (+ z 1))))
```

10^{13}
refactorings

Compress (MDL/Bayes objective)

```
map = (λ (f) (Y (λ (r 1) (if (nil? 1) nil  
                             (cons (f (car 1))  
                                   (r (cdr 1))))))
```

```
(Y (λ (r l) (if (nil? l) nil  
  (cons (+ (car l) (car l))  
        (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)
```

$\leq 10^6$

version spaces

```
(λ (z) (+ z z)))
```

```
(Y (λ (r l) (if (nil? l) nil  
  (cons (+ (car l) 1)  
        (r (cdr l))))))
```

refactor

```
((λ (f) (Y (λ (r l) (if (nil? l)
```

$\leq 10^6$


version spaces

```
(λ (z) (+ z 1)))
```


Compress (MDL/Bayes objective)

```
map = (λ (f) (Y (λ (r l) (if (nil? l) nil  
  (cons (f (car l))  
        (r (cdr l))))))
```

DreamCoder — Sleep-R (Experience Replay)



DreamCoder — Sleep-R (Dreaming)



List functions — Created & investigated by Lucas Morales

| Name | Input | Output |
|--------------------|----------------|---------------|
| repeat-3 | [7 0] | [7 0 7 0 7 0] |
| drop-3 | [0 3 8 6 4] | [6 4] |
| rotate-2 | [8 14 1 9] | [1 9 8 14] |
| count-head-in-tail | [1 2 1 1 3] | 2 |
| keep-div-5 | [5 9 14 6 3 0] | [5 0] |
| product | [7 1 6 2] | 84 |

Discovers 38 concepts, including 'filter'. With different tasks will also learn 'map', 'fold', 'unfold', etc. starting with 1950's Lisp



In the style of FlashFill (Gulwani 2012)

| Text Editing |
|--|
| $+106\ 769-438 \rightarrow 106.769.438$ $+83\ 973-831 \rightarrow 83.973.831$ $f(s) = (f_0\ \text{"."}\ \text{"-"}\ \text{" "}$ $(f_0\ \text{"."}\ \text{" "}\ \text{" "}$ $(\text{cdr}\ s)))$ |
| Temple Anna H \rightarrow TAH Lara Gregori \rightarrow LG $f(s) = (f_2\ s)$ |
| $f_0(s,a,b) = (\text{map}\ (\lambda\ (x)\$ $(\text{if}\ (= x\ a)\ b\ x))\ s)$ (f_0 : <i>Performs character substitution</i>) $f_1(s,c) = (\text{foldr}\ s\ s\ (\lambda\ (x\ a)\$ $(\text{cdr}\ (\text{if}\ (= c\ x)\ s\ a))))$ (f_1 : <i>Drop characters from s until c reached</i>) $f_2(s) = (\text{unfold}\ s\ \text{is-nil}\ \text{car}$ $(\lambda\ (z)\ (f_1\ z\ \text{" "}))$ (f_2 : <i>Abbreviates a sequence of words</i>) |

SyGuS problems: solves 3% before learning, vs 75% after learning.

Best prior work: 80%

List functions & Text editing: Learning curves on hold out tasks



Learning curves for DreamCoder both with (in orange) and without (in teal) the recognition model. Solid lines: % holdout testing tasks solved w/ 10m timeout. Dashed lines: Average solve time, averaged only over tasks that are solved.

Symbolic regression from visual input

Symbolic Regression



$$f(x) = (f_1 \ x)$$



$$f(x) = (f_6 \ x)$$



$$f(x) = (f_4 \ x)$$



$$f(x) = (f_3 \ x)$$

$$f_0(x) = (+ \ x \ \text{real})$$

$$f_1(x) = (f_0 \ (* \ \text{real} \ x))$$

$$f_2(x) = (f_1 \ (* \ x \ (f_0 \ x)))$$

$$f_3(x) = (f_0 \ (* \ x \ (f_2 \ x)))$$

$$f_4(x) = (f_0 \ (* \ x \ (f_3 \ x)))$$

(f_4 : 4th order polynomial)

$$f_5(x) = (/ \ \text{real} \ x)$$

$$f_6(x) = (f_5 \ (f_0 \ x))$$

(f_6 : rational function)

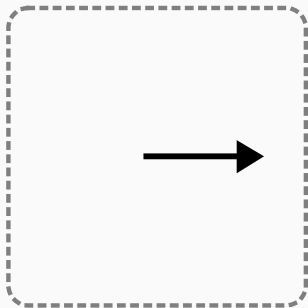
DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks

task : image

FW 1



DSL

OP ::= FW x | RT x | UP | DOWN | SET state

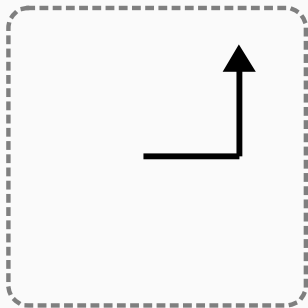
Tasks

task : image

```
FW 1
```

```
RT  $\frac{\pi}{2}$ 
```

```
FW 1
```




DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks

task : image

```
FW 1  
RT  $\frac{\pi}{2}$   
FW 1  
RT  $\frac{\pi}{2}$   
FW 1
```




DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks

task : image

```
for i in range(4)
> FW 1
> RT  $\frac{\pi}{2}$ 
```



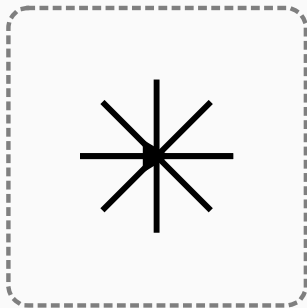
DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks

task : image

```
for i in range(8)
> FW 1
> SET origin
> RT  $\frac{2\pi}{8}$ 
```



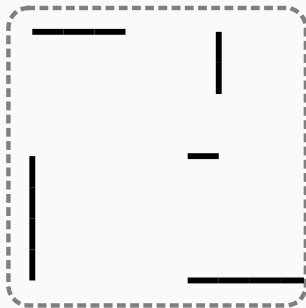
DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks

task : image

```
for i in range(8)
> PU
> FW  $\frac{i}{2}$ 
> PD
> FW  $\frac{i}{2}$ 
> RT  $\frac{\pi}{2}$ 
```




DSL

OP ::= FW x | RT x | UP | DOWN | SET state

Tasks


task : image

```
for i in range( $\infty$ )  
> FW  $\varepsilon$   
> RT  $\varepsilon$ 
```




NUM ::= 1 | π | ∞ | ε | + | - | * | /

Turtle graphics — Training tasks





Turtle graphics — Learning curves




- $\frac{\pi}{2}$ and $\frac{\pi}{4}$ from π , 2, + and /
- A line of length n followed with a right angle
- Loops of length n that uses the number n inside.
- Unit line then teleport back to origin
- ...

Turtle graphics — Dreams

Before training




Plateau 5 minutes




Plateau 2 hours


Turtle graphics — More dreams, 5 minutes, 1st iteration



Turtle graphics — More dreams, 5 minutes, last iteration



Turtle graphics — Dreaming from learned generative model



Turtle graphics — Dreaming from learned generative model




Turtle graphics — Dreaming from learned generative model



Turtle graphics — Dreaming from learned generative model



Turtle graphics — Dreaming from learned generative model




Turtle graphics — Dreaming from learned generative model



Turtle graphics — Dreaming from learned generative model



Turtle graphics — Dreaming from learned generative model




Tower building in blocks world

Control a hand that puts down blocks
(turtle: control a pen that puts down ink)




Tower building in blocks world

Control a hand that puts down blocks
(turtle: control a pen that puts down ink)




Tower building in blocks world: Learned concepts

Parametric planning primitives. Example pyramid concept:



Tower building in blocks world: Learned concepts

Parametric planning primitives. Example brickwall concept:



More human-like machine intelligence

- Acquiring a domain-specific representation (DSL)
- Learning to write programs (recognition model)

DreamCoder: an algorithm for jointly realizing these goals

```
f2(p,f,n,x) = (if (p x) nil
                  (cons (f x) (f2 (n x))))

(f2: unfold)

f3(i,l) = (if (= i 0) (car l)
              (f3 (f1 i) (cdr l)))


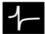


(f3: index)


f4(f,l,x) = (if (empty? l) x
                (f (car l) (f4 (cdr l))))

(f4: fold)

f5(f,l) = (if (empty? l) nil
              (cons (f (car l)) (f5 (cdr l))))

(f5: map)
```

| Symbolic Regression | |
|---|---|
|  |  |
| $f(x) = (f_1 \ x)$ | $f(x) = (f_6 \ x)$ |
|  |  |
| $f(x) = (f_4 \ x)$ | $f(x) = (f_3 \ x)$ |
| $f_0(x) = (+ \ x \ \text{real})$ $f_1(x) = (f_0 \ (* \ \text{real} \ x))$ $f_2(x) = (f_1 \ (* \ x \ (f_0 \ x)))$ $f_3(x) = (f_0 \ (* \ x \ (f_2 \ x)))$ $f_4(x) = (f_0 \ (* \ x \ (f_3 \ x)))$ (f_4 : 4th order polynomial) | |
| $f_5(x) = (/ \ \text{real} \ x)$ $f_6(x) = (f_5 \ (f_0 \ x))$ (f_6 : rational function) | |



More human-like machine intelligence

- Acquiring a domain-specific representation (DSL)
- Learning to write programs (recognition model)

DreamCoder: an algorithm for jointly realizing these goals

```
f2(p,f,n,x) = (if (p x) nil
                  (cons (f x) (f2 (n x))))

(f2: unfold)

f3(i,l) = (if (= i 0) (car l)
              (f3 (f1 i) (cdr l)))


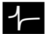


(f3: index)


f4(f,l,x) = (if (empty? l) x
                (f (car l) (f4 (cdr l))))

(f4: fold)

f5(f,l) = (if (empty? l) nil
              (cons (f (car l)) (f5 (cdr l))))

(f5: map)
```

| Symbolic Regression | |
|---|---|
|  |  |
| $f(x) = (f_1 \ x)$ | $f(x) = (f_6 \ x)$ |
|  |  |
| $f(x) = (f_4 \ x)$ | $f(x) = (f_3 \ x)$ |
| $f_0(x) = (+ \ x \ \text{real})$ | |
| $f_1(x) = (f_0 \ (* \ \text{real} \ x))$ | |
| $f_2(x) = (f_1 \ (* \ x \ (f_0 \ x)))$ | |
| $f_3(x) = (f_0 \ (* \ x \ (f_2 \ x)))$ | |
| $f_4(x) = (f_0 \ (* \ x \ (f_3 \ x)))$ | |
| <i>(f4: 4th order polynomial)</i> | |
| $f_5(x) = (/ \ \text{real} \ x)$ | |
| $f_6(x) = (f_5 \ (f_0 \ x))$ | |
| <i>(f6: rational function)</i> | |



The End.